

# **Econometrics Guide**

Xiang Ao

2026-08-10



# Table of contents

|                                                                                       |           |
|---------------------------------------------------------------------------------------|-----------|
| <b>Preface</b>                                                                        | <b>1</b>  |
| How to cite . . . . .                                                                 | 1         |
| <b>1 OLS</b>                                                                          | <b>3</b>  |
| 1.1 Fitness of OLS . . . . .                                                          | 3         |
| 1.2 Example . . . . .                                                                 | 4         |
| 1.3 The Geometry of Least Squares . . . . .                                           | 5         |
| 1.4 Properties of OLS estimator . . . . .                                             | 6         |
| <b>2 The Method of Maximum Likelihood</b>                                             | <b>9</b>  |
| 2.1 Basic Concepts . . . . .                                                          | 9         |
| 2.2 MLE for a linear model . . . . .                                                  | 9         |
| 2.3 Asymptotic Properties of MLE . . . . .                                            | 10        |
| 2.4 Example . . . . .                                                                 | 11        |
| 2.5 Hypotheses Testing . . . . .                                                      | 12        |
| 2.6 Test of Structural Change (Chow test) . . . . .                                   | 14        |
| 2.7 Heteroskedasticity and Autocorrelation Consistent Standard Errors . . . . .       | 18        |
| <b>3 GLS</b>                                                                          | <b>21</b> |
| 3.1 Introduction . . . . .                                                            | 21        |
| 3.2 The GLS Estimator . . . . .                                                       | 21        |
| 3.3 Weighted Least Squares . . . . .                                                  | 22        |
| <b>4 Endogeneity</b>                                                                  | <b>23</b> |
| 4.1 Instrumental Variables (IV) . . . . .                                             | 23        |
| 4.2 Control Function Approach . . . . .                                               | 27        |
| 4.3 Durbin-Wu-Hausman Test . . . . .                                                  | 29        |
| 4.4 Identification . . . . .                                                          | 30        |
| 4.5 Weak Instruments . . . . .                                                        | 31        |
| 4.6 When is endogeneity a problem and what can we say from an IV regression . . . . . | 33        |
| <b>5 General Method of Moments (GMM)</b>                                              | <b>35</b> |
| 5.1 The Method of Moments (MOM) . . . . .                                             | 35        |
| 5.2 OLS as a moment problem . . . . .                                                 | 35        |
| 5.3 IV as a moment problem . . . . .                                                  | 36        |
| 5.4 The Generalized Method of Moments . . . . .                                       | 36        |
| 5.5 GMM . . . . .                                                                     | 37        |
| 5.6 A few concepts of conditioning . . . . .                                          | 41        |
| <b>6 Censored, truncated or selected data</b>                                         | <b>45</b> |
| 6.1 Compare three different cases . . . . .                                           | 45        |

## Table of contents

|           |                                                                              |           |
|-----------|------------------------------------------------------------------------------|-----------|
| 6.2       | An illustration for truncated or censored data . . . . .                     | 45        |
| 6.3       | Truncated Models . . . . .                                                   | 46        |
| 6.4       | Censored Models . . . . .                                                    | 47        |
| 6.5       | Sample Selection . . . . .                                                   | 47        |
| 6.6       | Switching Regression (Treatment-Effects Model) . . . . .                     | 49        |
| <b>7</b>  | <b>Discrete and Limited Dependent Variables</b>                              | <b>51</b> |
| 7.1       | Binary Response Models . . . . .                                             | 51        |
| 7.2       | Models for More than Two Discrete Responses . . . . .                        | 52        |
| <b>8</b>  | <b>Count data models</b>                                                     | <b>55</b> |
| 8.1       | Poisson Model . . . . .                                                      | 55        |
| 8.2       | QMLE Poisson . . . . .                                                       | 56        |
| 8.3       | Fixed-effect Poisson model . . . . .                                         | 57        |
| 8.4       | Negative Binomial model . . . . .                                            | 57        |
| 8.5       | Models for truncated counts . . . . .                                        | 58        |
| 8.6       | The hurdle regression model . . . . .                                        | 58        |
| 8.7       | Zero-inflated count models . . . . .                                         | 59        |
| 8.8       | An example of extra-zero count data model comparison . . . . .               | 59        |
| 8.9       | Interpreting Coefficients and Marginal Effects in Nonlinear Models . . . . . | 60        |
| <b>9</b>  | <b>Panel Data Models</b>                                                     | <b>63</b> |
| 9.1       | Background . . . . .                                                         | 63        |
| 9.2       | Random Effect Methods . . . . .                                              | 64        |
| 9.3       | Fixed Effect Methods . . . . .                                               | 64        |
| 9.4       | Fixed effects or random effects: the Hausman test . . . . .                  | 65        |
| 9.5       | Collapse to unit means, or keep the long format? . . . . .                   | 66        |
| 9.6       | What the long format buys . . . . .                                          | 67        |
| <b>10</b> | <b>Survival Models</b>                                                       | <b>71</b> |
| 10.1      | Survival function . . . . .                                                  | 71        |
| 10.2      | Hazard function . . . . .                                                    | 72        |
| 10.3      | Log-likelihood function . . . . .                                            | 73        |
| 10.4      | Proportional Hazard Models . . . . .                                         | 74        |
| 10.5      | Discrete-time Survival Models . . . . .                                      | 78        |
| <b>11</b> | <b>Dynamic Panel Data</b>                                                    | <b>79</b> |
| 11.1      | When Is It a Problem . . . . .                                               | 79        |
| 11.2      | How Big Is the Bias . . . . .                                                | 79        |
| 11.3      | Anderson and Hsiao estimator . . . . .                                       | 80        |
| 11.4      | Arellano-Bond estimator . . . . .                                            | 80        |
| 11.5      | Blundell and Bond estimator . . . . .                                        | 80        |
| <b>12</b> | <b>Missing data</b>                                                          | <b>83</b> |
| 12.1      | Different cases under different assumptions . . . . .                        | 83        |
| 12.2      | Old methods . . . . .                                                        | 83        |
| 12.3      | Modern methods . . . . .                                                     | 84        |

|                                                                     |            |
|---------------------------------------------------------------------|------------|
| <b>13 Interpreting OLS: Effect Weights</b>                          | <b>87</b>  |
| 13.1 OLS as a weighted average . . . . .                            | 87         |
| 13.2 Deviation slopes $s_i$ versus true effects $b_i$ . . . . .     | 91         |
| 13.3 The general rule: FWL . . . . .                                | 94         |
| 13.4 Two causal decompositions: Angrist and Słoczyński . . . . .    | 96         |
| 13.5 Within versus between, in panels . . . . .                     | 100        |
| 13.6 Summary . . . . .                                              | 106        |
| <b>14 Interpreting OLS: Outcome Weights</b>                         | <b>107</b> |
| 14.1 The other reading of FWL . . . . .                             | 107        |
| 14.2 What the outcome weights look like . . . . .                   | 108        |
| 14.3 When the sign flips . . . . .                                  | 111        |
| 14.4 Out of OLS: the same weights beyond the linear model . . . . . | 116        |
| 14.5 The DiD setting . . . . .                                      | 121        |
| 14.6 Summary . . . . .                                              | 129        |
| <b>References</b>                                                   | <b>131</b> |
| Textbooks and lecture notes . . . . .                               | 131        |
| Core papers . . . . .                                               | 131        |
| Missing data . . . . .                                              | 133        |
| Software . . . . .                                                  | 133        |



# Preface

This is a study guide rather than a textbook. The chapters summarise what I have learned over the years from Davidson and MacKinnon's *Econometric Theory and Methods*, Chris Baum's *An Introduction to Modern Econometrics Using Stata*, and a scatter of papers and lecture notes that helped me make sense of the same material from different angles. The aim is a single place where the workhorse estimators of applied econometrics — OLS, MLE, GLS, IV and GMM, models for censored, discrete, count, panel, and survival data, dynamic panels, and missing-data methods — sit side by side with the assumptions they rely on and the small code snippets needed to fit them.

Examples are written in R or Stata, whichever has the cleaner implementation for the procedure at hand. The code is meant to be short enough to read and re-run rather than a packaged tutorial. For applied causal-inference material — difference-in-differences, matching, doubly robust methods, mediation, causal discovery, and the rest — see the companion books [Introduction to Causal Econometrics](#) and [Topics on econometrics and causal inference](#).

## How to cite

Ao, Xiang. *Econometrics Guide*. <https://xiangao.github.io/econometrics-guide/>.

```
@book{ao_econometrics_guide,  
  author = {Ao, Xiang},  
  title  = {Econometrics Guide},  
  url    = {https://xiangao.github.io/econometrics-guide/}  
}
```



# 1 OLS

A linear model with  $n$  observations and  $k$  regressors can be written as (in vector forms)

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \mathbf{u} \quad (1.1)$$

The idea of least-squares is to choose  $\boldsymbol{\beta}$  to minimize the residual sum of squares ( $SSR$ ),

$$SSR(\boldsymbol{\beta}) = (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})'(\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) = \mathbf{y}'\mathbf{y} - 2\boldsymbol{\beta}'\mathbf{X}'\mathbf{y} + \boldsymbol{\beta}'\mathbf{X}'\mathbf{X}\boldsymbol{\beta} \quad (1.2)$$

The first-order condition to minimize  $SSR$  are:

$$\frac{\partial(SSR)}{\partial\boldsymbol{\beta}} = -2\mathbf{X}'\mathbf{y} + 2\mathbf{X}'\mathbf{X}\boldsymbol{\beta} = \mathbf{0} \quad (1.3)$$

This generates the so-called normal equations

$$(\mathbf{X}'\mathbf{X})\boldsymbol{\beta} = \mathbf{X}'\mathbf{y} \quad (1.4)$$

Therefore, Ordinary Least Squares (OLS) estimator of  $\boldsymbol{\beta}$  is

$$\hat{\boldsymbol{\beta}} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y} \quad (1.5)$$

## 1.1 Fitness of OLS

We can see  $\mathbf{y}$  vector as the part explained by the regression and the unexplained (residual) part,

$$\mathbf{y} = \hat{\mathbf{y}} + \hat{\mathbf{u}} = \mathbf{X}\hat{\boldsymbol{\beta}} + \hat{\mathbf{u}} \quad (1.6)$$

where  $\hat{\mathbf{u}}$  is the OLS residual. Because the residual is orthogonal to the fitted values,  $\hat{\mathbf{y}}'\hat{\mathbf{u}} = \mathbf{0}$ , so the cross term drops:

$$\mathbf{y}'\mathbf{y} = (\hat{\mathbf{y}} + \hat{\mathbf{u}})'(\hat{\mathbf{y}} + \hat{\mathbf{u}}) = \hat{\mathbf{y}}'\hat{\mathbf{y}} + \hat{\mathbf{u}}'\hat{\mathbf{u}} = \hat{\boldsymbol{\beta}}'\mathbf{X}'\mathbf{X}\hat{\boldsymbol{\beta}} + \hat{\mathbf{u}}'\hat{\mathbf{u}} \quad (1.7)$$

Subtracting  $n\bar{y}^2$  ( $\bar{y}$  is the sample mean) from both sides, **and assuming the regression includes an intercept** (so that  $\bar{\hat{y}} = \bar{y}$  and the residuals sum to zero, which is what makes this centered decomposition and the  $R^2$  interpretation below hold),

$$\mathbf{y}'\mathbf{y} - n\bar{y}^2 = (\hat{\boldsymbol{\beta}}'\mathbf{X}'\mathbf{X}\hat{\boldsymbol{\beta}} - n\bar{y}^2) + \hat{\mathbf{u}}'\hat{\mathbf{u}} \quad (1.8)$$

We decompose the total sum of squares into two parts: sum of squares due to error (noise), and sum of squares explained by the linear regression.

## 1 OLS

The  $R^2$  is defined by

$$R^2 = 1 - \frac{SSR}{SST} \quad (1.9)$$

$SST$  is the total (centered) sum of squares of  $y$ ,  $\mathbf{y}'\mathbf{y} - n\bar{y}^2 = \sum_t (\mathbf{y}_t - \bar{y})^2$ .

$R^2$  is simply the proportion of the variation of  $Y$  that can be attributed to the variation of  $X$ .  $R^2$ , however, will never decrease with the addition of any variable to the set of regressors.  $R^2$  stays exactly the same only in the knife-edge case that the added variable is orthogonal to the current residuals; in practice it always rises a little, even for an irrelevant variable. The adjusted  $R^2$ , however, takes account of the addition of any new regressors:

$$\bar{R}^2 = 1 - \frac{SSR/(n-k)}{SST/(n-1)} \quad (1.10)$$

## 1.2 Example

Let's look at an example of OLS of car's mpg on disp, hp and wt.

```
lm1 <- lm(mpg~disp+hp+wt, data=mtcars)
summary(lm1)
```

Call:

```
lm(formula = mpg ~ disp + hp + wt, data = mtcars)
```

Residuals:

|  | Min    | 1Q     | Median | 3Q    | Max   |
|--|--------|--------|--------|-------|-------|
|  | -3.891 | -1.640 | -0.172 | 1.061 | 5.861 |

Coefficients:

|             | Estimate  | Std. Error | t value | Pr(> t )    |
|-------------|-----------|------------|---------|-------------|
| (Intercept) | 37.105505 | 2.110815   | 17.579  | < 2e-16 *** |
| disp        | -0.000937 | 0.010350   | -0.091  | 0.92851     |
| hp          | -0.031157 | 0.011436   | -2.724  | 0.01097 *   |
| wt          | -3.800891 | 1.066191   | -3.565  | 0.00133 **  |

---


Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.639 on 28 degrees of freedom

Multiple R-squared: 0.8268, Adjusted R-squared: 0.8083

F-statistic: 44.57 on 3 and 28 DF, p-value: 8.65e-11

### 1.3 The Geometry of Least Squares



For simplicity, let's assume there are two explanatory variables  $x_1$  and  $x_2$ .  $x_1$  and  $x_2$  form a plane. What OLS does is to project  $y$  onto this plane.

In mathematical terms, all linear combinations of these two vectors define a two-dimensional subspace of Euclidean space  $\mathbf{E}^n$ . This is called the column space of  $\mathbf{X}$ . The least-square principle is to choose  $\beta$  to make  $\hat{y}$ , which belongs to the subspace of  $\mathbf{X}$ , as close as possible to  $y$ .

When we estimate a linear regression model, we map the  $y$  into a vector of fitted values  $\mathbf{X}\hat{\beta}$  and a vector of residuals  $\hat{\mathbf{u}} = \mathbf{y} - \mathbf{X}\hat{\beta}$ . Geometrically, these are examples of orthogonal projections. A projection is a mapping that takes each point of  $\mathbf{E}^n$  into a point in a subspace of  $\mathbf{E}^n$ . An orthogonal projection maps any point into the point of the subspace that is closest to it.

An orthogonal projection can be performed by premultiplying the vector to be projected by a projection matrix. In the case of OLS, the two projection matrices that yield the vector of fitted values and the vector of residuals, are

$$\mathbf{P}_{\mathbf{X}} = \mathbf{X}(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}' \quad (1.11)$$

$$\mathbf{M}_{\mathbf{X}} = \mathbf{I} - \mathbf{P}_{\mathbf{X}} = \mathbf{I} - \mathbf{X}(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}' \quad (1.12)$$

## 1 OLS

From this, we see that the effects of two projection matrices.

$$\mathbf{P}_X \mathbf{y} = \mathbf{X} \hat{\boldsymbol{\beta}} = \hat{\mathbf{y}} \quad (1.13)$$

$$\mathbf{M}_X \mathbf{y} = (\mathbf{I} - \mathbf{P}_X) \mathbf{y} = \mathbf{y} - \mathbf{P}_X \mathbf{y} = \mathbf{y} - \mathbf{X} \hat{\boldsymbol{\beta}} = \hat{\mathbf{u}} \quad (1.14)$$

That is,  $\mathbf{P}_X \mathbf{y}$  projects  $\mathbf{y}$  onto  $\mathbf{X}$  and makes it  $\hat{\mathbf{y}}$ ;  $\mathbf{M}_X \mathbf{y}$  makes it  $\hat{\mathbf{u}}$ .

In the picture, that corresponds to the fact that the projection of  $\mathbf{y}$  onto the plane of  $\mathbf{X}$  creates two parts:  $\hat{\mathbf{y}}$  and  $\hat{\mathbf{u}}$ .

## 1.4 Properties of OLS estimator

### 1.4.1 Unbiasedness

In a linear model

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \mathbf{u}, \quad (1.15)$$

the OLS estimator is

$$\hat{\boldsymbol{\beta}} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y} \quad (1.16)$$

Since  $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \mathbf{u}$ ,

$$\hat{\boldsymbol{\beta}} = \boldsymbol{\beta} + (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{u}. \quad (1.17)$$

This makes

$$\mathbf{E}(\hat{\boldsymbol{\beta}}|\mathbf{X}) = \boldsymbol{\beta} + (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{E}(\mathbf{u}|\mathbf{X}). \quad (1.18)$$

The condition that makes the OLS estimator unbiased is:

$$\mathbf{E}(\mathbf{u}|\mathbf{X}) = \mathbf{0}, \quad (1.19)$$

that is, all explanatory variables which form the columns of  $\mathbf{X}$  are exogenous. This condition is weaker than the independence condition that  $u$  and  $X$  are independent. This says that given  $\mathbf{X}$ , the expected value of  $\mathbf{u}$  is zero; it implies that the model is correctly specified. That is,  $\mathbf{y}$  is a linear function of  $\mathbf{X}$ .

In the context of cross-sectional data, this assumption is plausible. However, when we have time series data, the assumption becomes strong, because it assumes that the entire series of  $\mathbf{X}$  has no relationship with the error term. In a time series context, this is hard to satisfy. The OLS estimator is biased if this condition is not satisfied.

For example, suppose we have a model

$$y_t = \beta_1 + \beta_2 y_{t-1} + u_t, \quad u_t \sim \text{IID}(0, \sigma^2). \quad (1.20)$$

In this simple model, even if we assume that  $y_{t-1}$  and  $u_t$  are uncorrelated, OLS estimator is still biased. That is because  $\mathbf{E}(\mathbf{u}|\mathbf{X}) = \mathbf{0}$  is not satisfied:  $y_{t-1}$  depends on  $u_{t-1}$ ,  $u_{t-2}$  and so on.

There is a weaker condition – weaker than the strict exogeneity above, and in fact the weakest of the conditions in this chapter:

$$E(\mathbf{X}\mathbf{u}) = \mathbf{0}, \quad (1.21)$$

Once we know that the model is correctly specified, then this equation can be used to derive results, such as GMM estimators.

### 1.4.2 Consistency

For OLS estimator to be consistent, a much weaker condition is needed:

$$E(u_t | X_t) = 0, \quad (1.22)$$

This condition is much weaker since it only assumes that the mean of current error term does not depend on the current predictors. Even a model with lagged dependent variable can easily satisfy this condition. This condition is called **contemporaneous exogeneity**. It is the weakest of three exogeneity conditions worth keeping apart, which is precisely why it is the one to invoke for consistency:

- **Strict exogeneity**,  $E[u_t | X_1, \dots, X_n] = 0$ : the error in every period is mean-independent of the regressors in *every* period. This is what unbiasedness requires, and what the lagged-dependent-variable example above violates.
- **Predeterminedness**,  $E[u_t | \mathcal{F}_{t-1}] = 0$  with  $X_t$  measurable with respect to  $\mathcal{F}_{t-1}$  – equivalently,  $X_t$  is orthogonal to the current and all *future* errors, while being allowed to correlate with *past* errors. This is the relevant condition in models with a lagged dependent variable.
- **Contemporaneous exogeneity**,  $E[u_t | X_t] = 0$ .
- **Contemporaneous orthogonality**,  $E[X_t u_t] = 0$  – mere zero correlation. This is weakest of all, and it is what consistency actually needs: together with a rank condition and a law of large numbers it gives  $\text{plim } n^{-1} \mathbf{X}'\mathbf{u} = \mathbf{0}$  and hence  $\text{plim } \hat{\beta} = \beta$ . Contemporaneous *exogeneity* is a strictly stronger, sufficient condition – it constrains the whole conditional mean rather than a single covariance – so the  $E[\mathbf{X}\mathbf{u}] = \mathbf{0}$  introduced under unbiasedness above sits at the *bottom* of this ladder, not between the others.

Each condition implies the next, because conditioning on a larger information set is a more restrictive requirement, not a less restrictive one. In particular, by the law of iterated expectations,

$$E[u_t | X_t] = E[E[u_t | \mathcal{F}_{t-1}] | X_t] = 0, \quad (1.23)$$

so predeterminedness delivers contemporaneous exogeneity but not conversely. In the time series example, then, the OLS estimator is biased, but can still be consistent, if we are willing to assume no contemporaneous correlation.

Note: general linear hypothesis testing, the Chow test for structural change, and heteroskedasticity- and autocorrelation-consistent (HAC) standard errors (White, Newey-West) are all standard OLS extensions, but are covered in the [Maximum Likelihood chapter](#) alongside this guide's other estimation-inference material rather than repeated here.



## 2 The Method of Maximum Likelihood

If we make the assumption that the error terms are normally distributed, the maximum likelihood estimators (MLE) coincide with the least square estimators. But MLE can also be applied to a wide variety of models other than regression models, and it generally generates estimators with excellent asymptotic properties. The major disadvantage of MLE is that it requires stronger distributional assumptions than does the method of moments.

### 2.1 Basic Concepts

Maximum likelihood estimation is to find the set of parameters which makes the current sample most likely given the statistical model.

Suppose we have a sample of  $n$  independent and identically distributed (IID) observations. Each of them comes from some distribution function  $f(\cdot, \boldsymbol{\theta})$ , where parameters  $\boldsymbol{\theta}$  is fixed. The joint distribution of all the data points would be

$$f(x_1, x_2, \dots, x_n | \boldsymbol{\theta}) = f(x_1 | \boldsymbol{\theta}) \times f(x_2 | \boldsymbol{\theta}) \times \dots \times f(x_n | \boldsymbol{\theta}). \quad (2.1)$$

It is referred to as the likelihood function of the model for the given data set.

Then it's natural to estimate  $\boldsymbol{\theta}$  by picking the  $\boldsymbol{\theta}$  that maximize this joint distribution function. A parameter vector  $\hat{\boldsymbol{\theta}}$  at which the likelihood takes on its maximum value is called a maximum likelihood estimate, or MLE, of the parameters.

Usually the log form of the likelihood function is preferred for computational purposes.

### 2.2 MLE for a linear model

Consider the classical normal linear model:

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \mathbf{u}, \quad \mathbf{u} \sim \mathbf{N}(\mathbf{0}, \sigma^2 \mathbf{I}) \quad (2.2)$$

In this case, we have two parameters to estimate,  $\sigma$  and  $\boldsymbol{\beta}$ . The distribution of  $\mathbf{y} | \mathbf{X}$  is

$$f(\mathbf{y} | \mathbf{X}, \boldsymbol{\beta}, \sigma) = (2\pi\sigma^2)^{-n/2} \exp \left( -\frac{(\mathbf{y} - \mathbf{X}\boldsymbol{\beta})'(\mathbf{y} - \mathbf{X}\boldsymbol{\beta})}{2\sigma^2} \right) \quad (2.3)$$

The log-likelihood function is

$$l(\mathbf{y}, \boldsymbol{\beta}, \sigma) = -\frac{n}{2} \log(2\pi) - \frac{n}{2} \log(\sigma^2) - \frac{1}{2\sigma^2} (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})'(\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) \quad (2.4)$$

The maximization of the log-likelihood function will lead to the same estimator as the OLS estimator, for  $\boldsymbol{\beta}$ . For  $\sigma^2$ , the MLE estimator is slightly different in limited sample, but asymptotically the same. This is strictly based on the assumption of normally distributed error term. Otherwise, these two estimators will be different.

## 2.3 Asymptotic Properties of MLE

### 1. Consistency

$$\text{plim}(\hat{\boldsymbol{\theta}}) = \boldsymbol{\theta}. \quad (2.5)$$

This is to say that MLE estimator can get arbitrarily close to the true parameter if the sample size gets to infinity.

### 2. Asymptotic Normality

$$\sqrt{n}(\hat{\boldsymbol{\theta}} - \boldsymbol{\theta}) \xrightarrow{d} N(\mathbf{0}, \mathbf{I}^{-1}(\boldsymbol{\theta})), \quad (2.6)$$

where  $\mathbf{I}(\boldsymbol{\theta})$  below is the *per-observation* information. Equivalently, in finite-sample terms,

$$\hat{\boldsymbol{\theta}} \approx N(\boldsymbol{\theta}, \mathbf{I}_n^{-1}(\boldsymbol{\theta})), \quad \mathbf{I}_n(\boldsymbol{\theta}) = n \mathbf{I}(\boldsymbol{\theta}), \quad (2.7)$$

with  $\mathbf{I}_n(\boldsymbol{\theta})$  the *full-sample* information. The  $\sqrt{n}$  scaling (equivalently, dividing the full-sample information by  $n$ ) is what makes the statement well-defined; writing  $\hat{\boldsymbol{\theta}} \sim N(\boldsymbol{\theta}, \mathbf{I}^{-1}(\boldsymbol{\theta}))$  without specifying which information is meant is ambiguous.

This states that the asymptotic distribution of  $\hat{\boldsymbol{\theta}}$  is normal with mean  $\boldsymbol{\theta}$  and variance given by the inverse of the information, where the per-observation information matrix  $\mathbf{I}(\boldsymbol{\theta})$  is defined using the single-observation log-likelihood contribution  $l_t$  (so that the full-sample log-likelihood is  $l = \sum_{t=1}^n l_t$ ):

$$\mathbf{I}(\boldsymbol{\theta}) = E\left[\left(\frac{\partial l_t}{\partial \boldsymbol{\theta}}\right)\left(\frac{\partial l_t}{\partial \boldsymbol{\theta}}\right)'\right] = -E\left[\left(\frac{\partial^2 l_t}{\partial \boldsymbol{\theta} \partial \boldsymbol{\theta}'}\right)\right] \quad (2.8)$$

The information matrix is the expected value of the outer product of the gradient of a single observation's log-likelihood contribution, or the negative value of the expected value of its Hessian (second derivative). Using the *full-sample*  $l$  here instead of  $l_t$  would give the full-sample information  $\mathbf{I}_n(\boldsymbol{\theta}) = n \mathbf{I}(\boldsymbol{\theta})$ , not  $\mathbf{I}(\boldsymbol{\theta})$  itself.

### 3. Asymptotic efficiency.

If  $\hat{\boldsymbol{\theta}}$  is the MLE estimator of  $\boldsymbol{\theta}$  and  $\mathbf{V}$  denotes its *asymptotic* variance matrix – the variance of the limiting distribution of  $\sqrt{n}(\hat{\boldsymbol{\theta}} - \boldsymbol{\theta})$ , not of  $\hat{\boldsymbol{\theta}}$  itself, the two differing by a factor of  $n$  exactly as with the information matrix above – then

$$\sqrt{n}(\hat{\boldsymbol{\theta}} - \boldsymbol{\theta}) \rightarrow^d N(\mathbf{0}, \mathbf{V}), \quad \mathbf{V} = \mathbf{I}^{-1}(\boldsymbol{\theta}). \quad (2.9)$$

If  $\tilde{\mathbf{V}}$  denotes the asymptotic variance matrix of any other *regular* consistent, asymptotically normal estimator, then  $\tilde{\mathbf{V}} - \mathbf{V}$  is a positive semidefinite matrix. That is, the MLE is the best estimator, in terms of asymptotic variance, in that class. The regularity qualifier is not pedantry: without it the statement is false, since superefficient estimators such as Hodges' can beat the bound at a single point of the parameter space, at the cost of much worse behaviour in a neighbourhood of it.

#### 4. Invariance.

If  $\hat{\theta}$  is the MLE of  $\theta$  and  $g(\theta)$  is a continuous function of  $\theta$ , then  $g(\hat{\theta})$  is the MLE of  $g(\theta)$ .

## 2.4 Example

Let's look at the example of regression of car's mpg on disp, hp and wt. This time we use MLE. (R has a `mle()` function which I cannot make it work on this example. So I am using `optim()`.)

```
ols.lf <- function(theta, y, X) {
  beta <- theta[-1]
  sigma2 <- theta[1]
  # Return a large penalty rather than NA: L-BFGS-B's line search can
  # transiently evaluate points that violate the box constraint, and
  # `fn` returning NA there causes a fatal "L-BFGS-B needs finite values
  # of 'fn'" error. A finite penalty keeps the optimizer on track instead.
  if (sigma2 <= 0) return(1e10)
  n <- nrow(X)
  e <- y - X%*%beta
  logl <- ((-n/2)*log(2*pi)) - ((n/2)*log(sigma2)) - ((t(e)%*%e)/(2*sigma2))
  return(-logl)
}
X <- cbind(1,mtcars[,c('disp','hp','wt')])
y <- mtcars$mpg

# Deliberately naive starting values (sigma2 = 1, slopes at +/-1): the point
# of the exercise is that the optimizer finds the OLS solution without being
# handed it. For harder likelihoods than this one, good starting values matter
# a great deal -- they reduce the risk of converging to a spurious local
# maximum, or of failing to converge at all.
optim(c(1,1,-1,-1,-1), method="L-BFGS-B", fn=ols.lf, lower=c(1e-6,-Inf,-Inf,-Inf,-Inf), upper=
```

```
$par
[1] 6.0935580702 37.1058998652 -0.0009352804 -0.0311594264 -3.8010038304

$value
[1] 74.32149

$counts
function gradient
```

```

126      126

$convergence
[1] 0

$message
[1] "CONVERGENCE: REL_REDUCTION_OF_F <= FACTR*EPSMCH"

```

```

# it should match the OLS estimator

lm1 <- lm(mpg~disp+hp+wt, data=mtcars)
summary(lm1)

```

```

Call:
lm(formula = mpg ~ disp + hp + wt, data = mtcars)

Residuals:
    Min       1Q   Median       3Q      Max
-3.891 -1.640 -0.172  1.061  5.861

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  37.105505   2.110815  17.579  < 2e-16 ***
disp         -0.000937   0.010350  -0.091  0.92851
hp           -0.031157   0.011436  -2.724  0.01097 *
wt           -3.800891   1.066191  -3.565  0.00133 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.639 on 28 degrees of freedom
Multiple R-squared:  0.8268,    Adjusted R-squared:  0.8083
F-statistic: 44.57 on 3 and 28 DF,  p-value: 8.65e-11

```

## 2.5 Hypotheses Testing

### 2.5.1 Linear Hypotheses Testing

Linear hypothesis can be framed as

$$\mathbf{H}_0 : \mathbf{R}\boldsymbol{\beta} - \mathbf{r} = \mathbf{0}, \quad (2.10)$$

where  $\mathbf{R}$  is a  $q \times k$  matrix of known constants, with  $q < k$ , and  $\mathbf{r}$  is a  $q$ -vector of known constants.

We assume that  $u_i$  are normally distributed.

$$\mathbf{u} \sim \mathbf{N}(\mathbf{0}, \sigma^2 \mathbf{I}) \quad (2.11)$$

Even if this assumption does not hold, we still have asymptotic normality for  $\hat{\beta}$ , so the tests below remain valid asymptotically.

Since linear combination of normal variables are also normally distributed,

$$\begin{aligned} \hat{\beta} &\sim \mathbf{N}(\beta, \sigma^2 (\mathbf{X}'\mathbf{X})^{-1}) \\ \mathbf{R}\hat{\beta} &\sim \mathbf{N}(\mathbf{R}\beta, \sigma^2 \mathbf{R}(\mathbf{X}'\mathbf{X})^{-1}\mathbf{R}') \\ \mathbf{R}(\hat{\beta} - \beta) &\sim \mathbf{N}(\mathbf{0}, \sigma^2 \mathbf{R}(\mathbf{X}'\mathbf{X})^{-1}\mathbf{R}') \end{aligned} \quad (2.12)$$

Under null hypothesis,

$$(\mathbf{R}\hat{\beta} - \mathbf{r}) \sim \mathbf{N}(\mathbf{0}, \sigma^2 \mathbf{R}(\mathbf{X}'\mathbf{X})^{-1}\mathbf{R}') \quad (2.13)$$

Therefore,

$$(\mathbf{R}\hat{\beta} - \mathbf{r})' [\sigma^2 \mathbf{R}(\mathbf{X}'\mathbf{X})^{-1}\mathbf{R}']^{-1} (\mathbf{R}\hat{\beta} - \mathbf{r}) \sim \chi^2(q) \quad (2.14)$$

Here we don't know  $\sigma^2$ . However, it can be shown that

$$\frac{\hat{\mathbf{u}}' \hat{\mathbf{u}}}{\sigma^2} \sim \chi^2(n - k) \quad (2.15)$$

Then we have

$$\frac{(\mathbf{R}\hat{\beta} - \mathbf{r})' [\mathbf{R}(\mathbf{X}'\mathbf{X})^{-1}\mathbf{R}']^{-1} (\mathbf{R}\hat{\beta} - \mathbf{r}) / \mathbf{q}}{\hat{\mathbf{u}}' \hat{\mathbf{u}} / (\mathbf{n} - \mathbf{k})} \sim F(q, n - k) \quad (2.16)$$

This is an example of the  $F$  test for any linear hypotheses testing.  $t$  test will be a special case of this.

### 2.5.2 The three classical tests: LR, Wald and LM

The  $F$  test above leans on normal errors. The general machinery for testing  $\mathbf{H}_0 : \mathbf{R}\theta - \mathbf{r} = \mathbf{0}$  after maximum likelihood needs only the asymptotic results already established, and it comes in three forms. Write  $\hat{\theta}_u$  for the unrestricted MLE,  $\hat{\theta}_r$  for the MLE subject to the restriction,  $s(\theta) = \partial\ell/\partial\theta$  for the score, and  $\mathbf{I}_n(\theta) = \mathbf{n}\mathbf{I}(\theta)$  for the full-sample information.

**Likelihood ratio.** Compare the two maximized log-likelihoods:

$$LR = 2[\ell(\hat{\theta}_u) - \ell(\hat{\theta}_r)] \xrightarrow{d} \chi^2(q). \quad (2.17)$$

**Wald.** Ask whether the *unrestricted* estimate is far from satisfying the restriction, measuring distance in units of its own estimated variance:

$$W = (\mathbf{R}\hat{\theta}_u - \mathbf{r})' [\mathbf{R}\mathbf{I}_n^{-1}(\hat{\theta}_u)\mathbf{R}']^{-1} (\mathbf{R}\hat{\theta}_u - \mathbf{r}) \xrightarrow{d} \chi^2(q). \quad (2.18)$$

**Lagrange multiplier (score).** Ask whether the score, evaluated at the *restricted* estimate, is far from zero – if the restriction is harmless, the restricted point should already be nearly a maximum:

$$LM = s(\hat{\theta}_r)' \mathbf{I}_n^{-1}(\hat{\theta}_r) s(\hat{\theta}_r) \xrightarrow{d} \chi^2(q). \quad (2.19)$$

The three are asymptotically equivalent under the null and have the same local power, so the choice among them is one of convenience: **Wald needs only the unrestricted fit, LM needs only the restricted fit, and LR needs both.** That is why LM is the natural choice for specification testing, where the restricted model is the easy one to estimate (White's heteroskedasticity test and the Breusch-Godfrey serial-correlation test are both LM tests run as auxiliary regressions), while Wald is what software reports by default, since the unrestricted fit is what was just computed.

In the classical normal linear model the three are monotone transformations of the same  $F$  statistic above, so nothing is lost by using  $F$  there. Two cautions otherwise. In finite samples the three can disagree, sometimes enough to straddle a conventional critical value, and there is a known ordering  $W \geq LR \geq LM$  in the linear model under normality. And the Wald statistic is **not invariant to nonlinear reparameterization**: testing  $\gamma = 1$  and testing  $\log \gamma = 0$  are the same hypothesis but give different Wald statistics, whereas  $LR$  is invariant because it compares likelihood *values*, which reparameterization does not change. When a hypothesis can be written several algebraically equivalent ways, prefer  $LR$ .

## 2.6 Test of Structural Change (Chow test)

Chow test is a test between two groups of observations. For example, we have data on consumption function of prewar period and postwar period. It is natural to think of a test on whether consumption function parameters differ between prewar and postwar periods. The null hypothesis is that there is no difference.

Let  $\mathbf{y}_i$ ,  $\mathbf{X}_i$  ( $i = 1, 2$ ) indicate the appropriate partitioning of the data. The unrestricted model may be written

$$\begin{bmatrix} \mathbf{y}_1 \\ \mathbf{y}_2 \end{bmatrix} = \begin{bmatrix} \mathbf{X}_1 & \mathbf{0} \\ \mathbf{0} & \mathbf{X}_2 \end{bmatrix} \begin{bmatrix} \beta_1 \\ \beta_2 \end{bmatrix} + u \quad u \sim N(0, \sigma^2 I) \quad (2.20)$$

The null hypothesis of no structural break is

$$\mathbf{H}_0 : \beta_1 = \beta_2. \quad (2.21)$$

Running OLS on two equations separately, we have

$$\begin{bmatrix} \hat{\beta}_1 \\ \hat{\beta}_2 \end{bmatrix} = \begin{bmatrix} \mathbf{X}_1' \mathbf{X}_1 & \mathbf{0} \\ \mathbf{0} & \mathbf{X}_2' \mathbf{X}_2 \end{bmatrix}^{-1} \begin{bmatrix} \mathbf{X}_1' \mathbf{y}_1 \\ \mathbf{X}_2' \mathbf{y}_2 \end{bmatrix} = \begin{bmatrix} (\mathbf{X}_1' \mathbf{X}_1)^{-1} \mathbf{X}_1' \mathbf{y}_1 \\ (\mathbf{X}_2' \mathbf{X}_2)^{-1} \mathbf{X}_2' \mathbf{y}_2 \end{bmatrix} \quad (2.22)$$

Restricted model under null hypothesis

$$\begin{bmatrix} \mathbf{y}_1 \\ \mathbf{y}_2 \end{bmatrix} = \begin{bmatrix} \mathbf{X}_1 \\ \mathbf{X}_2 \end{bmatrix} \boldsymbol{\beta} + \mathbf{u} \quad (2.23)$$

The test of the null hypothesis is given by

$$\mathbf{F} = \frac{(\hat{\mathbf{u}}'_* \hat{\mathbf{u}}_* - \hat{\mathbf{u}}' \hat{\mathbf{u}})/\mathbf{k}}{\hat{\mathbf{u}}' \hat{\mathbf{u}}/(\mathbf{n} - 2\mathbf{k})} \sim F(k, n - 2k) \quad (2.24)$$

where  $\hat{u}_*$  is the residual of the restricted model,  $\hat{u}$  is the stacked residual of two unrestricted models.

For example, for prewar period (suppose there is  $n_1$  observations):

$$y_1 = \beta_0 + \beta_1 x_1 + \beta_2 x_2 \quad (2.25)$$

For postwar period (suppose there is  $n_2$  observations):

$$y_2 = \gamma_0 + \gamma_1 x_1 + \gamma_2 x_2 \quad (2.26)$$

The null hypothesis of no structural break is

$$\mathbf{H}_0 : \beta_0 = \gamma_0, \beta_1 = \gamma_1, \beta_2 = \gamma_2. \quad (2.27)$$

Running OLS on two equations separately, we have two sums of squares of error ( $SSR_1$  and  $SSR_2$ ). The unrestricted error sum of squares is

$$SSR_U = SSR_1 + SSR_2 \quad (2.28)$$

Then run the regression on the stacked sample of  $n_1 + n_2$  observations and get  $SSR_R$ .

Then

$$\frac{(SSR_R - SSR_U)/k}{SSR_U/(n_1 + n_2 - 2k)} \sim F_{k, n_1 + n_2 - 2k}. \quad (2.29)$$

$k$  is number of restrictions; in this example, it is 3.

The other way to do the same test is easier: simply include interaction terms between group dummy and explanatory variables in the regression.

For our earlier example, we would estimate the model:

$$y = \alpha_0 + \alpha_1 x_1 + \alpha_2 x_2 + \eta_0 d + \eta_1 x_1 d + \eta_2 x_2 d \quad (2.30)$$

Here  $d$  is a dummy variable for postwar period. The Chow test is simply a test on the hypothesis that all the coefficients involving  $d$  are zero. A regression of the above model with  $n + m$  observations and a  $F$  test are easy to do.

$$\mathbf{H}_0 : \eta_0 = \eta_1 = \eta_2 = 0. \quad (2.31)$$

### 2.6.1 An example of Chow test in R

Here we use an example from the “strucchange” library.

```
## Example 7.4 from Greene (1993), "Econometric Analysis"
## Chow test on Longley data
data("longley")
library(strucchange)
## use structural change test from the library.
sctest(Employed ~ Year + GNP.deflator + GNP + Armed.Forces, data = longley,
       type = "Chow", point = 7)
```

Chow test

```
data:  Employed ~ Year + GNP.deflator + GNP + Armed.Forces
F = 3.9268, p-value = 0.06307
```

```
## which is equivalent to segmenting the regression via
fac <- factor(c(rep(1, 7), rep(2, 9)))
## here fac is the factor for segmenting it at point 7.
fm0 <- lm(Employed ~ Year + GNP.deflator + GNP + Armed.Forces, data = longley)
fm1 <- lm(Employed ~ fac/(Year + GNP.deflator + GNP + Armed.Forces), data = longley)
## Here anova is returning the F test (Chow test). Equivalent to sctest results.
anova(fm0, fm1)
```

Analysis of Variance Table

```
Model 1: Employed ~ Year + GNP.deflator + GNP + Armed.Forces
Model 2: Employed ~ fac/(Year + GNP.deflator + GNP + Armed.Forces)
  Res.Df    RSS Df Sum of Sq    F  Pr(>F)
1      11 4.8987
2       6 1.1466  5     3.7521 3.9268 0.06307 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
## estimates from Table 7.5 in Greene (1993)
summary(fm0)
```

Call:

```
lm(formula = Employed ~ Year + GNP.deflator + GNP + Armed.Forces,
    data = longley)
```

Residuals:

## 2.6 Test of Structural Change (Chow test)

| Min     | 1Q      | Median  | 3Q     | Max    |
|---------|---------|---------|--------|--------|
| -0.9058 | -0.3427 | -0.1076 | 0.2168 | 1.4377 |

Coefficients:

|              | Estimate   | Std. Error | t value | Pr(> t )   |
|--------------|------------|------------|---------|------------|
| (Intercept)  | 1.169e+03  | 8.359e+02  | 1.399   | 0.18949    |
| Year         | -5.765e-01 | 4.335e-01  | -1.330  | 0.21049    |
| GNP.deflator | -1.977e-02 | 1.389e-01  | -0.142  | 0.88940    |
| GNP          | 6.439e-02  | 1.995e-02  | 3.227   | 0.00805 ** |
| Armed.Forces | -1.015e-04 | 3.086e-03  | -0.033  | 0.97436    |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6673 on 11 degrees of freedom

Multiple R-squared: 0.9735, Adjusted R-squared: 0.9639

F-statistic: 101.1 on 4 and 11 DF, p-value: 1.346e-08

```
summary(fml)
```

Call:

```
lm(formula = Employed ~ fac/(Year + GNP.deflator + GNP + Armed.Forces),  
    data = longley)
```

Residuals:

| Min      | 1Q       | Median  | 3Q      | Max     |
|----------|----------|---------|---------|---------|
| -0.47717 | -0.18950 | 0.02089 | 0.14836 | 0.56493 |

Coefficients:

|                   | Estimate   | Std. Error | t value | Pr(> t )   |
|-------------------|------------|------------|---------|------------|
| (Intercept)       | 1.678e+03  | 9.390e+02  | 1.787   | 0.12413    |
| fac2              | 2.098e+03  | 1.786e+03  | 1.174   | 0.28473    |
| fac1:Year         | -8.352e-01 | 4.847e-01  | -1.723  | 0.13563    |
| fac2:Year         | -1.914e+00 | 7.913e-01  | -2.419  | 0.05194 .  |
| fac1:GNP.deflator | -1.633e-01 | 1.762e-01  | -0.927  | 0.38974    |
| fac2:GNP.deflator | -4.247e-02 | 2.238e-01  | -0.190  | 0.85576    |
| fac1:GNP          | 9.481e-02  | 3.815e-02  | 2.485   | 0.04747 *  |
| fac2:GNP          | 1.123e-01  | 2.269e-02  | 4.951   | 0.00258 ** |
| fac1:Armed.Forces | -2.467e-03 | 6.965e-03  | -0.354  | 0.73532    |
| fac2:Armed.Forces | -2.579e-02 | 1.259e-02  | -2.049  | 0.08635 .  |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4372 on 6 degrees of freedom

Multiple R-squared: 0.9938, Adjusted R-squared: 0.9845

F-statistic: 106.9 on 9 and 6 DF, p-value: 6.28e-06

## 2.7 Heteroskedasticity and Autocorrelation Consistent Standard Errors

### 2.7.1 Variance estimator under homoscedasticity

In OLS regression, the variance-covariance matrix of  $\hat{\beta}$  is

$$\text{var}(\hat{\beta}) = \mathbf{E}[(\hat{\beta} - \beta)(\hat{\beta} - \beta)'] = \mathbf{E}[(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{u}\mathbf{u}'\mathbf{X}(\mathbf{X}'\mathbf{X})^{-1}] = \sigma^2(\mathbf{X}'\mathbf{X})^{-1} \quad (2.32)$$

if

$$\text{var}(\mathbf{u}) = \sigma^2\mathbf{I}$$

{#eq-mle-33}.

That is, the error term has mean zero and constant variance (homoscedasticity). The pairwise correlation between error terms is always zero (no serial correlation).

The estimation of  $\sigma^2$ :

$$s^2 = \frac{\hat{\mathbf{u}}'\hat{\mathbf{u}}}{n - k} \quad (2.33)$$

is an unbiased estimator of  $\sigma^2$  (proof omitted).

The standard estimate of the variance-covariance matrix of the OLS parameter estimates under the assumption of IID errors is

$$\widehat{\text{Var}}(\hat{\beta}) = s^2(\mathbf{X}'\mathbf{X})^{-1} \quad (2.34)$$

### 2.7.2 White's estimator

We made strong assumption that the error terms of the regression model are IID when we estimate the variance-covariance matrix of OLS estimator. Under this assumption, the usual estimator of variance-covariance matrix of  $\hat{\beta}$  is consistent. Now let's relax this assumption to only independent but not identically distributed. Again, the linear regression models is

$$\mathbf{y} = \mathbf{X}\beta + \mathbf{u}, \quad \mathbf{E}(\mathbf{u}) = \mathbf{0}, \quad \mathbf{E}(\mathbf{u}\mathbf{u}') = \mathbf{\Omega}, \quad (2.35)$$

where  $\mathbf{\Omega}$  is the error variance-covariance matrix with diagonal elements being  $\sigma_t^2$  for  $t^{th}$  element, off-diagonal elements being zero. In other words, the error terms are heteroscedastic.

The variance-covariance matrix of the OLS estimator  $\hat{\beta}$  is equal to

$$\begin{aligned} \text{Var}(\hat{\beta}) &= \mathbf{E}[(\hat{\beta} - \beta)(\hat{\beta} - \beta)'] \\ &= [(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'(\mathbf{E}(\mathbf{u}\mathbf{u}'))\mathbf{X}(\mathbf{X}'\mathbf{X})^{-1}] \\ &= (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{\Omega}\mathbf{X}(\mathbf{X}'\mathbf{X})^{-1} \end{aligned} \quad (2.36)$$

If we know  $\sigma_t^2$ , then we would be able to estimate this “sandwich covariance matrix”. But we don't.

$$\text{Var}(\hat{\beta}) = \frac{1}{\mathbf{n}} \left[ \frac{1}{\mathbf{n}} (\mathbf{X}'\mathbf{X}) \right]^{-1} \left[ \frac{1}{\mathbf{n}} \mathbf{X}'\mathbf{\Omega}\mathbf{X} \right] \left[ \frac{1}{\mathbf{n}} (\mathbf{X}'\mathbf{X}) \right]^{-1} \quad (2.37)$$

Let  $y_t$  denote the  $t$ th observation on the dependent variable, and  $x'_t = [1 x_{2t} \cdots x_{kt}]$  denote the  $t$ th row of the  $\mathbf{X}$  matrix. Then

$$\mathbf{X}'\Omega\mathbf{X} = \sum_{t=1}^n \sigma_t^2 \mathbf{x}_t \mathbf{x}_t' \quad (2.38)$$

The White estimator replaces the unknown  $\sigma_t^2$  by  $\hat{u}_t^2$ , the estimated OLS residuals. This provides a consistent estimator of the variance matrix for the OLS coefficient vector and is particularly useful since it does not require any specific assumptions about the form of the heteroscedasticity.

Therefore,

$$\begin{aligned} \text{Var}(\hat{\beta}) &= \frac{1}{n} \left[ \frac{1}{n} (\mathbf{X}'\mathbf{X}) \right]^{-1} \left[ \frac{1}{n} \mathbf{X}'\hat{\Omega}\mathbf{X} \right] \left[ \frac{1}{n} (\mathbf{X}'\mathbf{X}) \right]^{-1} \\ \hat{\Omega} &= \text{diag}(\hat{u}_1^2, \hat{u}_2^2, \dots, \hat{u}_n^2) \end{aligned} \quad (2.39)$$

### 2.7.3 Newey-West estimator

White's estimator deals with the situation that we have heteroskedasticity (a diagonal  $\Omega$ ) of unknown form. When we have serial correlation of unknown form (a non-diagonal  $\Omega$ ), we can estimate the variance-covariance matrix by a heteroskedasticity and autocorrelation consistent, or HAC, estimator. Newey-West estimator is the most popular HAC estimator.

Given a time series data set, suppose we are interested in estimating the mean vector (suppose we have more than one variable) and its variance. We know that given IID data, we can apply central limit theorem: sample mean is a consistent estimator of the population mean and its variance can be calculated since asymptotically the sample mean conforms to a normal distribution and the variance can be estimated, relatively easily. However, in the case of time series data, autocorrelation usually exists. We may be concerned the CLT may not work in this case.

Fortunately, as proved in Hamilton (1994), if  $\mathbf{y}_t$  is a covariance-stationary (meaning that the covariance is not a function of time) vector process, then the sample mean satisfies:

$$\bar{\mathbf{y}}_t \rightarrow \boldsymbol{\mu}, \quad (2.40)$$

$$\mathbf{S} = \lim_{T \rightarrow \infty} \mathbf{T} \cdot \mathbf{E}[(\bar{\mathbf{y}}_T - \boldsymbol{\mu})(\bar{\mathbf{y}}_T - \boldsymbol{\mu})'] = \sum_{v=-\infty}^{\infty} \boldsymbol{\Gamma}_v. \quad (2.41)$$

where  $\boldsymbol{\Gamma}_v$  is the variance-covariance matrix for  $\mathbf{y}_t$  and  $\mathbf{y}_{t-v}$ .

The first one says for a covariance-stationary vector process, the law of large numbers still holds. The second one is used to calculate the standard error.

If the data were generated by a vector MA( $q$ ) process, then

$$\mathbf{S} = \sum_{v=-q}^q \boldsymbol{\Gamma}_v. \quad (2.42)$$

A natural estimate is

$$\hat{\mathbf{S}} = \hat{\boldsymbol{\Gamma}}_0 + \sum_{v=1}^q (\hat{\boldsymbol{\Gamma}}_v + \hat{\boldsymbol{\Gamma}}_v'), \quad (2.43)$$

## 2 The Method of Maximum Likelihood

where

$$\hat{\Gamma}_{\mathbf{v}} = (\mathbf{1}/\mathbf{T}) \sum_{\mathbf{t}=\mathbf{v}+1}^{\mathbf{T}} (\mathbf{y}_{\mathbf{t}} - \bar{\mathbf{y}})(\mathbf{y}_{\mathbf{t}-\mathbf{v}} - \bar{\mathbf{y}})'. \quad (2.44)$$

This gives a consistent estimate of  $\mathbf{S}$ ; however, it sometimes is not positive semidefinite.

Newey-West (1987) suggested putting in a weight:

$$\hat{\mathbf{S}} = \hat{\Gamma}_{\mathbf{0}} + \sum_{v=1}^q \left(1 - \frac{v}{q+1}\right) (\hat{\Gamma}_{\mathbf{v}} + \hat{\Gamma}'_{\mathbf{v}}), \quad (2.45)$$

where  $q$  is from the MA( $q$ ) process.

Consider a linear regression model:

$$y_t = \mathbf{x}'_t \boldsymbol{\beta} + u_t \quad (2.46)$$

Suppose we have the OLS estimator  $\mathbf{b}_{\mathbf{T}}$ , then

$$\sqrt{T}(\mathbf{b}_{\mathbf{T}} - \boldsymbol{\beta}) = [(1/T) \sum_{t=1}^T \mathbf{x}_t \mathbf{x}'_t]^{-1} [(1/\sqrt{T}) \sum_{t=1}^T \mathbf{x}_t \mathbf{u}_t] \quad (2.47)$$

The first term converges in probability to some constant. The second term is  $\sqrt{T}$  times the sample mean of the vector  $\mathbf{x}_t \mathbf{u}_t$ , which is exactly the form to which a central limit theorem applies.

Under general conditions,

$$\sqrt{T}(\mathbf{b}_{\mathbf{T}} - \boldsymbol{\beta}) \rightarrow^L N(0, Q^{-1} S Q^{-1}) \quad (2.48)$$

where  $Q = \text{plim } (1/T) \sum_{t=1}^T \mathbf{x}_t \mathbf{x}'_t = E[\mathbf{x}_t \mathbf{x}'_t]$ , and  $S$  can be estimated by

$$\hat{\mathbf{S}}_{\mathbf{T}} = \hat{\Gamma}_{\mathbf{0T}} + \sum_{v=1}^q \left(1 - \frac{v}{q+1}\right) (\hat{\Gamma}_{\mathbf{v},\mathbf{T}} + \hat{\Gamma}'_{\mathbf{v},\mathbf{T}}), \quad (2.49)$$

where

$$\hat{\Gamma}_{v,T} = (1/T) \sum_{t=v+1}^T (x_t \hat{u}_{t,T} \hat{u}_{t-v,T} x'_{t-v}), \quad (2.50)$$

where  $\hat{u}_{t,T}$  is the OLS residual for data  $t$  in a sample of size  $T$ .

Overall, the variance of  $\mathbf{b}_{\mathbf{T}}$  is approximated by

$$\hat{\Sigma}_{NW} = \left[ \sum_{t=1}^T x_t x'_t \right]^{-1} \left[ \sum_{t=1}^T \hat{u}_t^2 x_t x'_t + \sum_{v=1}^q \left(1 - \frac{v}{q+1}\right) \sum_{t=v+1}^T (x_t \hat{u}_{t,T} \hat{u}_{t-v,T} x'_{t-v} + x_{t-v} \hat{u}_{t-v,T} \hat{u}_{t,T} x'_t) \right] \left[ \sum_{t=1}^T x_t x'_t \right]^{-1} \quad (2.51)$$

This estimation obviously depends on the selection of  $q$ , the lag length beyond which we are willing to assume that the autocorrelation of  $x_t u_t$  and  $x_{t-v} u_{t-v}$  is essentially zero. The rule of thumb for the selection of  $q$  is  $0.75 \cdot T^{\frac{1}{3}}$ . Newey-West (1994) has suggested a way to automatically select the bandwidth  $q$ . Here we omitted the discussion. Both Stata and R now also implement Newey-West (1994) estimator, with no need to specify  $q$ .

## 3 GLS

### 3.1 Introduction

Since heteroscedasticity and serial correlation affect both linear and nonlinear regression models in the same way, there is no harm in limiting our attention to the simpler linear case.

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \mathbf{u}, \quad \mathbf{E}(\mathbf{u}) = \mathbf{0}, \quad \mathbf{E}(\mathbf{u}\mathbf{u}') = \boldsymbol{\Omega}, \quad (3.1)$$

where  $\boldsymbol{\Omega}$ , the variance-covariance matrix of the error term, is a positive definite  $n \times n$  matrix. If  $\boldsymbol{\Omega}$  is equal to  $\sigma^2\mathbf{I}$ , then it is just the linear regression model without heteroscedasticity and serial correlation. If  $\boldsymbol{\Omega}$  is diagonal with nonconstant diagonal elements, then the error terms are still uncorrelated, but there is heteroscedasticity. If  $\boldsymbol{\Omega}$  is not diagonal, then  $u_i$  and  $u_j$  are correlated. In next section, we obtain an efficient estimator for the vector  $\boldsymbol{\beta}$  by transforming the regression so that it satisfies the conditions of Gauss-Markov Theorem. This efficient estimator is called the Generalized Least Squares, or GLS, estimator.

### 3.2 The GLS Estimator

Since  $\boldsymbol{\Omega}$ , the variance-covariance matrix of the error term, is a positive definite  $n \times n$  matrix, there always exist full-rank  $n \times n$  matrices (usually triangular)  $\boldsymbol{\Psi}$  such that

$$\boldsymbol{\Omega}^{-1} = \boldsymbol{\Psi}\boldsymbol{\Psi}'. \quad (3.2)$$

Premultiplying (the equation above) by  $\boldsymbol{\Psi}'$  gives

$$\boldsymbol{\Psi}'\mathbf{y} = \boldsymbol{\Psi}'\mathbf{X}\boldsymbol{\beta} + \boldsymbol{\Psi}'\mathbf{u}. \quad (3.3)$$

The OLS estimator from regression (the equation above) is

$$\hat{\boldsymbol{\beta}}_{\text{GLS}} = (\mathbf{X}'\boldsymbol{\Psi}\boldsymbol{\Psi}'\mathbf{X})^{-1}\mathbf{X}'\boldsymbol{\Psi}\boldsymbol{\Psi}'\mathbf{y} = (\mathbf{X}'\boldsymbol{\Omega}^{-1}\mathbf{X})^{-1}\mathbf{X}'\boldsymbol{\Omega}^{-1}\mathbf{y} \quad (3.4)$$

The transformed error term has an identity variance-covariance matrix:

$$\mathbf{E}(\boldsymbol{\Psi}'\mathbf{u}\mathbf{u}'\boldsymbol{\Psi}) = \boldsymbol{\Psi}'\mathbf{E}(\mathbf{u}\mathbf{u}')\boldsymbol{\Psi} = \boldsymbol{\Psi}'\boldsymbol{\Omega}\boldsymbol{\Psi} = \mathbf{I} \quad (3.5)$$

Since the transformed model satisfies OLS assumptions, we have

$$\begin{aligned} \text{Var}(\hat{\boldsymbol{\beta}}_{\text{GLS}}) &= (\mathbf{X}'\boldsymbol{\Psi}\boldsymbol{\Psi}'\mathbf{X})^{-1} \\ &= (\mathbf{X}'\boldsymbol{\Omega}^{-1}\mathbf{X})^{-1} \end{aligned} \quad (3.6)$$

### 3.3 Weighted Least Squares

It is easy to obtain GLS estimates when the error terms are heteroscedastic but uncorrelated, which means  $\Omega$  is diagonal. Let  $\omega_t^2$  denote the  $t$ th diagonal element of  $\Omega$ . Then  $\Psi$  can be chosen as the diagonal matrix with  $t$ th diagonal element  $\omega_t^{-1}$ . For a typical observation, regression can be written as

$$\omega_t^{-1}y_t = \omega_t^{-1}\mathbf{x}_t'\beta + \omega_t^{-1}u_t. \quad (3.7)$$

This is called weighted least squares, or WLS. It is worth keeping two notions of “weight” distinct. The transformed variables  $y_t$  and  $\mathbf{x}_t'$  are *scaled* by  $\omega_t^{-1}$ ; equivalently, the WLS objective weights each squared residual by the inverse variance  $\omega_t^{-2}$  (the square of the scaling factor). Either way, observations whose error variance is large are given low weight, and observations for which it is small are given high weight.

#### 3.3.1 Feasible Generalized Least Squares

In many cases it is reasonable to suppose that  $\Omega$  depends in a known way on a vector of unknown parameters  $\gamma$ . If so, it may be possible to estimate consistently, so as to obtain  $\Omega(\hat{\gamma})$ . This type of procedure is called feasible generalized least squares, or feasible GLS. But we’ll have to specify the error term as some function of some known variables.

A typical FGLS procedure has three steps. First, run OLS and obtain the residuals  $\hat{u}_t$ . Second, use those residuals to estimate the parameters  $\gamma$  of the assumed error structure – for example:

- **Heteroscedasticity of known form**, e.g.  $\omega_t^2 = \exp(\mathbf{z}_t'\gamma)$ : regress  $\log(\hat{u}_t^2)$  on  $\mathbf{z}_t$  to estimate  $\gamma$ , then set  $\hat{\omega}_t^2 = \exp(\mathbf{z}_t'\hat{\gamma})$  and apply WLS as above. (The intercept of that auxiliary regression is inconsistent, because  $E[\log(u_t^2/\omega_t^2)] \neq 0$  – it is about  $-1.27$  for normal errors. This does not matter for WLS: a constant added to every  $\log \hat{\omega}_t^2$  scales all weights by a common factor, which leaves  $\hat{\beta}$  unchanged. It does matter if the fitted  $\hat{\omega}_t^2$  are wanted as variance estimates in their own right.)
- **AR(1) serial correlation**,  $u_t = \rho u_{t-1} + \epsilon_t$ : regress  $\hat{u}_t$  on  $\hat{u}_{t-1}$  to get  $\hat{\rho}$ , then quasi-difference each variable as  $\tilde{y}_t = y_t - \hat{\rho}y_{t-1}$  for  $t = 2, \dots, n$  (and similarly for  $X$ ), so that the transformed error is approximately i.i.d.

The first observation is where **Cochrane-Orcutt** and **Prais-Winsten** part company, and the  $\Psi$  framework above says which is right. For AR(1) errors  $\Omega = \frac{\sigma_\epsilon^2}{1-\rho^2} [\rho^{|i-j|}]$ , and its Cholesky factor makes the first row of  $\Psi'$  equal to  $\sqrt{1-\rho^2}$  times the first unit vector. Full GLS therefore keeps observation 1, rescaled as  $\tilde{y}_1 = \sqrt{1-\hat{\rho}^2} y_1$ ; that is Prais-Winsten. Cochrane-Orcutt simply drops observation 1. Both are consistent and asymptotically equivalent, but Cochrane-Orcutt throws away a data point, and the loss is not negligible in short time series or when  $|\rho|$  is close to 1 – precisely the case where the correction matters most. Prefer Prais-Winsten.

Third, run OLS on the transformed (weighted or quasi-differenced) data to get the FGLS estimator  $\hat{\beta}_{FGLS}$ . FGLS is consistent and asymptotically efficient as long as  $\hat{\gamma}$  is a consistent estimator of  $\gamma$ , even though it is generally biased in finite samples (since  $\hat{\Omega}$  depends on the first-stage OLS residuals).

## 4 Endogeneity

When we discuss the linear model we assume that the regressors are exogenous, meaning that they are independent of or uncorrelated with the error term. However, there could be reasons to believe that some regressors are correlated with the error term. In that case we call those regressors endogenous.

Under the classical assumptions OLS estimators are unbiased and consistent. One key assumption is that the regressors have to be uncorrelated with the error term. If this condition does not hold, OLS estimators are biased and inconsistent.

When one independent variable does not satisfy this condition, we say this variable is endogenous. When one variable is endogenous, the estimates of other coefficients will also be biased (or inconsistent).

The most popular cure for endogeneity is to use instrumental variables.

### 4.1 Instrumental Variables (IV)

The basic idea of IV is to use an exogenous variable (or exogenous variables) which is correlated with the endogenous independent variable as an “instrument” for the endogenous variable.

Suppose the linear regression model

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \mathbf{u}, \quad \mathbf{E}(\mathbf{u}\mathbf{u}') = \sigma^2\mathbf{I}, \quad (4.1)$$

at least one of the explanatory variables in the  $n \times k$  matrix  $\mathbf{X}$  is assumed not to be predetermined with respect to the error terms, or say, endogenous.

Suppose we have a set of variables  $\mathbf{Z}$ , an  $n \times l$  matrix of instruments, which satisfies the moment condition

$$\mathbf{E}[\mathbf{Z}'(\mathbf{y} - \mathbf{X}\boldsymbol{\beta})] = \mathbf{0}. \quad (4.2)$$

That is,  $\mathbf{Z}$  is uncorrelated with the error term.

Here is what we do for two-stage least squares (2sls):

Stage 1: Regress each of the variables in the  $\mathbf{X}$  matrix on  $\mathbf{Z}$  to obtain a matrix of fitted values  $\hat{\mathbf{X}}$ ,

$$\hat{\mathbf{X}} = \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\mathbf{X} = \mathbf{P}_\mathbf{Z}\mathbf{X} \quad (4.3)$$

This is essentially to get the part of  $\mathbf{X}$  that is correlated with  $\mathbf{Z}$ . Or to say, to project  $\mathbf{X}$  on to  $\mathbf{Z}$ .

Stage 2: Regress  $\mathbf{y}$  on  $\hat{\mathbf{X}}$  to obtain the estimated  $\boldsymbol{\beta}$

#### 4 Endogeneity

$$\hat{\beta}_{2sls} = (\hat{\mathbf{X}}' \hat{\mathbf{X}})^{-1} (\hat{\mathbf{X}}' \mathbf{y}) = (\mathbf{X}' \mathbf{P}_Z \mathbf{X})^{-1} (\mathbf{X}' \mathbf{P}_Z \mathbf{y}) = \hat{\beta}_{IV} \quad (4.4)$$

Standard errors:

$$Var[\hat{\beta}_{2sls}] = \hat{\sigma}^2 (\hat{\mathbf{X}}' \hat{\mathbf{X}})^{-1} = \hat{\sigma}^2 (\mathbf{X}' \mathbf{P}_Z \mathbf{X})^{-1} \quad (4.5)$$

where  $\hat{\sigma}^2 = \hat{\mathbf{u}}' \hat{\mathbf{u}} / N$ . This is the homoskedastic large-sample variance; finite-sample formulas typically apply a degrees-of-freedom correction ( $N - k$  in place of  $N$ ), and when errors are heteroskedastic or clustered the robust/clustered IV covariance matrix replaces  $\hat{\sigma}^2 (\mathbf{X}' \mathbf{P}_Z \mathbf{X})^{-1}$  with a sandwich form.

Note that  $\hat{\mathbf{u}} = \mathbf{y} - \mathbf{X} \hat{\beta}_{2sls}$ .

If we do it manually by two steps, the second step will report a wrong standard error. That is because the second stage regression will report standard errors based on  $\hat{\mathbf{u}} = \mathbf{y} - \hat{\mathbf{X}} \hat{\beta}_{2sls}$ . Therefore, it's always recommended to ask the statistical program to do a 2sls for you, since presumably that will give you correct standard errors.

Geometry Illustration:

Suppose the simplest case:  $y = \beta_0 + \beta_1 x + u$ . Take all variables in deviations from their means, so the intercept drops out and  $z' u = 0$ ; this keeps the algebra below about the slope alone. Split  $x$  into the part that lies along the error and the part that does not,

$$x = x^\perp + x^\parallel, \quad (4.6)$$

where  $x^\parallel$  is the orthogonal projection of  $x$  onto  $u$  and  $x^\perp$  is what is left over. Two facts hold *by construction* and are what make the picture work:  $x^\perp$  is perpendicular to  $u$ , and  $x^\perp$  is perpendicular to  $x^\parallel$ . Endogeneity is simply the statement that  $x^\parallel \neq 0$ . (This is a decomposition of  $x$ , not an extra assumption about it – every  $x$  admits it.)

Now let  $z$  be a vector perpendicular to  $u$  (exogeneity:  $z$  is uncorrelated with the error) with a non-zero projection onto  $x$  (relevance:  $z$  actually predicts the regressor). Project both  $y$  and  $x$  onto  $z$ , and write  $\hat{y}_z$  and  $\hat{x}_z$  for those two projections. The instrumental variables estimator is the ratio of their lengths,

$$\hat{\beta}_{IV} = \frac{\hat{y}_z}{\hat{x}_z} = \frac{z' y}{z' x}, \quad (4.7)$$

since  $z' y = \beta_1 z' x + z' u = \beta_1 z' x$ . Because the two right triangles involved are similar, this is the same number you would get from the *infeasible* regression of  $y$  on the exogenous part alone: that regression returns

$$\frac{x^{\perp'} y}{x^{\perp'} x^\perp} = \beta_1 \frac{x^{\perp'} x}{x^{\perp'} x^\perp} + \frac{x^{\perp'} u}{x^{\perp'} x^\perp} = \beta_1, \quad (4.8)$$

where the second term vanishes because  $x^\perp \perp u$  and the first collapses to  $\beta_1$  because  $x^\perp \perp x^\parallel$ , so  $x^{\perp'} x = x^{\perp'} x^\perp$ . That infeasible regression is the one we would run if we could observe the decomposition. We cannot – which is exactly why we need the instrument.

```
## DGP: data$y <- data$x + data$z + data$u
library(MASS)
set.seed(66)
nobs=10000
nDim = 3
sdxx = 1
sdww=1
sdzz=1

## here we have three variables x,z,w.
## z is the omitted variable, x and z are correlated, w is the instrument, which is correlated
crxz=.6
crzw=0
crxw=.5

covarMat = matrix( c(sdxx^2, crxz, crxw, crxz, sdzz^2, crzw, crxw, crzw, sdww^2 ), nrow=nDim,
covarMat
```

```
      [,1] [,2] [,3]
[1,]  1.0  0.6  0.5
[2,]  0.6  1.0  0.0
[3,]  0.5  0.0  1.0
```

```
data = data.frame(mvrnorm(n=nobs, mu=rep(0,nDim), Sigma=covarMat ))
names(data) <- c('x','z','w')
data$u <- rnorm(nobs,0,1)
# dgp
data$y <- data$x + data$z + data$u

lm <- lm(y~x, data=data)
lm.full <- lm(y~ x + z, data=data)
tsls.model <- ivreg(y ~ x | w, data=data)
# lm is biased
summary(lm)
```

Call:

```
lm(formula = y ~ x, data = data)
```

Residuals:

| Min     | 1Q      | Median | 3Q     | Max    |
|---------|---------|--------|--------|--------|
| -5.3506 | -0.8612 | 0.0123 | 0.8504 | 4.9096 |

Coefficients:

|             | Estimate  | Std. Error | t value | Pr(> t ) |
|-------------|-----------|------------|---------|----------|
| (Intercept) | -0.004284 | 0.012843   | -0.334  | 0.739    |

#### 4 Endogeneity

```
x          1.596114    0.013079 122.035    <2e-16 ***
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 1.284 on 9998 degrees of freedom

Multiple R-squared: 0.5983, Adjusted R-squared: 0.5983

F-statistic: 1.489e+04 on 1 and 9998 DF, p-value: < 2.2e-16

```
# lm.full is good
summary(lm.full)
```

Call:

```
lm(formula = y ~ x + z, data = data)
```

Residuals:

| Min     | 1Q      | Median  | 3Q     | Max    |
|---------|---------|---------|--------|--------|
| -3.7267 | -0.6633 | -0.0179 | 0.6742 | 3.6018 |

Coefficients:

|             | Estimate | Std. Error | t value | Pr(> t )   |
|-------------|----------|------------|---------|------------|
| (Intercept) | 0.001857 | 0.009973   | 0.186   | 0.852      |
| x           | 0.987181 | 0.012629   | 78.166  | <2e-16 *** |
| z           | 1.017559 | 0.012542   | 81.129  | <2e-16 *** |

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 0.997 on 9997 degrees of freedom

Multiple R-squared: 0.7578, Adjusted R-squared: 0.7577

F-statistic: 1.564e+04 on 2 and 9997 DF, p-value: < 2.2e-16

```
# tsls is good.
summary(tsls.model)
```

Call:

```
ivreg(formula = y ~ x | w, data = data)
```

Residuals:

| Min       | 1Q        | Median   | 3Q       | Max      |
|-----------|-----------|----------|----------|----------|
| -6.144530 | -0.962464 | 0.003414 | 0.944728 | 5.340560 |

Coefficients:

|             | Estimate | Std. Error | t value | Pr(> t )   |
|-------------|----------|------------|---------|------------|
| (Intercept) | 0.01153  | 0.01433    | 0.804   | 0.421      |
| x           | 0.95129  | 0.03007    | 31.637  | <2e-16 *** |

Diagnostic tests:

|                  | df1 | df2  | statistic | p-value    |
|------------------|-----|------|-----------|------------|
| Weak instruments | 1   | 9998 | 3074.7    | <2e-16 *** |
| Wu-Hausman       | 1   | 9997 | 807.8     | <2e-16 *** |
| Sargan           | 0   | NA   | NA        | NA         |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.431 on 9998 degrees of freedom

Multiple R-Squared: 0.5007, Adjusted R-squared: 0.5006

Wald test: 1001 on 1 and 9998 DF, p-value: < 2.2e-16

## 4.2 Control Function Approach

A second way to do an IV regression is also two-step approach: Regress  $\mathbf{X}$  (endogenous) on  $\mathbf{Z}$ , get the residual:  $\hat{\mathbf{v}} = \mathbf{X} - \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\mathbf{X}$ , then regress  $y$  on  $\mathbf{X}$  and  $\hat{\mathbf{v}}$  to get  $\hat{\beta}_{IV}$ .

The difference between this approach and the 2sls approach is that in 2sls we regress  $y$  on  $\hat{\mathbf{X}}$ ; in the control function approach, we regress  $y$  on  $\mathbf{X}$  and  $\hat{\mathbf{v}}$ . They should both give you the same coefficient estimates. There are advantages of using the control function approach.

An example to show how to do 2sls manually, or using control function approach.

```
## DGP: data$y <- data$x + data$z + data$u

lm1 <- lm(y~x, data=data)
lm2 <- lm(x~w, data=data)
tsls.manual <- lm(data$y ~ lm2$fitted.values)
summary(tsls.manual)
```

Call:

```
lm(formula = data$y ~ lm2$fitted.values)
```

Residuals:

|  | Min     | 1Q      | Median  | 3Q     | Max    |
|--|---------|---------|---------|--------|--------|
|  | -8.6485 | -1.3401 | -0.0203 | 1.3543 | 7.6974 |

Coefficients:

|                    | Estimate | Std. Error | t value | Pr(> t )   |
|--------------------|----------|------------|---------|------------|
| (Intercept)        | 0.01153  | 0.01977    | 0.583   | 0.56       |
| lm2\$fitted.values | 0.95129  | 0.04147    | 22.937  | <2e-16 *** |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.974 on 9998 degrees of freedom

#### 4 Endogeneity

Multiple R-squared: 0.04999, Adjusted R-squared: 0.04989  
F-statistic: 526.1 on 1 and 9998 DF, p-value: < 2.2e-16

```
# control function approach
tsls.control <- lm(data$y ~ data$x + lm2$residuals)
summary(tsls.control)
```

Call:

```
lm(formula = data$y ~ data$x + lm2$residuals)
```

Residuals:

| Min     | 1Q      | Median | 3Q     | Max    |
|---------|---------|--------|--------|--------|
| -5.1997 | -0.8329 | 0.0131 | 0.8370 | 4.7457 |

Coefficients:

|                | Estimate | Std. Error | t value | Pr(> t )   |
|----------------|----------|------------|---------|------------|
| (Intercept)    | 0.01153  | 0.01237    | 0.932   | 0.351      |
| data\$x        | 0.95129  | 0.02594    | 36.670  | <2e-16 *** |
| lm2\$residuals | 0.84313  | 0.02966    | 28.423  | <2e-16 *** |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.235 on 9997 degrees of freedom

Multiple R-squared: 0.6284, Adjusted R-squared: 0.6283

F-statistic: 8451 on 2 and 9997 DF, p-value: < 2.2e-16

This approach can be used to do a simple endogeneity test. First, you can do a simple test of endogeneity of **X**. For example:

```
reg x_endog x* z*
predict v_x, resid
reg y x_endog x* v_x
test v_x
```

Obviously, this test is based on **Z** being exogenous.

Residual inclusion extends to *some* non-linear models – logit, probit, Poisson – under additional control-function assumptions, but it is not automatically valid for “any glm”. The construction is model-specific: the appropriate first-stage residual (a plain linear residual or a generalized residual), the identification conditions, and the standard errors all differ for continuous, binary, and nonlinear endogenous regressors, and inserting a linear first-stage residual into an arbitrary GLM is generally not correct. The generated-regressor step also inflates uncertainty, so standard errors need to account for the estimated first stage (e.g. by bootstrap). For example, the “eteffects” procedure in Stata (version 14) uses a control-function approach to get endogenous treatment effects for several outcome types.

## 4.3 Durbin-Wu-Hausman Test

### 4.3.1 Idea

In econometric modeling, there are often questions on endogeneity. Do we know how to test whether an independent variable is endogenous statistically? The answer is: sort of, but not really. We cannot do endogeneity test without a valid instrument. Therefore, we have to have strong argument for a valid instrument first before we can do endogeneity test.

With endogenous variables on the right-hand side of the equation, we need to use instrumental variable (IV) regression for consistent estimation. However, with IV regression, we lose efficiency: the asymptotic variance of the IV estimator is larger, and can be much larger than the OLS estimator. Therefore, we gain consistency, but lose efficiency, by using IV estimator when there is an endogeneity problem.

Now we have a familiar scenario (if you are familiar with Hausman test for fixed effect and random effect estimator for panel data): Suppose we have the null hypothesis as the regressor being exogenous. We have an efficient estimator under null hypothesis yet inconsistent under alternative hypothesis (OLS estimator). We also have a consistent estimator under both null and alternative (IV estimator).

Similar to panel data setting, we have the Hausman test statistic as:

$$H = (\hat{\beta}_c - \hat{\beta}_e)' D^- (\hat{\beta}_c - \hat{\beta}_e) \quad (4.9)$$

where  $D = \text{Var}[\hat{\beta}_c] - \text{Var}[\hat{\beta}_e]$ ,  $-$  is the generalized inverse,  $\hat{\beta}_c$  is the consistent estimator (in this case the IV estimator) and  $\hat{\beta}_e$  is the efficient estimator (in this case OLS estimator).

$H$  conforms to  $\chi_k^2$  asymptotically, where  $k$  is the number of endogenous variables.

This test is to compare the IV estimator and the OLS estimator: if it's close, then OLS estimator is fine (fail to reject null that OLS is consistent, or say the variable is exogenous). If it's large, then IV estimator is needed, although we lose some efficiency. This test is based on the assumption that the instruments are exogenous. If that is in question, then it's pointless to do the test, since the IV estimator cannot guarantee consistency either.

### 4.3.2 Implementation in Stata

In Stata, there are different ways to do it:

Do a regular Hausman test:

```
ivreg y x1 (x2=x3 x4)
estimates store iv
reg y x1 x2
hausman iv ., constant sigmamore
```

Or, simply use “ivendog” in Stata.

## 4.4 Identification

Identification in a regression equation means that all parameters can be uniquely estimated. The relevant counting rule (the *order condition*) is most cleanly stated in terms of *excluded* instruments. Let  $k_e$  be the number of endogenous regressors and  $l_e$  the number of excluded instruments (instruments not already included as exogenous regressors). A necessary condition is  $l_e \geq k_e$ : there must be at least as many excluded instruments as endogenous regressors. (Equivalently, counting the included exogenous regressors as their own instruments, the total number of instruments must be at least the total number of regressors.) If  $l_e = k_e$  we have exact-identification. If  $l_e > k_e$ , we have over-identification.

Over-identification generates more efficient estimates, given the assumption of instruments being exogenous. The other advantage of over-identification is that over-identification tests can be done to test the adequacy of instruments.

Under the null hypothesis that all the instruments are uncorrelated with the error term, an LM statistic  $N \times R^2$  conforms to  $\chi^2(r)$  distribution,  $r = l_e - k_e$ , the number of excess instruments, or say, the number of overidentifying restrictions. If we reject the null, then we should be concerned about the exogeneity of the whole set of the instruments. This test is called Sargan's test in IV context, and (Hansen's) J test in GMM context.

What the J test or Sargan's test does is to test the whole set of instruments being exogenous or not. There is another test for testing exogeneity for a subset of instruments. It's call a C test or a difference-in-Sargan test. The idea is to calculate the difference between two Sargan's statistics (or Hansen's J in GMM setting); one is with the whole set of instruments, the other one without the suspected instruments. The null is that the suspect instruments are exogenous; or orthogonal to the error term. Obviously to conduct the C test, we'll have to have at least one extra instrument more than the number of endogenous variables.

To understand it better, we look at how to implement Sargan's test manually: For the 2SLS estimator, the test statistic is Sargan's statistic, typically calculated as  $N \times R^2$  from a regression of the IV residuals on the full set of instruments.

```
. ivregress 2sls rent pcturban (hsngval = faminc i.region)
. estat overid
```

Tests of overidentifying restrictions:

```
Sargan (score) chi2(3) = 11.2877 (p = 0.0103)
Basmann chi2(3)      = 12.8294 (p = 0.0050)
```

```
. predict res, residual

. reg res pcturban faminc i.region

. disp e(N)*e(r2)
11.287665
```

### 4.4.1 Implementation in Stata

In Stata, there are different ways to do over-identification test, `ivreg2` reports a comprehensive set of tests; `overid` command does the over-identification test after the `ivreg` command.

`ivreg2` with `gmm` option returns J test; it reports Sargan's test without this option.

`ivreg2` also reports C test statistic, with `ortho()`. If the C test rejects the null, and J test without the suspect instruments fail to reject null, then the suspect instruments are indeed the ones are not exogenous.

## 4.5 Weak Instruments

### 4.5.1 Problem with the cure

An instrument needs to satisfy two criteria: orthogonality and relevance. We need instruments to be orthogonal to the error term. We can verify the orthogonality condition by Sargan's test if there are extra instruments.

It turns out instrument relevance is important too: if instruments are weak, the usual large-sample approximations are poor. It helps to keep three cases apart. With fixed, nonzero first-stage relevance, IV remains *consistent* but can be badly *biased* in finite samples (and biased toward OLS), with the normal approximation working poorly. Under weak-instrument (local-to-zero) asymptotics, the limiting distribution is *nonstandard* and not centered in the usual way, so conventional standard errors and *t*-statistics mislead. Only in the limiting case of complete irrelevance ( $\Pi = 0$ ) is the parameter unidentified.

To see the problem, suppose

$$\mathbf{y} = \mathbf{X}\beta + \mathbf{u}, \quad \mathbf{E}(\mathbf{u}\mathbf{u}') = \sigma_u^2 \mathbf{I}, \quad (4.10)$$

$$\mathbf{X} = \mathbf{Z}\Pi + \mathbf{v}, \quad \mathbf{E}(\mathbf{v}\mathbf{v}') = \sigma_v^2 \mathbf{I}, \quad (4.11)$$

and

$$\mathbf{E}(\mathbf{Z}\mathbf{u}) = \mathbf{0}. \quad (4.12)$$

We can see here  $\mathbf{Z}$  is exogenous. However, the model does not say anything about relevance. To illustrate the problem caused by weak instrument, suppose we have only one endogenous variable and one instrument.

$$\hat{\beta}_{2sls} = \frac{\mathbf{Z}'\mathbf{y}}{\mathbf{Z}'\mathbf{X}} = \frac{\mathbf{Z}'(\mathbf{X}\beta + \mathbf{u})}{\mathbf{Z}'\mathbf{X}} = \beta + \frac{\mathbf{Z}'\mathbf{u}}{\mathbf{Z}'\mathbf{X}}. \quad (4.13)$$

If  $\mathbf{Z}$  is irrelevant, or,  $\Pi = 0$ , then

$$\hat{\beta}_{2sls} - \beta = \frac{\mathbf{Z}'\mathbf{u}}{\mathbf{Z}'\mathbf{v}} = \frac{\frac{1}{\sqrt{n}} \sum_{i=1}^n Z_i u_i}{\frac{1}{\sqrt{n}} \sum_{i=1}^n Z_i v_i} \xrightarrow{d} \frac{z_u}{z_v}, \quad (4.14)$$

where (normalizing  $\text{Var}(Z_i) = 1$ , which is harmless since any common scale factor cancels in the ratio)

$$\begin{bmatrix} z_u \\ z_v \end{bmatrix} \sim N(0, \begin{bmatrix} \sigma_u^2 & \sigma_{uv} \\ \sigma_{uv} & \sigma_v^2 \end{bmatrix}). \quad (4.15)$$

Therefore, if  $\mathbf{Z}$  is irrelevant,  $\beta_{2sls}$  is inconsistent. Also, the distribution of the bias is Cauchy-like (the ratio of correlated normals).

This is a case where the cure might be worse than the disease itself: the bias can be big comparing to the bias an OLS estimate suffers.

#### 4.5.2 Weak Instrument Tests

There are a variety of weak-instruments tests proposed. Most of them are based on so-called weak-instruments asymptotics and a new parameter called “concentration parameter”  $\mu^2 = \Pi' Z' Z \Pi / \sigma_v^2$ . Sample size only enters the distribution through  $\mu^2$ .

With weak-instruments asymptotics, IV estimators are no longer consistent, and they are not normal asymptotically. Most test statistics (J test, etc.) do not have normal or  $\chi^2$  distributions anymore.

Now I list the following tests in the order of recommended level by James Stock:

1. Moreira (2003) conditional likelihood ratio test (CLR).

Advantages of this test: a. Nearly optimal among invariant similar tests (Andrews, Moreira and Stock 2006). b. Implemented in Stata as `condivreg`.

Disadvantages:

- a. Complicated. Only developed so far for one endogenous variable case.
2. Stock-Yogo bias method and size method.

Stock and Yogo (2005) provide critical values for both methods: one is to control the size of bias, the other one is to control the size of a Wald test of  $\beta = \beta_0$ . Bias method is more frequently used. In the case of multiple endogenous variables, the Cragg-Donald statistics is used to compare with the critical values. It is implemented in Stata as part of the `ivreg2` command, but it’s only available for the situation there are at least two excluded variables (meaning the number of instruments minus the number of endogenous variables).

3. Anderson-Rubin confidence intervals.

#### 4.6 When is endogeneity a problem and what can we say from an IV regression

In the model of

$$\mathbf{y} = \mathbf{X}\beta + \mathbf{u}, \quad \mathbf{E}(\mathbf{u}\mathbf{u}') = \sigma_u^2 \mathbf{I}, \quad (4.16)$$

$$\mathbf{X} = \mathbf{Z}\Pi + \mathbf{v}, \quad \mathbf{E}(\mathbf{v}\mathbf{v}') = \sigma_v^2 \mathbf{I}, \quad (4.17)$$

The null hypothesis  $H_0 : \beta = \beta_0$ . Anderson-Rubin statistic is the F statistic in the regression of  $y - X\beta_0$  on  $Z$ , the F test that the coefficient vector on  $Z$  is zero (under the model that coefficient equals  $\Pi(\beta - \beta_0)$ , so it is zero exactly when  $\beta = \beta_0$ ):

$$AR(\beta_0) = \frac{(y - X\beta_0)' P_Z (y - X\beta_0) / l}{(y - X\beta_0)' M_Z (y - X\beta_0) / (N - l)} \quad (4.18)$$

where  $l$  is the number of instruments. The idea of AR confidence interval is to collect all values of  $\beta_0$  at which we fail to reject  $H_0 : \beta = \beta_0$ .

4. First-stage F test.

The rule of thumb for first-stage F test is  $F > 10$  for a single instrument case, the more instruments, the higher it gets.

5. Kleibergen's LM test.

This test is dominated by the CLR test, thus no longer the optimal test to use.

6. First-stage  $R^2$ , or partial  $R^2$ , etc., are not recommended.

### 4.6 When is endogeneity a problem and what can we say from an IV regression

Let's consider a model of housing price on no. of rooms and square footage.

$$P = \alpha R + \beta F + \epsilon \quad (4.19)$$

Should square footage really be there? We are not interested in, given square footage, how much money we need to pay for an additional room. We are generally interested in how much money we need to pay for an additional footage. However, if we do not include footage, then we "sort of" have endogeneity problem due to omitted variable.

Another example would be education's impact on wage. Say IQ is an omitted variable. Do we want to say something about education's impact on wage controlling on IQ, or not?

So it depends on what questions to ask. If I am a builder and ask what an extra bedroom on the already built house would bring to me, I need to control for footage. If I am generally asking what an extra bedroom would cost, then I am implying what that bedroom and associated footage would

## 4 Endogeneity

cost me. In that case, without controlling for footage is not a problem, since the question is: what an additional bedroom (AND the associated average square footage) would cost me.

In the wage and education example, if I am asking: what's the effect of education on wage? I implicitly ask what would be my pay raise if I go to college (given IQ, of course), for example. Then we are concerned that this model of

$$wage = \alpha edu + \epsilon \tag{4.20}$$

will generate biased estimate of  $\alpha$ , because we are not controlling for IQ. But if I am an employer and interested in hiring a person, what would I have to pay extra to hire a college graduate vs. a high school graduate? Then I don't have to include that IQ variable, because implicitly I am asking: what extra I need to pay for that college graduate who comes with a higher IQ, and everything else that is associated with higher education.

In addition, what an IV regression really gives us is part of the causal inference of  $X$  (suspected to be endogenous) on  $y$ , namely the part induced by  $Z$ . For example, we are interested in the relationship between college education and wage. If we use distance to college as an instrument, then our inference is the effect of college education from those who decide to go because of proximity of the college.

## 5 General Method of Moments (GMM)

### 5.1 The Method of Moments (MOM)

A population moment  $\gamma$  can be defined as the expectation of some continuous function  $g$  of a random variable  $x$ :

$$\gamma = E[g(x)] \quad (5.1)$$

On the other hand, a sample moment is the sample version of the population moment in a particular sample:

$$\hat{\gamma} = \frac{1}{n} \sum [g(x)] \quad (5.2)$$

### 5.2 OLS as a moment problem

Consider the simple linear regression

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \mathbf{u}, \quad \mathbf{u} \sim \text{IID}(\mathbf{0}, \sigma^2). \quad (5.3)$$

If the model is correctly specified, then

$$\mathbf{E}(\mathbf{X}'\mathbf{u}) = \mathbf{0}. \quad (5.4)$$

The MOM principle suggests that we replace the left-hand side with its sample analog  $\frac{1}{n}\mathbf{X}'(\mathbf{y} - \mathbf{X}\boldsymbol{\beta})$ .

Since we know that the true  $\boldsymbol{\beta}$  sets the population moment equal to zero in expectation, it seems reasonable to assume that a good choice of  $\hat{\boldsymbol{\beta}}$  would be one that sets the sample moment to zero. The MOM procedure suggests an estimate of  $\boldsymbol{\beta}$  that solves

$$\frac{1}{n}\mathbf{X}'(\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}) = \mathbf{0}. \quad (5.5)$$

The MOM estimator is

$$\hat{\boldsymbol{\beta}} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y}, \quad (5.6)$$

which is the same as the OLS estimator.

### 5.3 IV as a moment problem

Consider the simple linear regression

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \mathbf{u}, \quad \mathbf{u} \sim \text{IID}(\mathbf{0}, \sigma^2). \quad (5.7)$$

If one or more of the regressors is endogenous, then

$$\mathbf{E}(\mathbf{X}'\mathbf{u}) \neq \mathbf{0}. \quad (5.8)$$

We have to find an instrumental variable  $\mathbf{Z}$  which is

$$\mathbf{E}(\mathbf{Z}'\mathbf{u}) = \mathbf{0}. \quad (5.9)$$

Or,

$$\mathbf{E}(\mathbf{Z}'(\mathbf{y} - \mathbf{X}\boldsymbol{\beta})) = \mathbf{0}. \quad (5.10)$$

The sample analogy of this is

$$\frac{1}{n}\mathbf{Z}'(\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}) = \mathbf{0}. \quad (5.11)$$

When the model is *exactly identified* – the number of instruments equals the number of regressors, so  $\mathbf{Z}'\mathbf{X}$  is square and nonsingular – the sample moment condition has a unique solution, the simple IV estimator

$$\hat{\boldsymbol{\beta}} = (\mathbf{Z}'\mathbf{X})^{-1}\mathbf{Z}'\mathbf{y}. \quad (5.12)$$

When the model is *overidentified* – more instruments than regressors –  $\mathbf{Z}'\mathbf{X}$  is no longer square and we cannot set all sample moments to zero simultaneously. We instead minimize a quadratic form in the moments, which gives the two-stage least squares (2SLS) estimator

$$\hat{\boldsymbol{\beta}} = (\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1}\mathbf{X}'\mathbf{P}_Z\mathbf{y}, \quad \mathbf{P}_Z = \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}', \quad (5.13)$$

the special case of the GMM estimator below with weighting matrix  $(\mathbf{Z}'\mathbf{Z})^{-1}$ .

### 5.4 The Generalized Method of Moments

The expectation  $E(Y^r)$  for any  $r = 1, 2, \dots$  is called the  $r^{th}$  (raw) moment of  $Y$ . The expectation  $E[(Y - E(Y))^r]$  is called the  $r^{th}$  centered moment of  $Y$ .

The mean is the first raw moment.

The variance is the second centered moment.

The third centered moment measures the skewness of the distribution.

The fourth centered moment measures the kurtosis of the distribution. Interpreted as a measure of “fatness of tails”.

The standardized kurtosis is

$$k = \frac{E[(Y - E(Y))^4]}{E[(Y - E(Y))^2]^2}. \quad (5.14)$$

For a normal distribution,  $k = 3$ .

For a  $t$  distribution with  $v > 4$  degrees of freedom,  $k = 3 + 6/(v - 4) > 3$ , i.e., the  $t$  distribution has fatter tails than a normal distribution (the excess kurtosis  $6/(v - 4)$  shrinks toward 0 as  $v$  grows).

The distribution function of a random variable captures all information about the random variable. If the moment generating function exists in a neighborhood of zero, the full set of moments also determines the distribution (without such a condition this can fail: the lognormal distribution is not determined by its moments).

This matters directly for GMM: GMM only ever matches a finite set of moment conditions, and if the underlying distribution is not uniquely pinned down by its (even infinite) sequence of moments, then no amount of moment-matching – however many moment conditions we add – can recover the full data-generating distribution. This is precisely the sense in which GMM asks for less than MLE: MLE assumes we know the full distributional family and estimates its parameters (recovering the whole distribution when correctly specified), while GMM only requires correctly specified moment conditions and only ever identifies the parameters entering those moments, not the full distribution. When the distributional assumption behind MLE is correct, MLE is efficient and gives you more (the whole distribution); when it is not, GMM's weaker requirements make it more robust.

This distinction underlies the relative strengths and weaknesses of ML and GMM.

## 5.5 GMM

The statistical model takes the general form

$$E[m(Y_i; \theta_0)] = 0 \quad (5.15)$$

where -  $Y_1, \dots, Y_n$  are random variables from which the sample  $y_1, \dots, y_n$  is drawn, -  $m(Y, \theta)$  is a function specifying the model, -  $\theta_0$  is the “true value” of the parameter.

$E[m(Y_i; \theta_0)] = 0$  are called the population moment conditions.

Two ideas behind GMM:

1. Replace the population mean  $E[\cdot]$  with the sample mean calculated from the observed sample  $y_1, \dots, y_n$ .
2. Since  $E[m(Y_i; \theta_0)] = 0$ , choose  $\hat{\theta}_{GMM}$  to make  $\frac{1}{n} \sum_{i=1}^n m(y_i; \hat{\theta}_{GMM})$  as close to zero as possible.

Define the notation

$$\bar{m}(\theta) = \frac{1}{n} \sum_{i=1}^n m(y_i; \theta). \quad (5.16)$$

$\hat{\theta}_{GMM}$  is chosen to make  $\bar{m}(\theta)' \bar{m}(\theta)$  as close to zero as possible.

More generally,  $\hat{\theta}_{GMM}$  is chosen to minimize  $\bar{m}(\theta)' W \bar{m}(\theta)$  for some weighting matrix  $W$ .

The choice of  $W$  matters only under overidentification. With as many moment conditions as parameters, all sample moments can be driven to zero exactly and  $W$  drops out – which is why

## 5 General Method of Moments (GMM)

the just-identified IV estimator above did not need one. When there are more moments than parameters, the efficient choice is  $W = S^{-1}$ , where  $S$  is the asymptotic variance of the sample moments,  $S = \lim_{n \rightarrow \infty} \text{Var}(\sqrt{n} \bar{m}(\theta_0))$ . With that weighting,

$$\sqrt{n}(\hat{\theta}_{GMM} - \theta_0) \xrightarrow{d} N(0, (G' S^{-1} G)^{-1}), \quad G = E \left[ \frac{\partial m(Y_i; \theta_0)}{\partial \theta'} \right], \quad (5.17)$$

which is the GMM counterpart of the information-matrix bound of the previous chapter. Since  $S$  depends on  $\theta_0$ , in practice one starts from any consistent weighting (often the identity, or  $(\mathbf{Z}'\mathbf{Z})^{-1}$ ), estimates  $\hat{S}$  from the resulting residuals, and re-minimizes: two-step GMM. Stopping at  $W = (\mathbf{Z}'\mathbf{Z})^{-1}$  is exactly 2SLS, which is therefore efficient GMM only under homoskedasticity.

### 5.5.1 An example

Let's see an example with GMM, using the same simulated data as before. We have the same situation as before,  $X$  is endogenous. We are doing GMM version of 2sls.

Here I use R's "gmm" library, which makes things easy. Its first two arguments are **g** and **x**: with the formula interface, **g** is the model formula whose residual forms the moment function, and **x** is a formula giving the instruments. The moment condition in this example is  $E(\mathbf{z}'\mathbf{u}) = \mathbf{0}$ , where **z** is the instrument – the variable **w** in the simulated data – and  $u$  is the residual from the structural equation. Note that **W** above denotes the *weighting matrix* in  $\bar{m}'W\bar{m}$ , which is a different object entirely, so the instrument is written **z** here.

```
## DGP: data$y <- data$x + data$z + data$u

set.seed(66)
nobs=10000
nDim = 3
sdxx = 1
sdww=1
sdzz=1

## here we have three variables x,z,w.
## z is the omitted variable,x and z are correlated, w is the instrument, which is correlated
crxz=.6
crzw=0
crxw=.5

covarMat = matrix( c(sdxx^2, crxz, crxw, crxz, sdzz^2, crzw, crxw, crzw, sdww^2 ), nrow=nDim,
covarMat
```

```
      [,1] [,2] [,3]
[1,]  1.0  0.6  0.5
[2,]  0.6  1.0  0.0
[3,]  0.5  0.0  1.0
```

```

data = data.frame(mvrnorm(n=nobs, mu=rep(0,nDim), Sigma=covarMat ))
names(data) <- c('x','z','w')
data$u <- rnorm(nobs,0,1)
# dgp
data$y <- data$x + data$z + data$u

gmm.fit = gmm(y ~ x, x = ~w, data = data)
summary(gmm.fit)

```

Call:

```
gmm(g = y ~ x, x = ~w, data = data)
```

Method: twoStep

Kernel: Quadratic Spectral

Coefficients:

|             | Estimate   | Std. Error | t value    | Pr(> t )    |
|-------------|------------|------------|------------|-------------|
| (Intercept) | 1.1531e-02 | 1.4151e-02 | 8.1481e-01 | 4.1518e-01  |
| x           | 9.5129e-01 | 2.9971e-02 | 3.1740e+01 | 4.3442e-221 |

J-Test: degrees of freedom is 0

|              | J-test               | P-value |
|--------------|----------------------|---------|
| Test E(g)=0: | 1.16401567357162e-26 | *****   |

```

# It returns the same estimates as the 2sls results.
tsls.model <- ivreg(y ~ x | w, data=data)
summary(tsls.model)

```

Call:

```
ivreg(formula = y ~ x | w, data = data)
```

Residuals:

| Min       | 1Q        | Median   | 3Q       | Max      |
|-----------|-----------|----------|----------|----------|
| -6.144530 | -0.962464 | 0.003414 | 0.944728 | 5.340560 |

Coefficients:

|             | Estimate | Std. Error | t value | Pr(> t )   |
|-------------|----------|------------|---------|------------|
| (Intercept) | 0.01153  | 0.01433    | 0.804   | 0.421      |
| x           | 0.95129  | 0.03007    | 31.637  | <2e-16 *** |

Diagnostic tests:

| df1 | df2 | statistic | p-value |
|-----|-----|-----------|---------|
|-----|-----|-----------|---------|

## 5 General Method of Moments (GMM)

```
Weak instruments      1 9998      3074.7 <2e-16 ***
Wu-Hausman           1 9997       807.8 <2e-16 ***
Sargan               0  NA         NA      NA
```

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.431 on 9998 degrees of freedom

Multiple R-Squared: 0.5007, Adjusted R-squared: 0.5006

Wald test: 1001 on 1 and 9998 DF, p-value: < 2.2e-16

Two things about that agreement are worth being explicit about, because it is easy to over-generalise from it. First, GMM and 2SLS coincide *numerically* here because the model is just-identified: one instrument ( $\mathbf{w}$ ) for one endogenous regressor ( $\mathbf{x}$ ), so  $\mathbf{Z}'\mathbf{X}$  is square and the sample moments can be set to zero exactly, whatever the weighting matrix. The output says so – the  $J$ -test reports “degrees of freedom is 0”, and its value ( $\sim 10^{-26}$ ) is numerical noise rather than evidence. Second, `gmm()` is using an optimal two-step weighting with a heteroskedasticity-and-autocorrelation-consistent kernel (note Method: `twoStep` and Kernel: `Quadratic Spectral` in the output). Under *over*-identification that choice would generally **not** reproduce 2SLS, for exactly the reason given above:  $W = (\mathbf{Z}'\mathbf{Z})^{-1}$  is the efficient weighting only under homoskedasticity.

In OLS case, it would be  $E(Xu) = 0$ .

```
gmm.ols = gmm(y ~ x, x = ~x, data = data)
summary(gmm.ols)
```

Call:

```
gmm(g = y ~ x, x = ~x, data = data)
```

Method: `twoStep`

Kernel: `Quadratic Spectral`

Coefficients:

|             | Estimate   | Std. Error | t value     | Pr(> t )  |
|-------------|------------|------------|-------------|-----------|
| (Intercept) | -0.0042841 | 0.0126399  | -0.3389314  | 0.7346614 |
| x           | 1.5961141  | 0.0128364  | 124.3424084 | 0.0000000 |

J-Test: degrees of freedom is 0

|                 | J-test               | P-value |
|-----------------|----------------------|---------|
| Test $E(g)=0$ : | 2.66671052451529e-28 | *****   |

```
# It returns the same estimates as the OLS results.
ols <- lm(y ~ x, data=data)
summary(ols)
```

```

Call:
lm(formula = y ~ x, data = data)

Residuals:
    Min       1Q   Median       3Q      Max
-5.3506 -0.8612  0.0123  0.8504  4.9096

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.004284   0.012843  -0.334   0.739
x             1.596114   0.013079 122.035 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.284 on 9998 degrees of freedom
Multiple R-squared:  0.5983,    Adjusted R-squared:  0.5983
F-statistic: 1.489e+04 on 1 and 9998 DF,  p-value: < 2.2e-16

```

## 5.6 A few concepts of conditioning

### 5.6.1 Independence

If  $X$  and  $Y$  are independent then

$$f(x, y) = f(x)f(y) \quad (5.18)$$

and hence

$$f(y|x) = f(y). \quad (5.19)$$

If  $X$  and  $Y$  are independent then

$$E[g(X)h(Y)] = E[g(X)] \cdot E[h(Y)] \quad (5.20)$$

and hence

$$\text{Cov}[g(X), h(Y)] = 0. \quad (5.21)$$

i.e. all functions of  $X$  and  $Y$  are uncorrelated.

### 5.6.2 Law of Iterated Expectations

$$E[Y] = E[E(Y|X)]. \quad (5.22)$$

### 5.6.3 Dependence Concepts

$X, Y$  independent:

$$Cov[g(X), h(Y)] = 0 \quad (5.23)$$

$X, Y$  uncorrelated:

$$Cov[X, Y] = 0 \quad (5.24)$$

$E[Y|X] = 0$ :

$$Cov[g(X), Y] = 0 \quad (5.25)$$

These are listed in decreasing strength, and the implications run one way only: independence  $\Rightarrow$  mean independence  $\Rightarrow$  uncorrelatedness. The middle condition is what regression actually needs. Neither converse holds:  $X$  and  $Y$  can be uncorrelated with  $E[Y | X] \neq 0$  (take  $Y = X^2 - 1$  with  $X \sim N(0, 1)$ , so  $Cov[X, Y] = E[X^3] = 0$  while  $E[Y | X] = X^2 - 1$ ), and mean independence permits dependence through higher moments (any pure heteroskedasticity,  $Y = X\epsilon$  with  $\epsilon$  independent of  $X$  and mean zero). Keeping the three apart matters because the moment conditions below need only the middle one.

### 5.6.4 Regression

A regression model is a model of  $E[Y_i|X_i]$ . For example,

$$Y_i = \beta_0 + \beta_1 X_i + u_i \quad (5.26)$$

where  $E[u_i|X_i] = 0$ .

### 5.6.5 GMM regression

The regression model

$$Y_i = \beta_0 + \beta_1 X_i + u_i, \quad E[u_i|X_i] = 0 \quad (5.27)$$

implies the moment condition

$$E[u_i] = 0 \quad \text{and} \quad E[X_i u_i] = 0 \quad (5.28)$$

That is,

$$E[Y_i - \beta_0 - \beta_1 X_i] = 0 \quad (5.29)$$

$$E[X_i(Y_i - \beta_0 - \beta_1 X_i)] = 0 \quad (5.30)$$

The sample moment conditions are

$$\frac{1}{n} \sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i) = 0 \quad (5.31)$$

$$\frac{1}{n} \sum_{i=1}^n x_i (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i) = 0 \quad (5.32)$$

These are just normal equations for OLS.

A characteristic of GMM: the specification of the model generates the estimator. i.e. only  $E[Y_i|X_i] = \beta_0 + \beta_1 X_i$  is assumed.

Note there are no assumptions that  $u_i$  is homoscedastic, not autocorrelated or normally distributed. These properties affect the statistical properties of the GMM estimator, not its definition.



## 6 Censored, truncated or selected data

### 6.1 Compare three different cases

A sample is truncated if some observations are systematically excluded from the sample. For example, suppose we are interested in relationship between people's income( $y$ ) and education( $x$ ). If we have observations of both  $y$  and  $x$  only for people whose income is above \$20,000 per year, then we have a truncated sample. A sample is censored if no observations have been systematically excluded but some of information contained in them has been suppressed. In other words, in a truncated sample, you don't have observations which have been truncated; while in a censored sample, you do have observations but you don't have full information. The idea of "censoring" is that some data above or below a threshold are mis-reported at the threshold. If we use the same example of people's income, if people with income lower than \$20,000 per year do not appear in the sample at all—we have no observations of either  $y$  or  $x$  for them—then we have a truncated sample. If we have observations of  $x$  for people whose income above AND below that limit, but no information on exactly how much their income is for those whose income below the limit (only thing we know is that it's lower than \$20,000 per year), then we have a censored sample.

The problem of sample selection arises when no observations have been systematically excluded but some information has been suppressed, just like censoring. However, in sample selection problems, we have information on how the "selection" happens. In censoring, usually a constant "threshold" is set. Observations on dependent variable for those above (or below) that criteria are missing. Using the income example again, the threshold is set to be \$20,000, we have no information on exactly how much their income is for those whose income below the limit (only thing we know is that it's lower than \$20,000 per year). When we have a third variable that determines the censoring situation, then we call it sample selection. For example, we have observed only union member's information on income, but we observe information on education for everybody, including union members or non-members. In this example, union membership is an additional variable that "selects" the sample which are involved in the estimation of effect of education on income.

### 6.2 An illustration for truncated or censored data

Consider the model

$$y_t^0 = \beta_1 + \beta_2 x_t + u_t, \quad u_t \sim NID(0, \sigma^2), \quad (6.1)$$

where  $y_t^0$  is a latent variable. We actually observe  $y_t$ , which differs from  $y_t^0$ , because it's either truncated or censored.

Suppose that censorship or truncation occurs whenever  $y_t^0$  is less than 0. Clearly, the larger the error term  $u_t$ , the larger is  $y_t^0$  and thus the greater must be the probability that  $y_t^0 \geq 0$ . This probability also depends on  $x_t$ . So for the sample we observe,  $u_t$  does not have conditional mean 0

and is not uncorrelated with  $x_t$ . OLS using truncated or censored samples (OLS of  $y_t$  on  $x_t$ ) yields biased and inconsistent estimators. Normally we would like to draw inference on the population which is represented by the full sample. Shown in the figure below, ideally we would have the OLS regression line (if we had the full sample), which is very close to the “true” regression line (which is the mechanism generates the data). We would have the “OLS censored sample” line if we run OLS on the censored sample; we would have “OLS truncated sample” if we run OLS on the truncated sample. It shows that both of them are severely biased from the “true” regression line.

Here is a graph to illustrate from Davidson and MacKinnon.




Figure 6.1: Effects of truncation and censoring from Davidson and MacKinnon

### 6.3 Truncated Models

For truncated data, a consistent estimator comes from maximum likelihood estimator (MLE). If we assume the error terms in the latent variable model has a known distribution, then MLE estimator can be applied. The most popular choice is Gaussian.

$$\begin{aligned}
 \Pr(y_t^0 \geq 0) &= \Pr(\mathbf{X}_t\beta + u_t \geq 0) \\
 &= 1 - \Pr(u_t/\sigma < -\mathbf{X}_t\beta/\sigma) \\
 &= 1 - \Phi(-\mathbf{X}_t\beta/\sigma) = \Phi(\mathbf{X}_t\beta/\sigma)
 \end{aligned} \tag{6.2}$$

The density of  $y_t$  is proportional to the density of  $y_t^0$  when  $y_t^0 \geq 0$  and  $y_t$  is observed. It is 0 elsewhere. The factor of proportionality, which is needed to ensure that the density integrates to unity, is the inverse of the probability that  $y_t^0 \geq 0$ . The density of  $y_t$  can be written as

$$\frac{\sigma^{-1}\phi((y_t - \mathbf{X}_t\beta)/\sigma)}{\Phi(\mathbf{X}_t\beta/\sigma)} \tag{6.3}$$

This implies the log-likelihood function,

$$\ell(\mathbf{y}, \beta, \sigma) = -\frac{n}{2} \log(2\pi) - n \log(\sigma) - \frac{1}{2\sigma^2} \sum_{t=1}^n (y_t - \mathbf{X}_t \beta)^2 - \sum_{t=1}^n \log \Phi(\mathbf{X}_t \beta / \sigma). \quad (6.4)$$

This can be estimated by MLE. In Stata, `truncreg` is the command to do truncated regression model.

## 6.4 Censored Models

The most popular censored model is Tobit model.

$$\begin{aligned} y_t^0 &= \mathbf{X}_t \beta + u_t, \quad u_t \sim NID(0, \sigma^2) \\ y_t &= y_t^0 \text{ if } y_t^0 > 0; \quad y_t = 0 \text{ otherwise.} \end{aligned} \quad (6.5)$$

We see that

$$\begin{aligned} \Pr(y_t = 0) &= \Pr(y_t^0 \leq 0) = \Pr(\mathbf{X}_t \beta + u_t \leq 0) \\ &= \Pr(u_t / \sigma < -\mathbf{X}_t \beta / \sigma) \\ &= \Phi(-\mathbf{X}_t \beta / \sigma) \end{aligned} \quad (6.6)$$

The contribution to the log-likelihood function made by observations with  $y_t = 0$  is

$$\ell_t(y_t, \beta, \sigma) = \log \Phi(-\mathbf{X}_t \beta / \sigma). \quad (6.7)$$

If  $y_t$  is positive, the contribution to the log-likelihood is the logarithm of the density,

$$\log\left(\frac{1}{\sigma} \phi((y_t - \mathbf{X}_t \beta) / \sigma)\right). \quad (6.8)$$

The log-likelihood function of the tobit model is

$$\sum_{y_t=0} \log \Phi(-\mathbf{X}_t \beta / \sigma) + \sum_{y_t>0} \log\left(\frac{1}{\sigma} \phi((y_t - \mathbf{X}_t \beta) / \sigma)\right) \quad (6.9)$$

This can be estimated by MLE. In Stata, `tobit` or `intreg` can be used for censoring models.

## 6.5 Sample Selection

The sample selection models differ from censored model in that it involves a different variable (from  $y$  itself) to determine the censorship (selection).

Suppose that  $y_t^0$  and  $z_t^0$  are two latent variables, generated by the bivariate process

$$\begin{bmatrix} y_t^0 \\ z_t^0 \end{bmatrix} = \begin{bmatrix} \mathbf{X}_t \beta \\ \mathbf{W}_t \gamma \end{bmatrix} + \begin{bmatrix} u_t \\ v_t \end{bmatrix}, \quad \begin{bmatrix} u_t \\ v_t \end{bmatrix} \sim NID(\mathbf{0}, \begin{bmatrix} \sigma^2 & \rho\sigma \\ \rho\sigma & 1 \end{bmatrix}) \quad (6.10)$$

## 6 Censored, truncated or selected data

Here  $\rho$  is the correlation between  $u_t$  and  $v_t$ ; since  $\text{Var}(u_t) = \sigma^2$  and  $\text{Var}(v_t) = 1$ , their covariance is  $\rho\sigma$ .

We observe  $y_t$  and  $z_t$ :

$$y_t = y_t^0 \text{ if } z_t^0 > 0; y_t \text{ unobservable otherwise}; z_t = 1 \text{ if } z_t^0 > 0; z_t = 0 \text{ otherwise.} \quad (6.11)$$

There are two types of observations, ones we observe  $y_t = y_t^0$  and  $z_t = 1$ , along with both  $\mathbf{X}_t$  and  $\mathbf{W}_t$ , and ones we observe only  $z_t = 0$  and  $\mathbf{W}_t$ .

Each observation contributes to the likelihood function by

$$I(z_t = 0)\Pr(z_t = 0) + I(z_t = 1)\Pr(z_t = 1)f(y_t^0|z_t = 1), \quad (6.12)$$

Under the normality assumption,

$$u_t = \rho\sigma v_t + e_t \quad (6.13)$$

where  $e_t$  is independent of  $v_t \sim N(0, 1)$ , with  $\text{Var}(e_t) = \sigma^2(1 - \rho^2)$ . A useful fact about the standard normal distribution is that

$$E(v_t|v_t > -x) = \lambda(x) = \frac{\phi(x)}{\Phi(x)} \quad (6.14)$$

and the function  $\lambda(x)$  is called the inverse Mills ratio.

The log-likelihood function can be shown to be

$$\sum_{z_t=0} \log \Phi(-\mathbf{W}_t\gamma) + \sum_{z_t=1} \log\left(\frac{1}{\sigma}\phi((y_t - \mathbf{X}_t\beta)/\sigma)\right) + \sum_{z_t=1} \log \Phi\left(\frac{\mathbf{W}_t\gamma + \rho(y_t - \mathbf{X}_t\beta)/\sigma}{(1 - \rho^2)^{1/2}}\right) \quad (6.15)$$

So this model can be estimated by MLE.

It is also popular to use Heckman's two-step method.

Heckman's method is based on the fact that the original latent model can be rewritten as

$$y_t = \mathbf{X}_t\beta + \rho\sigma v_t + e_t \quad (6.16)$$

Here the error term  $u_t$  is divided into two parts, one perfectly correlated with  $v_t$ , and one independent of  $v_t$ .

In the first step, an ordinary probit model is used to obtain consistent estimates  $\hat{\gamma}$  of the parameters of the selection equation.

In the second step, the unobserved  $v_t$  is replaced by the selectivity regressor  $\frac{\phi(\mathbf{W}_t\hat{\gamma})}{\Phi(\mathbf{W}_t\hat{\gamma})}$ .

Therefore, the regression becomes

$$y_t = \mathbf{X}_t\beta + \rho\sigma \frac{\phi(\mathbf{W}_t\hat{\gamma})}{\Phi(\mathbf{W}_t\hat{\gamma})} + e_t \quad (6.17)$$

Or,

$$y_t = \mathbf{X}_t\beta + \rho\sigma\hat{\lambda}_t + e_t \quad (6.18)$$

(Strictly, conditional on selection the error also contains  $\rho\sigma(v_t - \lambda_t)$ , which is why the second-step OLS standard errors require correction; the coefficient  $\rho\sigma$  on  $\hat{\lambda}_t$  is unaffected.)

One thing to note in Heckman model is that in many situations,  $\hat{\lambda}$  is believed to be highly collinear with  $\mathbf{X}_t$ , if  $\mathbf{W}_t$  is the same as  $\mathbf{X}_t$ . However, the model can still be estimated due to the nonlinearity of the model. Nevertheless, it becomes a regular practice to require  $\mathbf{W}_t$  contains at least one extra variable than  $\mathbf{X}_t$ , which is sometimes called exclusion restriction. In many situations,  $\mathbf{W}_t$  contains all  $\mathbf{X}_t$  variables, and at least one more variable. The reason to contain all  $\mathbf{X}_t$  variables is because the selection is “endogenous” in the sense that  $y_t$  is a factor in determining selection; this is modeled by including all the  $\mathbf{X}_t$ ’s.

In Stata, `heckman` is the command to do sample selection models. It has options to do a Heckman two-step estimation or MLE estimation. Usually MLE is preferred.

## 6.6 Switching Regression (Treatment-Effects Model)

There is another situation that is similar to sample selection model: we observe  $y$  for  $z = 1$  and  $z = 0$ . In the example of effect of education on income, we observe both union member’s income and non-member’s income. In this case, we have switching regression model or treatment-effect model. The treatment effect model estimates the effect of an endogenous binary treatment  $z_t$  (treatment, program participation, etc.) on a continuous, fully-observed variable,  $y_t$ , conditional on the independent variables  $x_t$  and  $w_t$ .

Suppose  $z_t^0$  is a latent variable.

$$\begin{bmatrix} y_t \\ z_t^0 \end{bmatrix} = \begin{bmatrix} \mathbf{X}_t\beta \\ \mathbf{W}_t\gamma \end{bmatrix} + \delta \begin{bmatrix} z_t \\ 0 \end{bmatrix} + \begin{bmatrix} u_t \\ v_t \end{bmatrix}, \quad \begin{bmatrix} u_t \\ v_t \end{bmatrix} \sim NID(\mathbf{0}, \begin{bmatrix} \sigma^2 & \rho\sigma \\ \rho\sigma & 1 \end{bmatrix}) \quad (6.19)$$

We observe  $y_t$  and  $z_t$ :

$$z_t = 1 \text{ if } z_t^0 > 0; \quad z_t = 0 \text{ otherwise.} \quad (6.20)$$

Notice that the only difference between this model and the selection model is that we observe  $y_t$  when  $z_t = 0$  and  $z_t = 1$ . It’s not a selection, but a regime switching.

One caveat about the name. The specification above lets treatment shift only the *intercept*: there is a single  $\beta$  and a single  $\delta$ , so the treatment effect is  $\delta$  for everyone. That is the treatment-effects model. A **switching regression** in the fuller sense lets the whole coefficient vector differ by regime,

$$y_t = z_t(\mathbf{X}_t\beta_1 + u_{1t}) + (1 - z_t)(\mathbf{X}_t\beta_0 + u_{0t}), \quad z_t^0 = \mathbf{W}_t\gamma + v_t, \quad (6.21)$$

with  $u_{1t}$  and  $u_{0t}$  each allowed to correlate with  $v_t$  (two selection parameters,  $\rho_1$  and  $\rho_0$ , rather than one). This is the Roy model, and the treatment effect  $\mathbf{X}_t(\beta_1 - \beta_0)$  then varies with  $\mathbf{X}_t$  instead of being a constant. The distinction matters in practice: if the returns to education genuinely differ between union and non-union jobs, the common-slope model above will report a single  $\delta$  that is some average of them and will not reveal the heterogeneity. Fitting it requires two outcome equations rather than one intercept shift.

Both Stata (`etregress` in version 14) and SAS (`qlim`) can estimate the switching regression. The default is by MLE.



# 7 Discrete and Limited Dependent Variables

## 7.1 Binary Response Models

### 7.1.1 Probit and Logit

Let  $P_t$  denote the probability that  $y_t = 1$  conditional on the information set  $\Omega$ , which consists of exogenous and predetermined variables. A binary response model serves to model this conditional expectation. Since the values are 0 or 1, it is clear that  $P_t$  is also the expectation of  $y_t$  conditional on  $\Omega_t$ :

$$P_t \equiv \Pr(y_t = 1|\Omega_t) = E(y_t|\Omega_t). \quad (7.1)$$

Here  $X_t$  is a set of regressors. A linear probability model would set  $P_t = X_t\beta$  directly, which requires the fitted values to stay in  $[0, 1]$  – an awkward constraint, since the linear index  $X_t\beta$  is unbounded. Logit and probit avoid this problem by modeling  $P_t = F(X_t\beta)$ , mapping the unbounded index through a CDF-valued transformation into  $[0, 1]$ .

We ensure that  $0 \leq P_t \leq 1$  by specifying that

$$P_t \equiv \Pr(y_t = 1|\Omega_t) = F(\mathbf{X}_t\beta). \quad (7.2)$$

$F(x)$  is a transformation function, which has the same characteristics as the CDF of a probability distribution.

Two popular choices of  $F(x)$  are Gaussian (probit) and Logistic (logit).

The less familiar logistic function is

$$\Lambda(x) = \frac{e^x}{1 + e^x} \quad (7.3)$$

The logit model is most easily derived by assuming that

$$\log\left(\frac{P_t}{1 - P_t}\right) = \mathbf{X}_t\beta \quad (7.4)$$

which says the logarithm of the odds (the ratio of the two probabilities) is equal to  $\mathbf{X}_t\beta$ . Therefore,

$$P_t = \frac{\exp(\mathbf{X}_t\beta)}{1 + \exp(\mathbf{X}_t\beta)} = \Lambda(\mathbf{X}_t\beta) \quad (7.5)$$

### 7.1.2 MLE for binary data

The likelihood for an observation  $t$  is the probability that  $y_t = 1$  if  $y_t = 1$ , or the probability that  $y_t = 0$  if  $y_t = 0$ . The logarithm of the appropriate probability is then the contribution to the loglikelihood made by observation  $t$ . Therefore, if  $\mathbf{y}$  is an  $n$ -vector with typical element  $y_t$ , the loglikelihood function for  $\mathbf{y}$  can be written as

$$\ell(\mathbf{y}, \beta) = \sum_{t=1}^n (y_t \log F(\mathbf{X}_t \beta) + (1 - y_t) \log(1 - F(\mathbf{X}_t \beta))) \quad (7.6)$$

For the logit and probit models, this function is globally concave with respect to  $\beta$ . This implies that the first-order conditions, or likelihood equations, uniquely define the MLE estimator  $\hat{\beta}$ . These likelihood equations can be written as

$$\sum_{t=1}^n \frac{(y_t - F(\mathbf{X}_t \beta)) f(\mathbf{X}_t \beta) x_{ti}}{F(\mathbf{X}_t \beta)(1 - F(\mathbf{X}_t \beta))} = 0, \quad i = 1, \dots, k. \quad (7.7)$$

where  $f = F'$  is the density (e.g.  $\phi$  for probit). For logit,  $f = F(1 - F)$ , so the numerator cancels the denominator and the score reduces to  $\sum_t (y_t - F) x_{ti} = 0$ .

Newton's Method can be used to find  $\hat{\beta}$ .

## 7.2 Models for More than Two Discrete Responses

### 7.2.1 The Ordered Probit

Ordered Probit can be easily derived from a latent variable model.

$$y_t^0 = \mathbf{X}_t \beta + u_t, \quad u_t \sim NID(0, 1) \quad (7.8)$$

Suppose we observe  $y_t$  with three values.

$$\begin{cases} y_t = 0 & \text{if } y_t^0 < \gamma_1; \\ y_t = 1 & \text{if } \gamma_1 \leq y_t^0 < \gamma_2; \\ y_t = 2 & \text{if } y_t^0 \geq \gamma_2. \end{cases} \quad (7.9)$$

Therefore,

$$\begin{aligned} \Pr(y_t = 0) &= \Pr(y_t^0 < \gamma_1) = \Pr(\mathbf{X}_t \beta + u_t < \gamma_1) \\ &= \Pr(u_t < \gamma_1 - \mathbf{X}_t \beta) = \Phi(\gamma_1 - \mathbf{X}_t \beta) \end{aligned} \quad (7.10)$$

Similarly,

$$\Pr(y_t = 2) = 1 - \Phi(\gamma_2 - \mathbf{X}_t \beta) \quad (7.11)$$

$$\Pr(y_t = 1) = \Phi(\gamma_2 - \mathbf{X}_t \beta) - \Phi(\gamma_1 - \mathbf{X}_t \beta) \quad (7.12)$$

These probabilities depend solely on the value of the index function and on the two threshold parameters.

Two normalizations are needed before any of this can be estimated, and they are easy to pass over. First, the *location*: an intercept in  $\mathbf{X}_t\beta$  and both thresholds cannot be identified together, since adding a constant to the index and the same constant to  $\gamma_1$  and  $\gamma_2$  leaves every probability unchanged. Software resolves this by dropping the intercept or by fixing  $\gamma_1 = 0$ . Second, the *scale*: that is what setting  $\text{Var}(u_t) = 1$  above does, since multiplying the index, the thresholds and  $u_t$  by any positive constant is likewise unobservable. The analogous restriction for the multinomial logit below is  $\beta^0 = 0$ .

The loglikelihood function is

$$\ell(\beta, \gamma_1, \gamma_2) = \sum_{y_t=0} \log(\Phi(\gamma_1 - \mathbf{X}_t\beta)) + \sum_{y_t=2} \log(\Phi(\mathbf{X}_t\beta - \gamma_2)) + \sum_{y_t=1} \log(\Phi(\gamma_2 - \mathbf{X}_t\beta) - \Phi(\gamma_1 - \mathbf{X}_t\beta)) \quad (7.13)$$

### 7.2.2 The Multinomial Logit

When responses are unordered, two popular choices are multinomial logit and conditional logit. They are easy to conflate because they share the same softmax form, but they differ in what varies across the  $J + 1$  alternatives  $l = 0, \dots, J$ : in multinomial logit the *regressors* are individual-specific and only the *coefficients* vary by alternative; in conditional logit it's the reverse.

**Multinomial logit.** Here the regressor vector  $\mathbf{X}_t$  (individual-specific characteristics, e.g. age, income) is the same for every alternative, and it is the coefficient vector that is alternative-specific:

$$\Pr(y_t = l) = \frac{\exp(\mathbf{X}_t\beta^l)}{\sum_{j=0}^J \exp(\mathbf{X}_t\beta^j)} \quad \text{for } l = 0, \dots, J, \quad (7.14)$$

Here  $\mathbf{X}_t$  is a row vector of length  $k$  of individual-specific covariates, and  $\beta^l$  is a  $k$ -vector of alternative-specific parameters. For identification one category is taken as the base with  $\beta^0 = 0$  (so its numerator is 1).

**Conditional logit.** Here it is the regressors that vary by alternative (e.g. the price or characteristics of choice  $j$  itself), while the coefficient vector is common across alternatives:

$$\Pr(y_t = l) = \frac{\exp(\mathbf{W}_{tl}\beta)}{\sum_{j=0}^J \exp(\mathbf{W}_{tj}\beta)} \quad \text{for } l = 0, \dots, J, \quad (7.15)$$

Here  $\mathbf{W}_{tj}$  is a row vector of length  $k$  of alternative- $j$ -specific variables, and  $\beta$  is a single  $k$ -vector of parameters shared across all alternatives  $j = 0, \dots, J$ .

Writing *both* choice-varying regressors  $\mathbf{W}_{tl}$  and choice-specific coefficients  $\beta^l$  into the same formula, as some treatments do, conflates the two models. It is worth being precise about what that costs, because the usual complaint – that such a specification is not identified – is not true. Choice probabilities depend only on utility *differences*, so  $\Pr(y_t = l)$  is governed by  $\mathbf{W}_{tj}\beta^j - \mathbf{W}_{tl}\beta^l$ ; when the  $\mathbf{W}_{tj}$  vary independently across alternatives, those differences pin down each  $\beta^j$  separately, with no normalization required. It is a perfectly estimable model – it is what Stata's `asclogit` fits when alternative-specific coefficients are requested.

What the ambiguous formula really costs is clarity about *which normalization is in force*. Individual-specific regressors require one alternative's coefficients to be fixed ( $\beta^0 = 0$ ), since only differences across alternatives are identified; alternative-specific regressors entering with a common coefficient require no such restriction. A formula that does not say which of the two it means leaves the reader unable to tell what has been normalized, and therefore unable to interpret the coefficients.

## 8 Count data models

### 8.1 Poisson Model

If the interested variable is a count data variable, it is natural to model it as a Poisson process. The number of occurrence follows a Poisson process:

$$\Pr[Y = y] = \frac{e^{-\mu} \mu^y}{y!} \quad (8.1)$$

One of Poisson's properties is that it has equal variance as its mean:

$$E[Y] = \text{Var}[Y] = \mu \quad (8.2)$$

In a count data model such as Poisson regression model, we are modeling log of the expected counts:

$$\log(E(Y)) = \mathbf{X}_i' \boldsymbol{\beta} \quad (8.3)$$

The log-likelihood function is

$$\log L(\boldsymbol{\beta}) = \sum_{i=1}^N -\exp(X_i' \boldsymbol{\beta}) + y_i X_i' \boldsymbol{\beta} - \ln y_i! \quad (8.4)$$

To maximize it, the first order condition is

$$\sum_{i=1}^N (y_i - \exp(X_i' \boldsymbol{\beta})) X_i' = 0 \quad (8.5)$$

We can use Newton-Raphson or other optimization algorithms to find the solution for the first order condition.

## 8.2 QMLE Poisson

### 8.2.0.1 Model

Note that in the first-order condition above, if  $X'_i$  includes a constant, then  $(y_i - \exp(X'_i\beta))$  sums to zero. If  $E[y_i|X_i] = \exp(X'_i\beta)$ , then the summation on the left-hand side has expectation of zero. Hence the only specification needed to apply it is the conditional expectation of  $Y$  given  $X$ . Even if the data is not Poisson-distributed, the estimator is still consistent. Therefore, this estimator is called the quasi-ML (QML) Poisson estimator.

Although the QML Poisson estimator is consistent under relatively weak condition, the regular variance estimator for the coefficients are not valid anymore.

If a stronger assumption is made that the data follows a Poisson distribution, then the error term has a variance which equals to the mean of  $Y$ . The estimator  $\hat{\beta}_P$  (ML estimator) follows a normal distribution asymptotically with a variance matrix

$$\text{Var}_{ML}[\hat{\beta}_P] = \left( \sum_{i=1}^N \mu_i X_i X'_i \right)^{-1} \quad (8.6)$$

On the other hand, the variance matrix for the QML Poisson estimator  $\hat{\beta}_{QML}$  is

$$\text{Var}_{QML}[\hat{\beta}_{QML}] = \left( \sum_{i=1}^N \mu_i X_i X'_i \right)^{-1} \left( \sum_{i=1}^N \omega_i X_i X'_i \right) \left( \sum_{i=1}^N \mu_i X_i X'_i \right)^{-1} \quad (8.7)$$

where  $\omega_i = \text{Var}[y_i|X_i]$  is the conditional variance of  $y_i$ .

### 8.2.1 Implementation

Poisson estimation is implemented in almost every statistical package. However, some of them may not work if you have a continuous dependent variable.

Stata's implementation of the Poisson model: "poisson" and "xtpoisson" do take a continuous dependent variable. Standard errors reported without any "vce" specification assume the Poisson variance equals the mean and are likely too small under overdispersion. For a cross-sectional Poisson treated as QMLE, the robust (sandwich) standard errors – the "robust" option – are exactly the adjustment needed for a misspecified variance, so they are valid even when the equidispersion assumption fails; the conditional mean only needs to be correctly specified. In the panel fixed-effects case, serial correlation within a unit is a separate issue that requires *clustered* standard errors. Older versions of "xtpoisson" did not offer a cluster-robust option, which is why the user-written program `xtpqml` was written to supply the cluster-robust standard errors of Wooldridge (1999) after `xtpoisson`.

### 8.3 Fixed-effect Poisson model

We specify a panel data count model as:

$$E[y_{it}|\alpha_i, X_{it}] = \alpha_i \exp(X'_{it}\beta) = \exp(\gamma_i + X'_{it}\beta). \quad (8.8)$$

It turns out that for fixed-effect Poisson model, just like in linear regression case, there is NO incidental parameter problem. Poisson MLE is the same for conditional and unconditional likelihood. (See Cameron and Trivedi, 2005, Page 805).

Therefore, the coefficient estimates are the same for conditional or unconditional fixed-effect Poisson model. That is, in Stata, “xtpoisson, fe” will return the same results as “xi: poisson i.group”, is the same as “xtpqml, fe”. The only difference is that “xtpqml, fe” returns “correct” standard errors for QMLE Poisson. You can also calculate standard errors with the other commands, using clustered standard errors, or bootstrapped standard errors.

Conditional fixed effect negative binomial model is also possible to be consistent for some specifications. But Poisson fixed effect is more popular since it is QMLE.

### 8.4 Negative Binomial model

The Poisson model has a restrictive property that the conditional variance equals the conditional mean. Often we’ll see “overdispersion”, that is, the variance exceeds the mean, in real data. The negative binomial model is a generalization of poisson model which allows for overdispersion by introducing an unobserved heterogeneity term. In a negative binomial model,

$$E[Y|X, \tau] = \mu\tau = e^{X'\beta} \cdot \tau \quad (8.9)$$

This extra  $\tau$  term follows a *Gamma*( $\theta, \theta$ ) distribution, with  $E(\tau) = 1$  and  $Var(\tau) = 1/\theta$ .

Conditional on  $X$  and  $\tau$ , the distribution of  $Y$  is still poisson:

$$\Pr[Y = y|X, \tau] = \frac{e^{-\mu\tau}(\mu\tau)^y}{y!} \quad (8.10)$$

Conditional on  $X$  only, the distribution of  $Y$  is negative binomial:

$$\Pr[Y = y|X] = \frac{\theta^\theta \mu^y \Gamma(\theta + y)}{\Gamma(y + 1) \Gamma(\theta) (\mu + \theta)^{\theta + y}} \quad (8.11)$$

This distribution still has conditional mean of  $\mu$ , but variance is  $\mu(1 + (1/\theta)\mu)$ .

When  $\alpha = 1/\theta$  approaches zero, the negative binomial model converges to poisson model.

## 8.5 Models for truncated counts

In some situations, all we observe are non-zero counts, because of the way data was collected. For example, we only observe people's visits to a park, only if they visit. To model how many times they visit, based on this kind of data set, we need to use zero-truncated count models.

$$\Pr[Y = y|y > 0, X] = \frac{\Pr(y|X)}{\Pr(y > 0|X)} = \frac{\Pr(y|X)}{1 - \exp(-\mu)} \quad (8.12)$$

where the last equality holds for the Poisson case, since  $\Pr(y = 0|X) = e^{-\mu}$ ; for the negative binomial,  $\Pr(y = 0|X) = (\theta/(\theta + \mu))^\theta$ , so the truncation factor is  $1 - (\theta/(\theta + \mu))^\theta$  instead. Thus we are modeling the counts given that we have a positive outcome. Both zero-truncated poisson ("ztp" in Stata) and zero-truncated negative binomial ("ztnb" in Stata) are estimated by Maximum-likelihood method.

We need to pay extra attention to the "over-dispersion" problem when dealing with zero-truncated data. When data is not truncated, coefficients estimated by poisson model is consistent, even if data is "overdispersed". However, if we have zero-truncated sample, then "over-dispersion" results in inconsistent coefficient estimation (Grogger and Carson, 1991). Therefore, "over-dispersion" needs to be tested before using truncated poisson model or truncated negative binomial model.

## 8.6 The hurdle regression model

Although negative binomial model relaxes the assumption of equal mean and variance, some researchers prefer modeling the process of generating zeros differently from the process of generating other values. Suppose zero counts are generated by a binary process:

$$\Pr[y = 0|X] = \frac{\exp(X\gamma)}{1 + \exp(X\gamma)} = \pi \quad (8.13)$$

Positive counts are generated by a truncated count process.

In this model, zero is a "hurdle" to get past before reaching positive counts. The hurdle model is estimated by two separate equations. It is easy to estimate. However, the marginal effect is tricky to calculate, since an independent variable, if appearing in both equations, will have effect through both equations.

$$E(y|X) = [\pi \times 0] + (1 - \pi) \times E(y|y > 0, X) = (1 - \pi) \times E(y|y > 0, X) \quad (8.14)$$

In stata, there are commands such as "hprobit", which is a hurdle model with first equation logit, and second equation truncated poisson. It is the same as two separate estimations; namely "logit" first, then "ztp" on positive counts.

The difference between a hurdle model and a Heckman selection model is that in a Heckman model the two equations are jointly estimated and the outcome equation must account for selection bias. An exclusion restriction – a variable that affects selection but not the outcome – is strongly recommended for the selection equation. Strictly speaking, under joint normality the model is

identified by the nonlinearity of the inverse Mills ratio even without an excluded variable, but identification then rests entirely on functional form and is fragile, so an exclusion restriction is recommended in practice. On the other hand, a hurdle model assumes the two processes are separate; there is no selection.

## 8.7 Zero-inflated count models

A zero-inflated count model also models two different processes. One is a binary process: generating zero or non-zeros. The second one is a count model. Note there is a difference between zero-inflated and hurdle model: In hurdle model, the second process is a truncated count model (positive counts). In a zero-inflated model, a zero can come from either the binary process, or from the count process.

A second difference is that the estimation of zero-inflated model is a joint estimation of the two processes.

$$y_i \sim \begin{cases} 0 & \text{with probability } \psi \\ g(y_i|X_i) & \text{with probability } 1 - \psi \end{cases} \quad (8.15)$$

Then

$$P(Y_i = y_i | X_i, Z_i) = \begin{cases} \psi(\gamma' Z_i) + (1 - \psi(\gamma' Z_i))g(0|X_i) & \text{if } y_i = 0 \\ (1 - \psi(\gamma' Z_i))g(y_i|X_i) & \text{if } y_i > 0 \end{cases} \quad (8.16)$$

This says that every non-zero observation follows a count process  $g(y_i|X_i)$ . Every zero observation has two possible sources: from a binary process (with probability  $\psi(\gamma' Z_i)$ ) and from a count process (with probability  $1 - \psi(\gamma' Z_i)$  into the count process and then with probability of  $g(0|X_i)$  to be a zero). The count process can be a Poisson process (“zip” in Stata) or a Negative Binomial process (“zinb” in Stata).

## 8.8 An example of extra-zero count data model comparison

A companion simulation study (see the detailed write-up in this guide’s companion [blog\\_book](#), chapter “Which count data model to use”) compares plain Poisson, Negative Binomial, zero-inflated Poisson (ZIP, both with a correctly and an incorrectly specified zero-process), and hurdle models across sample sizes and zero-inflation levels. The main findings: with large samples, plain Poisson QMLE is consistent even with unmodeled excess zeros, as long as the conditional mean of the count process is correctly specified and independent of the zero-generating process – the excess zeros only inflate the standard errors (via overdispersion), not the point estimates, so robust/sandwich standard errors should be used rather than switching models. Zero-inflated Poisson with a correctly specified zero process performs best when that specification is available; hurdle and misspecified-zero-process ZIP perform similarly to plain Poisson; log-linear regression on  $\log(y + 1)$  is consistently the worst performer, since it targets a different (log-scale) conditional expectation than the models above.

## 8.9 Interpreting Coefficients and Marginal Effects in Nonlinear Models

This section is broader than count data: it covers how to read coefficients and compute marginal effects in nonlinear models generally, including the binary discrete-choice models of the **Discrete and Limited Dependent Variables** chapter as well as the count models above.

### 8.9.1 Binary Response Models

In a linear model

$$y_i = X_i' \beta + \epsilon, \quad (8.17)$$

$$E(y_i) = X_i' \beta \quad (8.18)$$

When  $y_i$  is binary (meaning it takes two values, 1 or 0), We can also do an OLS regression on it, but a more popular way is to use a non-linear model. The most popular choices for modeling binary response are logit model and probit model. Both of them use the same idea: use a link function to map the binary variable into a continuous variable which is a linear function of the predictors.

Suppose  $p_i$  is the probability that  $y_i = 1$ . Then

$$p_i = E(y_i) = g^{-1}(X_i' \beta), \quad (8.19)$$

where  $g^{-1}()$  is a function that maps from the linear predictor to  $y_i$ , which is often called index function or inverse link function.

This says that the expectation of  $y_i$  (equivalently,  $p_i$ ) is not a linear function of  $x_i$ 's, but by using the link function, we are still able to model a transformed  $p_i$  as a linear function of  $x_i$ 's:

$$g(p_i) = X_i' \beta. \quad (8.20)$$

$g^{-1}()$  can be a normal CDF then we have a probit model:

$$g^{-1}(z) = \Phi(z) = \int_{-\infty}^z \frac{1}{\sqrt{2\pi}} \exp[-(t^2/2)] dt, \quad (8.21)$$

or it can be a logit:

$$g^{-1}(z) = \Lambda(z) = e^z / (1 + e^z). \quad (8.22)$$

### 8.9.2 Marginal Effects

Often times we are interested in the marginal effects of the predictors. In a linear model, it's easy: the coefficient on  $x_1$  means when  $x_1$  increase by one unit, how much will  $y$  increase.

In a binary response model,

$$\frac{\partial p}{\partial x} = \frac{\partial g^{-1}(X'_i \beta)}{\partial x} \quad (8.23)$$

In the case of probit:

$$\frac{\partial p}{\partial x} = \frac{\partial \Phi(X'_i \beta)}{\partial x} = \phi(X'_i \beta) \beta, \quad (8.24)$$

where  $\phi()$  is the density function of normal distribution.

In the case of logit:

$$\frac{\partial p}{\partial x} = \frac{\partial \Lambda(X'_i \beta)}{\partial x} = \gamma(X'_i \beta) \beta, \quad (8.25)$$

where  $\gamma()$  is the density function for logistic distribution:

$$\gamma(X'_i \beta) = \Lambda(X'_i \beta)(1 - \Lambda(X'_i \beta)) = p(1 - p). \quad (8.26)$$

Therefore for logit, a quick approximation of the marginal effect is to plug in the sample proportion  $\bar{p}$  and coefficient estimate  $\hat{\beta}$ :

$$\frac{\hat{\partial p}}{\partial x} \approx \bar{p}(1 - \bar{p})\hat{\beta} \quad (8.27)$$

This evaluates the effect at the mean probability; it is a rough rule of thumb, not the average marginal effect  $\frac{1}{n} \sum_i \hat{p}_i(1 - \hat{p}_i)\hat{\beta}$ , which is what software reports.

The logit model can also be formulated as

$$\log\left(\frac{p}{1-p}\right) = \mathbf{X}'_i \beta \quad (8.28)$$

which says the logarithm of the odds (the ratio of the two probabilities) is equal to  $\mathbf{X}'_i \beta$ . Therefore,

$$p = \frac{\exp(X'_i \beta)}{1 + \exp(X'_i \beta)} = \Lambda(X'_i \beta) \quad (8.29)$$

Often times we have discrete valued predictors, such as gender, or other dummy variables. In that case, it makes more sense to ask what's the effect of gender being 1 vs. 0 (female vs. male). What this question is can be formulated as:

$$E(p|x_1 = 1) - E(p|x_1 = 0). \quad (8.30)$$

In empirical analysis, it's often calculated by calculating the predicted probability by setting  $x_1$  to 1 and by setting  $x_1$  to 0. Then calculate the difference.

### 8.9.3 Count Data Models

In a count data model such as Poisson regression model or Negative Binomial model, we are modeling log of the expected counts:

$$\log(E(y)) = \mathbf{X}_i' \boldsymbol{\beta} \quad (8.31)$$

Therefore the marginal effect of  $x_1$ , for example, on  $E(y)$  is

$$\frac{\partial \log E(y)}{\partial x} = \hat{\beta}, \quad (8.32)$$

which means  $\beta$  is the marginal effect of  $x$  on log of the expected counts. Suppose the expected counts at  $x$  is  $\mu_x$ , then

$$\frac{(\partial \mu_x)/\mu_x}{\partial x} = \hat{\beta}, \quad (8.33)$$

which says that  $\beta$  represents the marginal effect of percentage change in  $\mu_x$  with respect to  $x$ .

Another way to formulate this is to use the Incidence Rate Ratio Interpretation. Suppose the expected counts at  $x$  is  $\mu_x$ , and expected counts at  $x + 1$  is  $\mu_{x+1}$ .

$$\log(\mu_x) = x' \beta. \quad (8.34)$$

Therefore,

$$\mu_x = \exp(x' \beta), \quad (8.35)$$

and

$$\frac{\mu_{x+1}}{\mu_x} = \exp(\beta). \quad (8.36)$$

This says that the exponentiated  $\beta$  is the incidence rate ratio of the expected counts, if  $x$  increases by 1.

To calculate the marginal effect of  $x$  on  $y$  (or  $E(y)$ ), we need to calculate

$$\frac{\partial \mu_x}{\partial x} = \hat{\beta} \mu_x = \hat{\beta} \exp(X_i' \hat{\beta}), \quad (8.37)$$

if we evaluate  $\mu_x$  at predicted value.

### 8.9.4 How to calculate marginal effects in Stata

Stata's margins command is a powerful command to get the predicted value or marginal effects after an estimation command. Check Stata's manual for more details.

# 9 Panel Data Models

## 9.1 Background

Panel data are repeated observations on the same “unit” (we call that cross-sectional units) over time. Panel data models are used since it has advantages over cross-sectional models or time-series models. The major advantage it has over cross-sectional models is: **Controlling for individual heterogeneity**.

For example, Baltagi and Levin (1992) studies cigarette demand across 46 states for the years 1963-88. Consumption is modeled as a function of lagged consumption, price and income. These variables vary across states and time. However, there are other variables that may be time-invariant ( $z_i$ ) or individual-invariant ( $w_t$ ). Examples of  $z_i$  are religion and, in a panel of adults, completed education. For example, Utah, as a Mormon state, has very low cigarette demand due to religious reason. Generally we consider it does not change over time or change very little over time. Examples of  $w_t$  include cigarette commercials on national TV or radio. Panel data models are able to control for individual heterogeneity (or cross-sectional heterogeneity) while cross-sectional models are not. In many (or most) social science studies, there is “unobserved effects” which is embedded in the error term. In other words, we have “omitted variable” problem. Without controlling for it, the estimation results are generally biased and inconsistent.

Panel data give more informative data, more variability and less collinearity among variables. Often time series data suffers from multicollinearity, while panel data has more variability from its cross-sectional units.

The basic unobserved effects model can be written as:

$$y_{it} = \mathbf{x}_{it}\boldsymbol{\beta} + c_i + u_{it} \quad (9.1)$$

where  $i$  indexes “units” (can be people, firms, or households, etc.), and  $t$  indexes time periods.

Three kinds of regressor appear in this chapter, and it pays to name them once.  $\mathbf{x}_{it}$  varies across units *and* time — price, income, and in most applications the treatment itself.  $z_i$  varies across units only, so within a unit it is fixed over time; religion and completed education above are the examples.  $w_t$  varies across time only, so in any period it is common to every unit; the national TV campaign is one. Which group a variable falls into is what decides almost everything that follows, including which coefficients fixed effects can estimate and how OLS weights the data.

Instruments do not appear until the panel turns dynamic. When they do, they are written  $\mathbf{Z}$ , which is a different object from the time-invariant  $z_i$  here — see the [Endogeneity](#), [GMM](#) and [Dynamic Panel Data](#) chapters.

Comparing to regular cross-sectional models, there is an extra term  $c_i$ .  $c_i$  can be treated as a random effect or fixed effect. Traditionally it is distinguished by whether it is estimated as a random

variable or a parameter. Modern econometrics tends to distinguish by the correlation between  $c_i$  and  $\mathbf{x}_{it}$ . If there is no correlation between  $c_i$  and  $\mathbf{x}_{it}$ , then it's a random effect; otherwise, it's a fixed effect.

It is true that in some cases it is hard to justify that there is no correlation between  $c_i$  and  $\mathbf{x}_{it}$ , which is one reason that people in economics tend to use fixed-effect models. However, fixed-effect models have its own limitations. One of them is: It is hard to justify that ALL the cross-sectional units' characteristics other than those already in the model do not change over time.

## 9.2 Random Effect Methods

A random effect model puts  $c_i$  into the error term, then estimate by FGLS (feasible GLS). Random effect model is sometimes called "Error-components model" since the overall error term is divided into an individual level error term ( $c_i$ ) and individual-time level error term ( $u_{it}$ ).

Under a random effect model, the variance-covariance matrix of the error term becomes:

$$\Omega = \begin{bmatrix} \sigma_c^2 + \sigma_u^2 & \sigma_c^2 & \cdots & \sigma_c^2 \\ \sigma_c^2 & \sigma_c^2 + \sigma_u^2 & \cdots & \sigma_c^2 \\ \vdots & \vdots & \ddots & \vdots \\ \sigma_c^2 & \sigma_c^2 & \cdots & \sigma_c^2 + \sigma_u^2 \end{bmatrix} \quad (9.2)$$

The random effect estimator is:

$$\hat{\beta}_{RE} = \left( \sum_{i=1}^N \mathbf{X}_i' \hat{\Omega}^{-1} \mathbf{X}_i \right)^{-1} \left( \sum_{i=1}^N \mathbf{X}_i' \hat{\Omega}^{-1} \mathbf{y}_i \right) \quad (9.3)$$

It is a special case of GLS. Before calculating  $\hat{\beta}_{RE}$ ,  $\sigma_c^2$  and  $\sigma_u^2$  need to be estimated. This is generally done by using residuals from within and between regressions (the Swamy-Arora method, the default in most software), or from pooled OLS.

## 9.3 Fixed Effect Methods

Random effect methods assume that  $c_i$  be orthogonal to  $\mathbf{x}_{it}$ . In many applications, the whole point of using panel data is to allow  $c_i$  be correlated with  $\mathbf{x}_{it}$ . We need fixed-effect models in those cases.

For the same model as in the equation above, the fixed-effect methods try to eliminate  $c_i$  using fixed effects transformation, or "within transformation". The FE transformation is to "de-mean" each observation by subtracting the group mean. Basically subtracting the group-mean equation below from the unit-level equation above:

$$\bar{y}_i = \bar{\mathbf{x}}_i \beta + c_i + \bar{u}_i \quad (9.4)$$

where  $\bar{y}_i$ , etc., means group means. For example, if we have 50 states with state level data of cigarette consumption across years, then groups means states, and we have 50 group means which are state level average cigarette consumption across years.

The reason for “demeaning” is to remove  $c_i$  from final estimation. The “demeaned” regression equation is the final regression used in estimation:

$$(y_{it} - \bar{y}_i) = (\mathbf{x}_{it} - \bar{\mathbf{x}}_i)\beta + (u_{it} - \bar{u}_i). \quad (9.5)$$

The unit-specific effect  $c_i$  cancels because it is constant over  $t$  ( $c_i - \bar{c}_i = 0$ ). Writing the within-transformed variables as  $\ddot{y}_{it} = y_{it} - \bar{y}_i$  and  $\ddot{\mathbf{x}}_{it} = \mathbf{x}_{it} - \bar{\mathbf{x}}_i$ , OLS on the demeaned data gives the **fixed-effects (within) estimator** (here  $\mathbf{x}_{it}$  is a  $1 \times k$  row vector, so  $\ddot{\mathbf{x}}'_{it}\ddot{\mathbf{x}}_{it}$  is the  $k \times k$  outer product)

$$\hat{\beta}_{FE} = \left( \sum_i \sum_t \ddot{\mathbf{x}}'_{it}\ddot{\mathbf{x}}_{it} \right)^{-1} \sum_i \sum_t \ddot{\mathbf{x}}'_{it}\ddot{y}_{it}. \quad (9.6)$$

This is consistent for  $\beta$  even when  $c_i$  is correlated with  $\mathbf{x}_{it}$ , which is the key advantage of FE over random effects. The price is that any time-invariant regressor (the  $z_i$  above) is also swept out by the demeaning and cannot be estimated. The usual variance estimator is  $\hat{V}(\hat{\beta}_{FE}) = \hat{\sigma}_u^2 (\sum_i \sum_t \ddot{\mathbf{x}}'_{it}\ddot{\mathbf{x}}_{it})^{-1}$  with  $\hat{\sigma}_u^2$  using degrees of freedom  $NT - N - k$ ; in practice one clusters by unit to allow serial correlation within  $i$ .

## 9.4 Fixed effects or random effects: the Hausman test

Random effects is more efficient (it uses both within and between variation) but is consistent only if  $c_i$  is uncorrelated with  $\mathbf{x}_{it}$ ; fixed effects is consistent either way. The Hausman test compares the two estimators:

$$H = (\hat{\beta}_{FE} - \hat{\beta}_{RE})' [\hat{V}(\hat{\beta}_{FE}) - \hat{V}(\hat{\beta}_{RE})]^{-1} (\hat{\beta}_{FE} - \hat{\beta}_{RE}) \sim \chi_k^2 \quad (9.7)$$

under the null that random effects is consistent. A large  $H$  (small  $p$ -value) favors fixed effects. In R this is `phtest(fe, re)` from the `plm` package; in Stata, `hausman fe re`.

**Caveat: this classic form requires RE to be efficient.** The formula above relies on  $\hat{V}(\hat{\beta}_{FE}) - \hat{V}(\hat{\beta}_{RE})$  being positive semi-definite, which in turn requires the RE estimator to actually be the *efficient* estimator under the null – true only under the RE model’s homoscedasticity and no-serial-correlation assumptions. If you are using robust or cluster-robust covariance matrices (the usual practice when errors are heteroscedastic or serially correlated), RE is generally no longer efficient, the covariance difference can fail to be positive definite, and the classic Hausman test statistic can be negative or otherwise ill-behaved. In that case, use a regression-based Hausman test instead: augment the RE (or pooled) regression with the within-unit means of the time-varying regressors (Mundlak’s approach) or with the demeaned regressors themselves (Wooldridge’s regression-based test), and test the added coefficients jointly with a standard (robust) Wald test – this is valid under heteroscedasticity/clustering and reduces to the same idea as the classic Hausman test under its assumptions.

## 9.5 Collapse to unit means, or keep the long format?

A practical question comes up whenever we hold a panel. Should we average each unit over time and run one regression on the  $N$  resulting rows, or keep all  $NT$  rows and cluster the standard errors by unit? Many people treat this as a trade-off and assume the long format must carry more information. It usually does not. The answer turns entirely on where the regressors live.

Take the case where every regressor is time-invariant, so the model is  $y_{it} = \mathbf{z}_i\gamma + c_i + u_{it}$ . With a balanced panel, pooled OLS on the  $NT$  rows and OLS on the  $N$  unit means give the same  $\hat{\gamma}$ . Not close — identical. Random effects returns the same number as well. Wooldridge (2010) states the reason plainly: because  $c_i$  is there, new observations within a unit carry no information about  $\gamma$  beyond how they move the unit average  $\bar{y}_i$ . In effect we only have  $N$  useful pieces of information. Averaging discards nothing, because  $\mathbf{z}_i$  was already constant over  $t$  and the within variation was never available to  $\gamma$  in the first place.

What does differ is the degrees of freedom. The unit-mean regression uses  $N - k$ ; pooled OLS reports  $NT - k$ , and that is the mistake. Moulton (1990) is the classic illustration, and Donald and Lang (2007) show that  $t_{N-k}$  is the reference distribution that keeps the size right. Clustering by unit repairs the standard error when  $N$  is large, but it does not repair the degrees of freedom when  $N$  is small.

Now put a time-varying  $\mathbf{x}_{it}$  into the model. The two options stop agreeing, and not only for the coefficient on  $\mathbf{x}_{it}$  — every coefficient moves, the one on  $\mathbf{z}_i$  included. Averaging to unit means throws away all within variation, so what survives is the between estimator. Pooled OLS keeps both sources and forces a single coefficient on them, so it returns a variance-weighted average of the within and the between relationship. The two agree only when the within and between coefficients are equal, which is precisely the null the Hausman test examines. So this is not a new problem. It is the FE-versus-RE question in different clothes.

| Situation                                     | Collapsed vs long                                              |
|-----------------------------------------------|----------------------------------------------------------------|
| All regressors time-invariant, balanced panel | Same coefficients; only the degrees of freedom differ          |
| Any time-varying regressor                    | All coefficients differ, including those on $\mathbf{z}_i$     |
| Unbalanced panel                              | Differ regardless, since pooled OLS weights each unit by $T_i$ |

The reconciliation is Mundlak's device, the same one used above for the regression-based Hausman test. Add the unit means  $\bar{\mathbf{x}}_i$  to the long regression alongside  $\mathbf{x}_{it}$ , and the within and between variation carry separate coefficients:  $\mathbf{x}_{it}$  picks up the within estimate,  $\bar{\mathbf{x}}_i$  picks up the *difference* between the between and within estimates, so the two coefficients sum to what the collapsed regression reports. That second coefficient is the Hausman contrast, and its  $t$ -statistic is the Hausman test. We keep the within variation and recover the between answer in one regression.

The useful way to remember it: collapse freely when every regressor is constant within the group, and never when one is not.

## 9.6 What the long format buys

The section above is mostly a warning. The other side is worth seeing directly, so we simulate panels where we know the truth. We use `fixest` here; in Stata the same regressions are `regress`, `xtreg`, `be` and `xtreg, fe`.

### 9.6.1 The equivalence, and where it fails

Take 1000 units over 6 periods, one time-invariant regressor  $z_i$ , and a true  $\gamma = 2$ .

```
set.seed(42)
N <- 1000; T <- 6

z <- rnorm(N)
bal <- data.frame(id = rep(1:N, each = T), t = rep(1:T, N),
                  z = rep(z, each = T), c = rep(rnorm(N), each = T))
bal$y <- 2 * bal$z + bal$c + rnorm(N * T)
agg <- aggregate(cbind(y, z) ~ id, data = bal, FUN = mean)

print(c(long      = coef(feols(y ~ z, bal))["z"],
        collapsed = coef(feols(y ~ z, agg))["z"]), digits = 12)
```

```
      long      collapsed
1.99351677619 1.99351677619
```

The two agree to every digit printed. Now break the balance by dropping the later periods for units with high  $z_i$ , so  $T_i$  varies, and varies with the regressor.

```
unb <- bal[!(bal$z > 0.5 & bal$t > 2), ]
agg2 <- aggregate(cbind(y, z) ~ id, data = unb, FUN = mean)

print(c(long      = coef(feols(y ~ z, unb))["z"],
        collapsed = coef(feols(y ~ z, agg2))["z"]), digits = 12)
```

```
      long      collapsed
2.00682070731 2.01481052322
```

They no longer agree. Pooled OLS weights each unit by  $T_i$ ; the collapsed regression counts each unit once. Neither is wrong — they answer different questions, and we should know which one we asked. That  $T_i$  is a special case of a general rule, set out in the next section.

### 9.6.2 What collapsing costs

Now the case that matters. Let  $x_{it}$  vary over time and make  $c_i$  correlated with  $\bar{x}_i$ , which is the situation fixed effects exists for. The truth is  $\beta = 1$ .

```
mu <- rnorm(N)
d <- data.frame(id = rep(1:N, each = T),
                x = rep(mu, each = T) + rnorm(N * T),
                c = rep(mu + rnorm(N), each = T))
d$y <- 1 * d$x + d$c + rnorm(N * T)
d$xbar <- ave(d$x, d$id)
dagg <- aggregate(cbind(y, x) ~ id, data = d, FUN = mean)

between <- feols(y ~ x, data = dagg)
pooled <- feols(y ~ x, data = d, cluster = ~id)
within <- feols(y ~ x | id, data = d)
mundlak <- feols(y ~ x + xbar, data = d, cluster = ~id)

pull <- function(m, v) c(coef(m)[[v]], se(m)[[v]])
tab <- rbind(`between (collapsed)` = pull(between, "x"),
            `pooled OLS` = pull(pooled, "x"),
            `within / FE` = pull(within, "x"),
            `Mundlak: x` = pull(mundlak, "x"),
            `Mundlak: xbar` = pull(mundlak, "xbar"))
colnames(tab) <- c("estimate", "std. error")
round(tab, 3)
```

|                     | estimate | std. error |
|---------------------|----------|------------|
| between (collapsed) | 1.845    | 0.035      |
| pooled OLS          | 1.484    | 0.022      |
| within / FE         | 0.995    | 0.014      |
| Mundlak: x          | 0.995    | 0.014      |
| Mundlak: xbar       | 0.850    | 0.036      |

Read the table down. The collapsed regression returns about 1.85, badly biased, because it is the between estimator and  $\bar{x}_i$  carries  $c_i$  with it. Pooled OLS returns about 1.48 — closer, since some within variation is mixed in, but still inconsistent. The within estimator returns about 0.99, and only the long format can compute it: once we collapse there is nothing left to demean. Mundlak delivers both at once, exactly as described above — 0.99 on  $x_{it}$ , 0.85 on  $\bar{x}_i$ , and  $0.99 + 0.85$  recovers the between estimate.

So the argument for the long format is not precision. It is that demeaning, the Hausman comparison, unit-specific slopes, lagged dependent variables and anything else built on within variation stop existing once the panel is collapsed. Dynamic models are the extreme case, and they get their own chapter, [Dynamic Panel Data](#).

The rule behind that  $T_i$ , and behind the variance weighting in the between-versus-within comparison, is not specific to panels. It is set out in [Interpreting OLS: Effect Weights](#) and its companion, [Outcome Weights](#).



# 10 Survival Models

## 10.1 Survival function

Survival analysis is also called time to event analysis or event history analysis. We use survival models to analyze data that have the following characteristics:

1. the dependent variable is the waiting time till the occurrence of some event;
2. some observations are censored; that is, some units are not in the data set anymore for some reason before an event happens. For some units we observe the event and the exact waiting time; for others it has not occurred, and all we know is that the waiting time exceeds the observation time. However, we believe if they stayed in the study, sooner or later an event would happen.
3. we are trying to explain (predict) the occurrence of an event with some predictors (explanatory variables).

Let  $T$  denote the time variable, a nonnegative, continuous random variable with PDF  $f(t)$  and CDF  $F(t)$ , where  $t$  is a realization of  $T$ . The survival function is defined by

$$S(t) = Pr(T > t) = 1 - F(t) \quad (10.1)$$

The graph of  $S(t)$  against  $t$  is known as the survival curve. If there is no censoring, then the survival function is easy to estimate by the ratio of the number of units surviving beyond  $t$  to the total number of units in the sample. The most common method to estimate the survival function with a survival data set containing censored observations is the Kaplan-Meier method. The basic idea is to use the product of a series of conditional probabilities.

First sort the event times from the smallest to the largest:

$$t_{(1)} \leq t_{(2)} \leq \dots \leq t_{(r)} \quad (10.2)$$

Then the survival curve is estimated by:

$$\hat{S}(t) = \prod_{j|t_{(j)} \leq t} (1 - \frac{d_j}{r_j}), \quad (10.3)$$

where  $r_j$  is the number of units at risk at  $t_{(j)}$  and  $d_j$  is the number of “failure” or events at  $t_{(j)}$ . Note that units censored at  $t_{(j)}$  are included in  $r_j$ .

The variance of this estimator can be estimated by:

$$\widehat{var}(S(t)) = (\hat{S}(t))^2 \sum_{j|t_{(j)} \leq t} \frac{d_j}{r_j(r_j - d_j)}. \quad (10.4)$$

## 10.2 Hazard function

We are often interested in which periods have the highest or lowest change of “death” or “failure” or other events. That is, we may be interested in the probability that a state ends between  $t$  and  $t + \Delta t$  conditional on having reached  $t$  in the first place.

The probability is

$$\Pr(t < T \leq t + \Delta t | T \geq t) = \frac{F(t + \Delta t) - F(t)}{S(t)} \quad (10.5)$$

The hazard function is defined by

$$h(t) = \lim_{\Delta t \rightarrow 0} \frac{\Pr(t < T \leq t + \Delta t | T \geq t)}{\Delta t} = \frac{f(t)}{S(t)} = \frac{f(t)}{1 - F(t)} \quad (10.6)$$

One of the simplest functional forms is the exponential distribution

$$f(t, \theta) = \theta e^{-\theta t} \quad (10.7)$$

Therefore, the hazard function is

$$h(t) = \theta \quad (10.8)$$

A much more flexible functional form is Weibull, with CDF, survival, and density

$$F(t, \theta, \alpha) = 1 - \exp(-(\theta t)^\alpha), \quad S(t) = \exp(-(\theta t)^\alpha), \quad f(t) = \alpha \theta^\alpha t^{\alpha-1} \exp(-(\theta t)^\alpha). \quad (10.9)$$

Hazard function can be shown to be

$$h(t) = \alpha \theta^\alpha t^{\alpha-1} \quad (10.10)$$

However, Weibull distribution does not allow for the possibility that the hazard may first increase then decrease over time. Log-normal distribution does allow this.




Figure 10.1: Various hazard functions from Davidson and MacKinnon

### 10.3 Log-likelihood function

Suppose we have  $n$  units with a survival function  $S(t)$ , density  $f(t)$ , and hazard  $h(t)$ . If for unit  $i$ , the event happens at  $t_i$ , its contribution to the likelihood function is the density at that point:

$$L_i = f(t_i) = S(t_i)h(t_i) \quad (10.11)$$

If no event at  $t_i$ , all we know is that the lifetime exceeds  $t_i$ . The probability is

$$L_i = S(t_i), \quad (10.12)$$

which is the contribution to the likelihood function from a censored observation.

Therefore, if  $U$  denotes the set of uncensored observations, the log-likelihood function for the entire sample can be written as

$$\ell(t, \theta) = \sum_{i \in U} \log h(t_i | \mathbf{X}_i, \theta) + \sum_{i=1}^n \log S(t_i | \mathbf{X}_i, \theta) \quad (10.13)$$

## 10.4 Proportional Hazard Models

### 10.4.1 Model and likelihood

One class of models that is quite widely used is the proportional hazard models (based on Cox (1972)).

$$h(t|\mathbf{x}_i) = h_0(t) \exp(\mathbf{x}'_i \beta), \quad (10.14)$$

In this model  $h_0(t)$  is called baseline hazard rate; it describes the risk for units with  $\mathbf{x}_1 = \mathbf{0}$ .  $\exp(\mathbf{x}'_i \beta)$  represents the relative risk, a proportionate increase or decrease in risk, due to  $\mathbf{x}_i$ . The interpretation of  $\beta$  is that  $\exp(\beta_i)$  gives the relative risk change associated with an increase of one unit in  $x_i$ , all other explanatory variables remaining constant.

$\beta$  is estimated by maximizing a partial likelihood function. It is partial likelihood function since the baseline hazard is factored out. Suppose one event happens at time  $t_j$ . Conditional on this event the probability that case  $i$  dies (meaning this event happened to case  $i$ , not other cases in the risk set) is

$$L_i = \frac{h_0(t_j) \exp(x'_i \beta)}{\sum_l I(T_l \geq t_j) h_0(t_j) \exp(x'_l \beta)} = \frac{\exp(x'_i \beta)}{\sum_l I(T_l \geq t_j) \exp(x'_l \beta)} \quad (10.15)$$

We can see here that Cox's model is for continuous time survival data. It is assumed there is no ties at time  $t_j$  (in continuous time is not possible for two events to happen at the same time). In reality, we observe ties of event time, because in many situations, we observe grouped data. For example, we observe case  $m$  and  $n$  have events at the same time  $j$ . In that case, in the denominator, whether we include case  $m$  in the calculation of the likelihood of case  $n$  is ambiguous. In that case, special treatment is needed. A standard approximation (Breslow and Peto) is to let

$$L_{m \in D(t_j)} \simeq \frac{\prod_{m \in D(t_j)} \exp(x'_m \beta)}{[\sum_{l \in R(t_j)} \exp(x'_l \beta)]^{d_j}} \quad (10.16)$$

where  $D(t_j)$  is the set of cases that die at time  $t_j$  and  $d_j$  denotes the number of cases that die at time  $t_j$ . This approximation works well with small number of ties relative to total number of cases at risk.

One important feature (disadvantage?) for the proportional hazard model is that for all units, the effect is the same at all time  $t$ . To see this:

$$\frac{h_1(t)}{h_2(t)} = \frac{h_0(t) \exp(\mathbf{x}'_1 \beta)}{h_0(t) \exp(\mathbf{x}'_2 \beta)} = \frac{\exp(\mathbf{x}'_1 \beta)}{\exp(\mathbf{x}'_2 \beta)}, \quad (10.17)$$

which does not depend on  $t$ .

One advantage of Cox's model is that it is considered as a semi-parametric model since it does not have to specify the distribution of the survival time. The estimation process relies only on the order in which events occur, not the exact time they occur. Parameter estimates are derived assuming continuous survival times.

### 10.4.2 Interpretation

In the equation above, if we would like to calculate the partial effect of  $x_j$  on  $h$ , we have:

$$\partial h(t|\mathbf{x}_i)/\partial x_j = h_0(t) \exp(\mathbf{x}_i' \beta) \beta_j = \beta_j h(t|\mathbf{x}_i), \quad (10.18)$$


If we do it the other way, plug in  $x_j + 1$  will generate

$$h_0(t) \exp(\mathbf{x}_i' \beta + \beta_j) = \exp(\beta_j) h(t|\mathbf{x}_i) \quad (10.19)$$

Therefore, a one unit change in  $x_j$  multiplies the hazard by  $\exp(\beta_j)$  (a proportionate change of  $\exp(\beta_j) - 1$ ).

### 10.4.3 Example in R

```
library(ggfortify)
library(survival)
fit <- survfit(Surv(time, status) ~ sex, data = lung)
autoplot(fit)
```



Censored observations are denoted by crosses. This is the so-called Kaplan-Meier curve.

Now let's do a Cox model example.

```
fit2 <- coxph(Surv(time, status) ~ sex + age, data = lung, method="breslow")
summary(fit2)
```

Call:

```
coxph(formula = Surv(time, status) ~ sex + age, data = lung,
      method = "breslow")
```

n= 228, number of events= 165

|     | coef      | exp(coef) | se(coef) | z      | Pr(> z )   |
|-----|-----------|-----------|----------|--------|------------|
| sex | -0.512565 | 0.598957  | 0.167462 | -3.061 | 0.00221 ** |
| age | 0.017013  | 1.017158  | 0.009222 | 1.845  | 0.06506 .  |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

|     | exp(coef) | exp(-coef) | lower .95 | upper .95 |
|-----|-----------|------------|-----------|-----------|
| sex | 0.599     | 1.6696     | 0.4314    | 0.8316    |
| age | 1.017     | 0.9831     | 0.9989    | 1.0357    |

Concordance= 0.603 (se = 0.025 )

Likelihood ratio test= 14.08 on 2 df, p=9e-04

Wald test = 13.44 on 2 df, p=0.001

Score (logrank) test = 13.69 on 2 df, p=0.001

Examples of parametric survival models: exponential, Weibull, or log-logistic:

```
exp_fit <- survreg(Surv(time, status) ~ sex + age, data = lung, dist="exponential")
summary(exp_fit)
```

Call:

```
survreg(formula = Surv(time, status) ~ sex + age, data = lung,
      dist = "exponential")
```

|             | Value    | Std. Error | z     | p      |
|-------------|----------|------------|-------|--------|
| (Intercept) | 6.35967  | 0.63547    | 10.01 | <2e-16 |
| sex         | 0.48093  | 0.16709    | 2.88  | 0.004  |
| age         | -0.01562 | 0.00911    | -1.72 | 0.086  |

Scale fixed at 1

Exponential distribution

Loglik(model)= -1156.1 Loglik(intercept only)= -1162.3

Chisq= 12.48 on 2 degrees of freedom, p= 0.002

Number of Newton-Raphson Iterations: 4

n= 228

```
weibull <- survreg(Surv(time, status) ~ sex + age, data = lung, dist="weibull")
summary(weibull)
```

Call:

```
survreg(formula = Surv(time, status) ~ sex + age, data = lung,
        dist = "weibull")
```

|             | Value    | Std. Error | z     | p       |
|-------------|----------|------------|-------|---------|
| (Intercept) | 6.27485  | 0.48137    | 13.04 | < 2e-16 |
| sex         | 0.38209  | 0.12748    | 3.00  | 0.0027  |
| age         | -0.01226 | 0.00696    | -1.76 | 0.0781  |
| Log(scale)  | -0.28230 | 0.06188    | -4.56 | 5.1e-06 |

Scale= 0.754

Weibull distribution

Loglik(model)= -1147.1    Loglik(intercept only)= -1153.9

Chisq= 13.59 on 2 degrees of freedom, p= 0.0011

Number of Newton-Raphson Iterations: 5

n= 228

```
loglog <- survreg(Surv(time, status) ~ sex + age, data = lung, dist="loglogistic")
summary(loglog)
```

Call:

```
survreg(formula = Surv(time, status) ~ sex + age, data = lung,
        dist = "loglogistic")
```

|             | Value    | Std. Error | z     | p       |
|-------------|----------|------------|-------|---------|
| (Intercept) | 5.92232  | 0.53269    | 11.12 | < 2e-16 |
| sex         | 0.47751  | 0.14036    | 3.40  | 0.00067 |
| age         | -0.01401 | 0.00771    | -1.82 | 0.06946 |
| Log(scale)  | -0.56991 | 0.06543    | -8.71 | < 2e-16 |

Scale= 0.566

Log logistic distribution

Loglik(model)= -1152.9    Loglik(intercept only)= -1160.9

Chisq= 16.07 on 2 degrees of freedom, p= 0.00032

Number of Newton-Raphson Iterations: 4

n= 228

One caveat when reading these next to the Cox fit above. `survreg` estimates *accelerated failure time* models, parameterized on log survival time, whereas `coxph` estimates proportional hazards. The coefficients therefore carry *opposite* signs for the same underlying effect: a covariate that raises the hazard shortens expected survival time. A sign flip between the two tables is the expected result of the differing parameterizations, not a contradiction.

## 10.5 Discrete-time Survival Models

Cox (1972) is the *continuous-time* proportional hazards model; the discrete-time analogue models the odds of an event at time  $t_j$  given survival up to that point. This proportional-odds specification for the discrete hazard,

$$\frac{h(t_j|\mathbf{x}_i)}{1 - h(t_j|\mathbf{x}_i)} = \frac{h_0(t_j)}{1 - h_0(t_j)} \exp(\mathbf{x}_i' \beta), \quad (10.20)$$

is a discrete-time *proportional-odds* hazard model rather than a direct discretization of Cox's continuous-time proportional hazards model. (The complementary log-log link is the discrete-time form that corresponds more directly to a grouped continuous-time proportional hazards model.) If we take logs on both sides, we have

$$\text{logit}(h(t_j|\mathbf{x}_i)) = \alpha_j + \mathbf{x}_i' \beta, \quad (10.21)$$

This looks very similar to a logit model. In fact, we can fit a discrete-time proportional hazard model by running a logistic regression on a set of pseudo observations generated by “filling in” the observations between the start time to the event time or time at the end of the study period (censored). The outcome of this process is 0 unless there is an event, then it turns 1. It is treated as independent Bernoulli observations with probability given by hazard  $h_{ij}$  for individual  $i$  at time  $t_j$ . The likelihood function for the discrete-time survival model coincides with the likelihood of a sequence of independent Bernoulli trials (Singer and Willett, 1993, has more details).

### 10.5.1 Do it in Stata

To do a discrete-time survival model in Stata, your data set needs to be set up as by unit-time. For example, observations by company month. The event (or failure) variable needs to be set to zero until the event happens. If no event until the end of the study period, then it is censored. A data set by unit-time can be analyzed by logit model, optionally with clustered standard error. To set up a structure like this, user-written programs called *dthaz* and *prsnperd* (means person-period) can be used.

To do a continuous-time survival model in Stata, your data can be either set up as discrete-time case (in which case you can use a time-varying explanatory variable) or a single-record-per-person data set. Stata has built-in command *stset* to set up a survival analysis structure. Then you can run *stcox* for Cox's proportional hazard model.

# 11 Dynamic Panel Data

## 11.1 When Is It a Problem

Model setup:

$$y_{i,t} = \gamma y_{i,t-1} + X'_{i,t} \beta + C_i + \epsilon_{i,t} \quad (11.1)$$

Pooled regression is biased and inconsistent. The problem is:

$$Cov(y_{i,t-1}, (C_i + \epsilon_{i,t})) \approx \frac{\sigma_c^2}{(1 - \gamma)} \quad (11.2)$$

Random effects model is biased and inconsistent, for the same reason that  $y_{i,t-1}$  is necessarily correlated with  $C_i$ , which is part of the composite error term in the random effect model.

Fixed effect model setup:

$$y_{i,t} - \bar{y}_i = (X_{i,t} - \bar{X}_i)' \beta + \gamma(y_{i,t-1} - \overline{y_{i,-1}}) + (\epsilon_{i,t} - \bar{\epsilon}_i) \quad (11.3)$$

where  $\overline{y_{i,-1}}$  is the unit mean of the *lagged* dependent variable. This is not exactly  $\bar{y}_i$  (the mean of current  $y_{i,t}$ ): the two differ at the first and last usable periods, though they coincide asymptotically and under the usual balanced-sample convention.

Nickell (1981) shows that (see also Anderson and Hsiao 1981)

$$Cov((y_{i,t-1} - \overline{y_{i,-1}}), (\epsilon_{i,t} - \bar{\epsilon}_i)) \approx -\frac{\sigma_\epsilon^2}{T(1 - \gamma)^2} \left[ \frac{(T - 1) - T\gamma + \gamma^T}{T} \right] \quad (11.4)$$

which indicates that the demeaned lagged  $y$  is correlated with the demeaned error term. The correlation may be large if  $T$  is small, but when  $T$  is big, then the correlation goes down to near zero. When  $T$  is big, we don't have a problem with fixed effect model.

## 11.2 How Big Is the Bias

When  $T$  is small, which is typically the case in many microeconomic settings, then we say the fixed effect model with a lagged dependent variable is inconsistent. But how bad is it?

The limit of  $\hat{\gamma} - \gamma$  is approximately  $-\frac{(1+\gamma)}{T-1}$ . When  $T = 10$  and  $\gamma = .5$ , the bias is about  $-0.167$  which is  $1/3$  of the true value.

But a simulation study (Judson and Owen 1999) shows that the bias can be as big as 20% of the true coefficient even when  $T = 30$ .

### 11.3 Anderson and Hsiao estimator

Anderson and Hsiao (1981) suggested looking at the first difference estimator:

$$y_{i,t} - y_{i,t-1} = (X_{i,t} - X_{i,t-1})'\beta + \gamma(y_{i,t-1} - y_{i,t-2}) + (\epsilon_{i,t} - \epsilon_{i,t-1}) \quad (11.5)$$

or

$$\Delta y_{i,t} = \Delta X_{i,t}\beta + \gamma\Delta y_{i,t-1} + \Delta\epsilon_{i,t} \quad (11.6)$$

This does not solve the endogeneity problem, since  $\Delta y_{i,t-1}$  is still correlated with the error term. AH's idea is to instrument  $\Delta y_{i,t-1}$  with the past level  $y_{i,t-2}$ , or past difference  $y_{i,t-2} - y_{i,t-3}$ . This estimator is consistent, since neither of these instruments is correlated with  $\Delta\epsilon_{i,t}$ , assuming error is not auto-correlated.

### 11.4 Arellano-Bond estimator

Arellano and Bond (1991) expanded the idea by using additional lags of the dependent variable as instruments. For example, both  $y_{i,t-2}$  and  $y_{i,t-3}$  can be used as instruments. In fact, as  $t$  increases, the number of instruments available also increases. In period 3 only  $y_{i,1}$  is available. In period 4  $y_{i,1}$  and  $y_{i,2}$  are available. In period 5  $y_{i,1}$  and  $y_{i,2}$  and  $y_{i,3}$  are available, and so on. In other words, we'll have an instrument matrix with one row for each time period that we are instrumenting:

$$Z_i = \begin{bmatrix} y_{i,1} & 0 & 0 & 0 & 0 & 0 & \cdots & 0 & 0 & 0 \\ 0 & y_{i,1} & y_{i,2} & 0 & 0 & 0 & \cdots & 0 & 0 & 0 \\ 0 & 0 & 0 & y_{i,1} & y_{i,2} & y_{i,3} & \cdots & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \cdots & \vdots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & 0 & \cdots & 0 & y_{i,1} & \cdots & y_{i,T-2} \end{bmatrix} \quad (11.7)$$

This is the so called difference GMM estimator.

### 11.5 Blundell and Bond estimator

AB model has a problem: when the instruments are weak, the estimator is not good. That will happen when  $y$  follows random walk or near random walk. In that case the past levels won't be a good predictor of the future changes.

BB add a second set of moment conditions in *levels*. The basic idea is that lagged differences  $\Delta y_{i,t-1}$  are uncorrelated with the composite error  $C_i + \epsilon_{i,t}$  under an additional initial-condition/stationarity restriction, so they can instrument the level equation.

This is implemented by so-called *system GMM*, which stacks two equations and estimates them jointly:

- the **differenced equation** ( $\Delta y_{i,t}$ ), instrumented by lagged *levels*  $y_{i,t-2}, y_{i,t-3}, \dots$  (the Arellano–Bond moments); and
- the **level equation** ( $y_{i,t}$  in levels, keeping  $C_i$ ), instrumented by lagged *differences*  $\Delta y_{i,t-1}$ .

The level-equation moments are valid only under the extra assumption that the initial deviations from the long-run mean are uncorrelated with  $C_i$  (mean stationarity of the process). System GMM is more efficient than difference GMM when  $y$  is persistent (near a random walk), and it allows time-invariant regressors. A practical warning: the instrument count grows with  $T^2$ , so *instrument proliferation* can overfit the endogenous regressors and weaken the Hansen  $J$  test; lags should be collapsed or capped in practice.



## 12 Missing data

Data set like:

| y | x | z | w | q |
|---|---|---|---|---|
| 1 | 4 | 2 | ? | ? |
| 4 | ? | 1 | 2 | ? |
| 2 | ? | ? | ? | ? |
| ? | 2 | 1 | ? | ? |
| ? | ? | ? | ? | ? |

### 12.1 Different cases under different assumptions

1. Missing completely at Random (MCAR): the occurrence of missing data is not related to the missing value, the values of any other variables, or the pattern of missingness in other variables. Too good to be true situation.
2. Missing at Random (MAR): the occurrence of missing values for a variable is random, contingent on the value or missingness of observable variables. Or, the missingness can be modeled.
3. Missing Not at Random (MNAR): the occurrence of missing values is systematically related to the unobserved value itself, even after conditioning on observed data. This is the hardest case: there is no model-free fix, but it is not hopeless. It can be addressed with selection models, pattern-mixture models, or sensitivity analysis, all of which rely on untestable assumptions about the missingness mechanism.

The MAR case is the one we are interested in.

For MCAR: you lose efficiency if you simply drop the data. For MAR: if you don't model it, you suffer bias and efficiency.

If we model missingness, for MCAR, you gain efficiency. For MAR, you correct bias and gain efficiency.

### 12.2 Old methods

1. Listwise deletion
2. Mean imputation.

## 12 Missing data

Problems: \* understates variability in the imputed variable \* does not recover associations between variables.

So standard errors are in general under-estimated.

### 3. Regression-based imputation

Use a model such as  $x = \alpha_0 + \alpha_1 z + \alpha_2 y$ . But that introduce extra noise that are not accounted for (similar to generated variable problem).

### 4. Interpolation of panel data

That is, to use the observation from last period or a linear interpolation for the same unit.

## 12.3 Modern methods

- Account for uncertainty in imputed variable
- Use a model to predict missing observation.
- Instead of picking one, pick many
- uncertainty is represented by VCV matrix of the coefficients used to predict missing values.

### 12.3.1 Basic ideas

Imputation model:  $x = \alpha_0 + \alpha_1 z + \alpha_2 y$

Main model:  $y = \beta_0 + \beta_1 x + \beta_2 z$

1. Pick  $M$  values of  $\alpha$  out of the asymptotic distribution, the multivariate normal  $N(\hat{\alpha}, \hat{\Sigma})$ , using the estimate  $\hat{\alpha}$  and its VCV  $\hat{\Sigma}$  as the mean and VCV of the distribution.
2. For each of the  $M$  draws, predict the missing values *and add a random residual* drawn from the imputation model's estimated error distribution, creating  $M$  data sets. Both draws are needed: the draw of  $\alpha$  carries parameter uncertainty, and the residual draw carries the irreducible scatter around the prediction. Dropping the residual draw gives what is called "improper" imputation, and reintroduces precisely the understated-variability problem that sank mean imputation and single regression imputation above.
3. calculate  $M$  new estimates and combine them by averaging,  $\tilde{\beta} = \frac{1}{M} \sum_{m=1}^M \tilde{\beta}_m$ , using each of the imputed data sets. then calculate the variance of  $\tilde{\beta}$ :

$$V_{\beta} = W + (1 + 1/M)B$$

where  $W = \frac{1}{M} \sum_{m=1}^M s_m^2$ ,  $B = \frac{1}{M-1} \sum_{m=1}^M (\tilde{\beta}_m - \tilde{\beta})^2$ , in other words, the within-imputation and between-imputation variation. The standard error is  $\sqrt{V_{\beta}}$ . These formulas (Rubin's rules) are written here for a single scalar coefficient  $\beta$ ; for the full coefficient vector,  $W$  is the average of the within-imputation covariance matrices,  $B$  is the between-imputation covariance matrix  $\frac{1}{M-1} \sum_m (\tilde{\beta}_m - \tilde{\beta})(\tilde{\beta}_m - \tilde{\beta})'$ , and  $V_{\beta} = W + (1 + 1/M)B$  is the total covariance matrix.

### 12.3.2 MI through Chained Equations (MICE) (by Buuren)

1. Discard observations with all missing.
2. Fill in the missing data with random draws from the observed values.
3. Move through the columns and perform single-variable imputation using some method.
4. Replace the original replacements with the fitted replacements. Repeat 3 for a large number of times, or with a convergence criteria.
5. Do 1-4  $m$  times to create  $m$  imputed data sets.

Many ways to implement MICE.

- Regression (linear, logit, multinomial), get  $\hat{w}$  or  $f(\hat{w})$  sample.
- The default is Predictive Mean Matching (PMM)
  - a. create predicted value
  - b. pick a small set of donor cases that have the closest predicted values (the number of donors is software-specific; e.g., three or five).
  - c. randomly choose one of those donors' observed values to impute.

### 12.3.3 Bayesian Data Augmentation

- MICE is Markov Chain
- Build a missing data model into a Bayesian model, treating missing values as another parameter to estimate by drawing out of its posterior distribution.

$$f(\beta, y_{miss} | y_{obs}) \propto f(y_{obs} | \beta, y_{miss}) f(\beta, y_{miss})$$

- We integrate out the missing values by sampling from the total distribution, then averaging out the beta distributions over the space of missing data points.
- As long as data is MAR, the likelihood of missingness is not related to  $\beta$  (ignorability), this is fine.

### 12.3.4 FIML

If there are no missing data, the likelihood function is

$$L(\mu, \Sigma) = \prod_{i=1} f(y_i | \mu, \Sigma)$$

If there are missing value,

$$L(\mu, \Sigma) = \prod_{i=1} f(y_i | \mu_i, \Sigma_i)$$

### 12.3.5 mi in stata

```
log using mi_10_model2.log, replace

clear
set more off

* set matsize is obsolete since Stata 16 (matrix sizes are now dynamic)
* and can be safely omitted on modern Stata; kept here only because the
* original script targeted an older Stata version.
set matsize 4000

use "CROSSED_ContestUserActivity FEB 2014 temp.dta", clear
*sample 10

mi set flong
*local dummies "x1 x2"
local continuous "CulturalDist_KS5D_ctr tight_ctr targ_tight_ctr user_country_opennessValue_ctr"

mi register imputed `continuous'
mi impute chained (regress) `continuous' , add(10) rseed(1)

mi estimate, cmdok: heckprob has_won c.CulturalDist_KS5D_ctr##c.tight_ctr targ_tight_ctr gender
```

This is a sample code. In stata, you have to do “mi set”, then “mi register” to register variables you need to impute. Then “mi impute chained” if you have multiple variables to impute. Then the imputation will take all variables you registered (imputed or regular) in the prediction model unless you specify otherwise. Then do “mi estimate”.

# 13 Interpreting OLS: Effect Weights

An OLS regression returns one number,  $\hat{\beta}$ . What that number estimates when effects are heterogeneous, and how each observation contributes to it, depends on the geometry of the sample and on what else is in the model. This chapter derives the implicit weights OLS uses, and follows them across cross-sectional and panel settings.

Two companion notes work through related applications in causal inference: [Weights in OLS in causal inference](#) and [OLS cannot be ATE](#).

## 13.1 OLS as a weighted average

Start with a simple regression of  $y$  on  $x$  with an intercept:

$$y_i = \alpha + \beta x_i + e_i \quad (13.1)$$

The OLS slope is the familiar ratio of a covariance to a variance:

$$\hat{\beta} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} \quad (13.2)$$

Now give each observation its own slope through the sample centroid  $(\bar{x}, \bar{y})$ :

$$s_i \equiv \frac{y_i - \bar{y}}{x_i - \bar{x}} \quad (13.3)$$

Write  $d_i \equiv x_i - \bar{x}$  for the horizontal distance from the centroid, and  $w_i \equiv d_i^2$ . Multiply and divide each term of Equation 13.2 by  $d_i$  and it rearranges into


$$\hat{\beta} = \frac{\sum_{i=1}^n d_i^2 s_i}{\sum_{i=1}^n d_i^2} = \frac{\sum_{i=1}^n w_i s_i}{\sum_{i=1}^n w_i} \quad (13.4)$$

So  $\hat{\beta}$  is an exact weighted average of the deviation slopes  $s_i$ , and the weight each observation carries is its squared distance from the mean of  $x$ . Nothing is assumed here; Equation 13.4 is algebra.

### 13.1.1 What the weights look like


The figures below draw the weights  $w_i$  of Equation 13.4 as the thickness of each ray from the centroid. They grow with the square of horizontal distance.

```
show <- function(x, y, main = "") {
  xb <- mean(x); yb <- mean(y); d <- x - xb; s <- (y - yb) / d; w <- d^2
  plot(x, y, pch = 16, cex = 1.5, xlim = c(-0.5, 11.5), ylim = c(0, 9),
       xlab = "x", ylab = "y", main = main)
  abline(v = xb, lty = 3, col = "grey55")
  segments(xb, yb, x, y, lwd = 0.6 + 9 * w / sum(w), col = "#9E431899")
  abline(lm(y ~ x), lwd = 2, col = "#0E6473")
  points(x, y, pch = 16, cex = 1.5)
  points(xb, yb, pch = 21, bg = "white", col = "#0E6473", lwd = 2.5, cex = 1.7)
  text(x, y, labels = sprintf("%.1f%%", 100 * w / sum(w)),
       pos = ifelse(y >= yb, 3, 1), cex = .85, col = "#9E4318")
}
show(c(1, 3), c(2, 8), "Two Points")
```




With two points the weights are equal, 50% each, and both  $s_i$  are just the line through the pair.

```
show(c(1, 3, 4), c(2, 8, 5), "Three Points")
```



Add a third point at  $x = 4$  and  $\bar{x}$  moves to 2.67. The point at  $x = 3$  now sits almost on top of  $\bar{x}$ , so it takes only 2.4% of the weight even though its deviation slope is the largest of the three ( $s_2 = 9.0$ ).

```
show(c(1, 3, 4, 9), c(2, 8, 5, 6), "Four Points")
```




Add an extreme point at  $x = 9$  and  $\bar{x}$  moves to 4.25. That one observation now carries 65% of the total weight, while  $x = 4$  falls to 0.18% because it has landed near the new centroid. The weights are relative to the whole sample: add data anywhere and the weights everywhere change.

**Simulation — a 100-point sample.** The concentration does not wash out in larger samples. We draw  $n = 100$  points,  $x_i \sim N(5, 2^2)$  and  $y_i = 2x_i + \varepsilon_i$  with  $\varepsilon_i \sim N(0, 3^2)$ , and compare the weight held by the ten points nearest  $\bar{x}$  against the ten furthest.

```
set.seed(3)
n <- 100
x <- rnorm(n, 5, 2)
y <- 2 * x + rnorm(n, 0, 3)

xb <- mean(x); yb <- mean(y); w <- (x - xb)^2
o <- order(w, decreasing = TRUE)

plot(x, y, pch = 16, cex = 1, col = "grey40", xlab = "x", ylab = "y", main = "N = 100 Points")
abline(v = xb, lty = 3, col = "grey55")
segments(xb, yb, x, y, lwd = 0.3 + 10 * w / max(w), col = "#9E431870")
abline(lm(y ~ x), lwd = 2.5, col = "#0E6473")
points(xb, yb, pch = 21, bg = "white", col = "#0E6473", lwd = 2.5, cex = 1.8)
```



```
vtab(c(`Weight 10 Nearest Mean (%)` = 100 * sum(w[o[91:100]]) / sum(w),
      `Weight 10 Furthest (%)`      = 100 * sum(w[o[1:10]]) / sum(w),
      `OLS Estimate`                = coef(lm(y ~ x))["x"]), 3)
```

|                            | value  |
|----------------------------|--------|
| Weight 10 Nearest Mean (%) | 0.138  |
| Weight 10 Furthest (%)     | 35.360 |
| OLS Estimate               | 1.865  |

The 10% of points furthest from  $\bar{x}$  hold 35% of the weight; the 10% closest hold 0.14%. A point sitting exactly at  $\bar{x}$  has an undefined  $s_i$  and zero weight: the wild value arrives paired with a weight that annihilates it.

## 13.2 Deviation slopes $s_i$ versus true effects $b_i$

Now let each unit have its own structural slope  $b_i$ :

$$y_i = \alpha + b_i x_i + e_i \quad (13.5)$$

$b_i$  is a parameter we never see.  $s_i$  from Equation 13.3 is computed from the data. The two are not the same object, and the gap between them is the point of this section.

### 13.2.1 Decomposing $s_i$

Let  $\bar{b} = n^{-1} \sum_i b_i$  and  $C = n^{-1} \sum_i (b_i - \bar{b})(x_i - \bar{x})$ , the sample covariance between  $b$  and  $x$ . Substituting  $x_i = d_i + \bar{x}$  into Equation 13.5,

$$s_i = b_i + \frac{\bar{x}(b_i - \bar{b}) - C}{d_i} + \frac{e_i - \bar{e}}{d_i} \quad (13.6)$$

The middle term is the trouble. When effects are heterogeneous ( $b_i \neq \bar{b}$ ) it does not vanish, and dividing it by  $d_i$  makes  $s_i$  blow up for points near  $\bar{x}$  — before any noise enters. Taking expectations over  $e_i$ ,

$$E[s_i | x, b] = b_i + \frac{\bar{x}(b_i - \bar{b}) - C}{d_i} \neq b_i \quad (13.7)$$

So  $s_i$  is biased for  $b_i$ , observation by observation.

### 13.2.2 Aggregating up to $\hat{\beta}_{OLS}$

Take the  $d_i^2$ -weighted expectation of  $s_i$ , combining Equation 13.4 and Equation 13.7:

$$E[\hat{\beta}_{OLS} | x, b] = \frac{\sum_i d_i^2 E[s_i | x, b]}{\sum_i d_i^2} = \frac{\sum_i d_i^2 b_i + \bar{x} \sum_i d_i (b_i - \bar{b}) - C \sum_i d_i}{\sum_i d_i^2} \quad (13.8)$$

Because  $\sum_i d_i = 0$  and  $\sum_i d_i (b_i - \bar{b}) = nC$ , the numerator collapses:

$$\sum_i d_i^2 E[s_i | x, b] = \sum_i d_i^2 b_i + n\bar{x}C = \sum_i b_i d_i (d_i + \bar{x}) = \sum_i b_i d_i x_i \quad (13.9)$$

And since  $\sum_i d_i^2 = \sum_i d_i x_i$ , we get the exact expectation:

$$E[\hat{\beta}_{OLS} | x, b] = \frac{\sum_{i=1}^n d_i x_i b_i}{\sum_{i=1}^n d_i x_i} = \frac{\sum_{i=1}^n (x_i - \bar{x}) x_i b_i}{\sum_{i=1}^n (x_i - \bar{x})^2} \quad (13.10)$$

The individual biases  $E[s_i] - b_i$  of Equation 13.7 do not cancel. They reshape the weighting instead: from  $d_i^2$  on the noisy  $s_i$  to  $d_i x_i$  on the true slopes  $b_i$ . Each  $s_i$  is wrong in a specific, permanent way, yet combined they are exactly right, because the wrongness moves into the weights rather than vanishing.

### 13.2.3 So what?

Equation 13.10 has three consequences for practice.

1. **OLS is not the ATE under heterogeneity.**  $\hat{\beta}_{OLS}$  is not the plain average effect  $\bar{b} = E[b_i]$ . It is a variance-weighted parameter, and it equals the ATE only when the  $b_i$  are uncorrelated with the leverage  $d_i x_i = (x_i - \bar{x})x_i$ .
2. **Leverage annihilates the individual bias.** Equation 13.7 shows  $s_i$  is a hopeless estimate of  $b_i$  precisely where  $d_i \approx 0$ . The weight  $w_i = d_i^2$  zeroes out exactly those unstable points, turning biased observation-level slopes into the exact parameter of Equation 13.10.
3. **The same mechanism runs through the causal estimators.** With a binary regressor it puts the most weight where treatment is near a coin flip (Angrist 1998); grouped by treatment status it reverses the propensity weights, favouring the smaller group (Słoczyński 2022). Both are Equation 13.10 read at different granularities, as the rest of the chapter shows.

**Simulation — deviation slopes versus true effects.** We check the algebra directly, drawing  $n = 2000$  with heterogeneous slopes  $b_i \sim N(1, 0.5^2)$ , regressor  $x_i \sim N(0, 2^2)$ , and  $y_i = b_i x_i + \varepsilon_i$ ,  $\varepsilon_i \sim N(0, 2^2)$ .

```
set.seed(6)
n <- 2000
b <- rnorm(n, 1, 0.5)
x <- rnorm(n, 0, 2)
y <- b * x + rnorm(n, 0, 2)
xb <- mean(x); d <- x - xb; w <- d^2
s <- (y - mean(y)) / d

mtab(cbind(value = c(
  `cor(s_i, b_i), Pearson`      = cor(s, b),
  `cor(s_i, b_i), Spearman`     = cor(s, b, method = "spearman"),
  `sd(s_i)`                     = sd(s),
  `Weighted Avg of s_i`        = sum(w * s) / sum(w),
  `OLS Estimate`               = unname(coef(lm(y ~ x))[2]),
  `E[beta | x, b] (eq-exact)`   = sum(d * x * b) / sum(d * x))), 3)
```

|                           | value   |
|---------------------------|---------|
| cor(s_i, b_i), Pearson    | 0.019   |
| cor(s_i, b_i), Spearman   | 0.258   |
| sd(s_i)                   | 117.420 |
| Weighted Avg of s_i       | 0.950   |
| OLS Estimate              | 0.950   |
| E[beta   x, b] (eq-exact) | 0.991   |

The  $s_i$  are wildly dispersed ( $\text{sd}(s_i) > 100$ ) and barely correlated with the true  $b_i$  — Pearson 0.02, though the rank correlation is 0.26, higher once the explosive tails near  $\bar{x}$  are set aside. Their leverage-weighted average equals  $\hat{\beta}_{OLS}$  exactly; that is Equation 13.4, an identity. The same leverage

applied to the true  $b_i$  gives 0.99, the value  $\hat{\beta}$  estimates *in expectation* (Equation 13.10) — and this sample's  $\hat{\beta} = 0.95$  sits about two standard errors below it, a reminder that Equation 13.10 governs the expectation, not any one draw. The signal is buried, not absent.

Noise is not what distorts  $s_i$ . Turn it off entirely, set  $y_i = b_i x_i$ , and ask how well each  $s_i$  tracks its own  $b_i$ :

```
y0 <- b * x
s0 <- (y0 - mean(y0)) / (x - xb)
vtab(c(`cor(s_i, b_i), no noise` = cor(s0, b),
      `Weighted Avg of s_i`      = sum(w * s0) / sum(w),
      `Weighted Avg of b_i`      = sum(w * b) / sum(w),
      `OLS, noiseless`           = unname(coef(lm(y0 ~ x))[2])), 4)
```

|                         | value  |
|-------------------------|--------|
| cor(s_i, b_i), no noise | 0.1436 |
| Weighted Avg of s_i     | 0.9914 |
| Weighted Avg of b_i     | 0.9919 |
| OLS, noiseless          | 0.9914 |

The correlation stays at 0.14 with no noise in the data at all. A ray's slope reflects where every point sits, not just its own, so when the  $b_i$  genuinely differ one point's ray cannot recover that point's own slope. Noise makes it worse; it is not the cause.

### 13.2.4 Under homogeneity, the weights buy precision

If  $b_i = b$  for all  $i$ , Equation 13.6 reduces to  $s_i = b + (e_i - \bar{e})/d_i$ . With homoskedastic errors  $\text{Var}(e_i) = \sigma^2$ ,

$$E[s_i | X] = b, \quad \text{Var}(s_i | X) = \frac{\sigma^2(1 - 1/n)}{d_i^2} \propto \frac{1}{d_i^2} \quad (13.11)$$

Now  $w_i = d_i^2$  is exactly the inverse-variance weight. Any weighting of the  $s_i$  — equal,  $d_i^2$ , anything — is unbiased for  $b$ ; what the weights buy is precision, not correctness, and  $d_i^2$  is the choice that minimises variance. That is OLS as BLUE. The argument is conditional on  $X$ , where  $d_i$  is a known constant and the weights are functions of  $x$  alone.

## 13.3 The general rule: FWL

The panel chapter met two weightings — by  $T_i$  when a panel is unbalanced, and by variance when a regression mixes within and between variation. Neither is a panel quirk. Both follow from one rule.

Regress  $y$  on a regressor of interest  $V$  and a control matrix  $X$ , and let  $\tilde{V}$  be the residual from projecting  $V$  on  $X$ . By the Frisch–Waugh–Lovell theorem,

$$\hat{\beta} = \frac{\sum_{i=1}^n \tilde{V}_i y_i}{\sum_{i=1}^n \tilde{V}_i^2} \quad (13.12)$$

Be careful what  $\tilde{V}_i$  is. A single observation has no variation of its own; variation belongs to a collection. What it has is a *deviation* — how far its regressor sits from what the rest of the model predicts for it. With only an intercept, that is distance from the grand mean. With unit fixed effects, it is distance from the unit's own mean. Which mean gets subtracted is the most consequential detail in this section, and we come back to it.

Equation 13.12 reads two ways. This chapter takes the first; the second opens the companion chapter [Interpreting OLS: Outcome Weights](#).

**Weights on effects.** Under a heterogeneous model  $y_i = \beta_i V_i + X_i \gamma + e_i$ ,

$$\hat{\beta} = \frac{\sum_{i=1}^n w_i \beta_i}{\sum_{i=1}^n w_i}, \quad w_i = \tilde{V}_i V_i \quad (13.13)$$

When  $X$  holds saturated indicators or fixed effects,  $\sum_i \tilde{V}_i V_i = \sum_i \tilde{V}_i^2 \geq 0$ . This is the reading that answers *whose* effect  $\hat{\beta}$  estimates. In Equation 13.4 the weight  $d_i^2$  sat on  $s_i$ , because  $s_i$  was what we were averaging; here  $w_i = \tilde{V}_i V_i$  sits on  $\beta_i$ . The two differ observation by observation and coincide only when  $x$  is centred, though they always sum to the same total, since  $\sum_i d_i x_i = \sum_i d_i^2$ .

The same identity also weights the observed outcomes directly,  $\hat{\beta} = \sum_i \omega_i y_i$  with  $\omega_i = \tilde{V}_i / \sum_j \tilde{V}_j^2$  — a signed combination that says which rows drive the number and whether any enters with the wrong sign. That is a different object with its own uses, and it is the subject of the next chapter.

### 13.3.1 Where the square comes from

The square in  $w_i$  is derived, not assumed. Substitute  $y_i = \beta_i V_i + X_i \gamma + e_i$  into the numerator of Equation 13.12 and it becomes  $\sum_i \beta_i \tilde{V}_i V_i$ ; over the whole sample,

$$\sum_{i=1}^n \tilde{V}_i V_i = \sum_{i=1}^n \tilde{V}_i^2 \quad (13.14)$$

because  $\tilde{V}$  is orthogonal to the fitted part of  $V$ .

In a panel, several rows share a unit, and it is natural to ask whether the same cancellation holds *within a single unit* — summed over that unit's own  $t$ 's — and not only over the whole sample. Write  $V_{it}$  for the observation and  $\bar{V}_i$  for the unit's own mean. Demeaning gives the within-unit version by direct expansion,

$$\sum_{t=1}^T (V_{it} - \bar{V}_i) V_{it} = \sum_{t=1}^T V_{it}^2 - T \bar{V}_i^2 = \sum_{t=1}^T (V_{it} - \bar{V}_i)^2 \quad (13.15)$$

while a coarser control matches only in total, not unit by unit. Hazlett and Shinkre (2024) make this within-unit cancellation the centre of their treatment.

### 13.3.2 Special cases of the effect weights

In general  $\sum_t \tilde{V}_{it}^2$  within a unit does not factor into anything tidier —  $\tilde{V}_{it}$  carries both subscripts even when  $V$  does not. What changes from model to model is only whether it happens to simplify:

| Model                                | Regressor $V$           | Within-unit weight $w_i$          |
|--------------------------------------|-------------------------|-----------------------------------|
| Panel, group-level regressor         | $z_i$ (unit-level)      | $T_i \tilde{z}_i^2$               |
| Panel, within-unit variation         | $x_{it}$ (time-varying) | $(T-1) \widehat{\text{Var}}_i(x)$ |
| Binary treatment, saturated controls | $D_i \in \{0, 1\}$      | $n_x \hat{D}(X)(1 - \hat{D}(X))$  |

Add a time-varying  $x_{it}$  alongside  $z_i$  in the first row and the shortcut is gone: residualising  $z_i$  on something that moves within the unit makes  $\tilde{z}$  move within the unit too. The rule survives, the shortcut does not, and the  $T_i$  form is the one-regressor special case. The last row is Angrist (1998), whose strata weights  $\hat{D}(x)(1 - \hat{D}(x))$  peak where treatment is closest to a coin flip. Same rule, different regressor.

## 13.4 Two causal decompositions: Angrist and Słoczyński

Both results are usually met as free-standing facts about treatment effects. They are the rule above, read at two granularities. Equation 13.13 says the coefficient is a weighted average of effects with weights  $w_i = \tilde{V}_i V_i$ ; everything below only evaluates those weights. Angrist reads them at the stratum level. Słoczyński reads the same weights one level up, grouped by treatment status, where they come out reversed relative to the natural shares of the ATE. The two subsections are the same calculation, twice.

Let  $D \in \{0, 1\}$  be a binary treatment and  $X$  discrete strata  $x \in \{1, \dots, K\}$ .

### 13.4.1 Angrist (1998): weighting by stratum

Residualising  $D_i$  on a full set of stratum indicators gives  $\tilde{D}_i = D_i - \hat{D}(x)$ , where  $\hat{D}(x) = E[D \mid X = x]$ . Sum the effect weights  $w_i = \tilde{D}_i D_i$  within a stratum:

$$w_x = \sum_{i \in x} \tilde{D}_i D_i = \sum_{i \in x} D_i^2 - \hat{D}(x) \sum_{i \in x} D_i = n_x \hat{D}(x) - n_x \hat{D}(x)^2 = n_x \hat{D}(x)(1 - \hat{D}(x)) \quad (13.16)$$

Nothing new has been added: it is  $w_i = \tilde{V}_i V_i$  from Equation 13.13 with a binary  $V$ , where  $D_i^2 = D_i$  makes the sum collapse. The maximum at  $\hat{D}(x) = 1/2$  is just the fact that a binary variable has most variance when it is a coin flip.

Note what plays the role of  $s_i$  here. A single observation with a binary regressor has no slope; you need one treated and one control unit to form a difference. The smallest unit carrying an effect is therefore the *stratum*, and  $s_i$  becomes the stratum difference in means. The regressor sets the granularity, we do not choose it.

Let  $\hat{\delta}_x$  be the stratum difference-in-means. Substituting Equation 13.16 into Equation 13.13,

$$\hat{\beta}_{OLS} = \frac{\sum_{x=1}^K n_x \hat{D}(x)(1 - \hat{D}(x)) \hat{\delta}_x}{\sum_{x=1}^K n_x \hat{D}(x)(1 - \hat{D}(x))} \quad (13.17)$$

so strata with propensity near  $\hat{D}(x) = 0.5$  get the most weight per observation.

**Simulation — Angrist and Słoczyński weights.** Five covariate strata with stratum effects  $(0, 1, 2, 3, 4)$  and  $y = \tau_x D + X + \varepsilon$ ; the treatment probability varies by stratum, run once balanced ( $\rho \approx 0.5$ ) and once unbalanced ( $\rho \approx 0.09$ ).

```
set.seed(21)
n <- 400000
X <- sample(1:5, n, TRUE)

compare <- function(px) {
  d <- rbinom(n, 1, px[X]); ta <- c(0, 1, 2, 3, 4)[X]
  yy <- ta * d + X + rnorm(n)
  dt <- resid(lm(d ~ factor(X)))
  nx <- as.vector(table(X)); dh <- as.vector(tapply(d, X, mean))
  DIM <- as.vector(tapply(yy[d == 1], X[d == 1], mean) - tapply(yy[d == 0], X[d == 0], mean))
  rho <- mean(d); ATT <- weighted.mean(ta, d); ATU <- weighted.mean(ta, 1 - d)
  c(`rho` = rho,
    `Max |Weight Diff|` = max(abs(as.vector(tapply(dt * d, X, sum)) - nx*dh*(1-dh))),
    `OLS Saturated` = unname(coef(lm(yy ~ d + factor(X)))[["d"]]),
    `Angrist Weighted` = sum(nx*dh*(1-dh) * DIM) / sum(nx*dh*(1-dh)),
    `Sloczynski Approx` = (1 - rho) * ATT + rho * ATU,
    `ATE` = mean(ta))
}

res <- cbind(`Balanced (rho ~ 0.5)` = compare(c(.08, .20, .50, .75, .93)),
             `Unbalanced (rho ~ 0.09)` = compare(c(.02, .04, .08, .12, .20)))
mtab(res, 4)
```

|                   | Balanced (rho ~ 0.5) | Unbalanced (rho ~ 0.09) |
|-------------------|----------------------|-------------------------|
| rho               | 0.4913               | 0.0909                  |
| Max  Weight Diff  | 0.0000               | 0.0000                  |
| OLS Saturated     | 2.0121               | 2.8740                  |
| Angrist Weighted  | 2.0121               | 2.8740                  |
| Sloczynski Approx | 2.0293               | 2.8577                  |
| ATE               | 1.9978               | 1.9978                  |

The two weight vectors — the stratum sums  $\sum \tilde{D}_i D_i$  and Angrist's  $n_x \hat{D}(x)(1 - \hat{D}(x))$  — are the same, below  $10^{-9}$  (the Max |Weight Diff| row, 0 to four decimals), and the Angrist-weighted average reproduces the OLS coefficient exactly. In the balanced design the coefficient is 2.012 against an ATE of 1.998: close, because balance is where the weighting does least. Push the treated share to

0.09 and the coefficient jumps to 2.874 against the same 1.998. Nothing about the effects changed. Only the treated share did.

### 13.4.2 Słoczyński (2022): aggregating to two groups

Angrist sorts by covariate stratum. Słoczyński sorts by treatment status instead, collapsing everything into two numbers — the effect on the treated and the effect on the untreated. Aggregating the stratum effects  $\hat{\delta}_x$  of Equation 13.17 with each stratum's count of treated or untreated units gives

$$\text{ATT} = \frac{\sum_x n_x \hat{D}(x) \hat{\delta}_x}{\sum_x n_x \hat{D}(x)}, \quad \text{ATU} = \frac{\sum_x n_x (1 - \hat{D}(x)) \hat{\delta}_x}{\sum_x n_x (1 - \hat{D}(x))} \quad (13.18)$$

and the ATE combines them with their natural shares,  $\rho = P(D = 1)$ :

$$\text{ATE} = \rho \text{ATT} + (1 - \rho) \text{ATU} \quad (13.19)$$

The OLS weights of Equation 13.17 go as the propensity *variance*  $\hat{D}(x)(1 - \hat{D}(x))$ ; the group weights of Equation 13.18 go as the propensity itself or its complement. The relation between the two is Słoczyński's result. Write  $p = \hat{D}(X)$  for the cell propensity. Because the design is saturated — the controls are a full set of stratum dummies — residualising  $D$  on the strata leaves exactly  $\tilde{D} = D - p$ , so the coefficient is  $E[\tilde{D}Y]/E[p(1 - p)]$ . Since  $p$  is constant within a stratum,  $Y$  enters only through the group means  $m_j(p) = E[Y \mid p, D = j]$ , and, using  $\tilde{D} = 1 - p$  for treated and  $-p$  for untreated units,

$$\beta_{OLS} = \frac{\rho E[(1 - p) m_1(p) \mid D = 1] - (1 - \rho) E[p m_0(p) \mid D = 0]}{E[p(1 - p)]} \quad (13.20)$$

Substitute the within-group linear projections  $m_j(p) = \alpha_j + \gamma_j p + \varepsilon_j(p)$ . The residual term  $E[p(1 - p)(\varepsilon_1 - \varepsilon_0)]/E[p(1 - p)]$  vanishes when the group means are linear in  $p$ , leaving

$$\beta_{OLS} = (\alpha_1 - \alpha_0) + (\gamma_1 - \gamma_0) \frac{E[p^2(1 - p)]}{E[p(1 - p)]} \quad (13.21)$$

The last ratio is moment algebra — an identity in the first three moments of  $p$ , exact for any distribution. With  $\mu_j = E[p \mid D = j]$ ,  $V_j = \text{Var}(p \mid D = j)$ , and  $w_1, w_0$  the weights defined in Equation 13.23 below,

$$\frac{E[p^2(1 - p)]}{E[p(1 - p)]} = w_1 \mu_1 + w_0 \mu_0 \quad (13.22)$$

holds identically in the distribution of  $p$ . Under linearity  $\text{ATT} = (\alpha_1 - \alpha_0) + (\gamma_1 - \gamma_0)\mu_1$  and  $\text{ATU} = (\alpha_1 - \alpha_0) + (\gamma_1 - \gamma_0)\mu_0$ , so Equation 13.20 through Equation 13.22 combine into the main result,

$$\hat{\beta}_{OLS} = w_1 \text{ATT} + w_0 \text{ATU}, \quad w_1 = \frac{(1 - \rho)V_0}{\rho V_1 + (1 - \rho)V_0}, \quad w_0 = 1 - w_1 \quad (13.23)$$

where  $V_j$  is the variance of the propensity score within group  $j$ . This is exact for the effects of the within-group linear projections; for ATT and ATU it is exact when the group means are linear in  $p$  and an approximation otherwise.

Two things about Equation 13.23 matter. First, the weight  $w_1$  on ATT falls as the treated share  $\rho$  rises: the larger the treated group, the smaller its say. That is the whole of “smaller groups get more weight.” Second, when the propensity has equal variance in the two groups,  $V_1 = V_0$ , the weights collapse to the reversed shares,

$$\hat{\beta}_{OLS} \approx (1 - \rho) \text{ATT} + \rho \text{ATU} \quad (13.24)$$

the same two quantities as the ATE in Equation 13.19, with the weights **swapped**. At  $\rho = 0.1$  the coefficient puts 0.9 on the effect for the treated tenth; the ATE puts 0.1 there.

One caveat. Equation 13.24 is exact under equal propensity variance across groups — not “no covariates,” though no covariates is one trivial way to get there. Słoczyński calls the condition unlikely to hold exactly and offers  $\rho$  as a rule of thumb, cleanest near  $\rho = 0\%$ ,  $50\%$ , or  $100\%$ . In the simulation above the raw- $\rho$  form gives 2.029 against an exact 2.012, and 2.858 against an exact 2.874 — within about a percent, but the exact weights are what to report.

### 13.4.3 The three, side by side

| Decomposition     | Sorted by             | Smallest effect unit                | Exact when                | Reference      |
|-------------------|-----------------------|-------------------------------------|---------------------------|----------------|
| Observation rule  | nothing               | $s_i$                               | always (identity)         | Equation 13.4  |
| Angrist (1998)    | covariate stratum $x$ | stratum difference $\hat{\delta}_x$ | saturated controls        | Equation 13.17 |
| Słoczyński (2022) | treatment status $D$  | group averages<br>ATT, ATU          | equal propensity variance | Equation 13.24 |

Read down the last column. The observation identity always holds and says little. Angrist’s needs saturation and gives an interpretable weight. Słoczyński’s needs more and gives the sharpest statement of the three: the coefficient and the ATE are the same two numbers with the weights reversed. Underneath all of them, Hazlett and Shinkre (2024) show these weights are exactly the bias from fitting a constant effect when effects vary, and argue for estimators that tolerate heterogeneity rather than for diagnosing the weights after the fact.

## 13.5 Within versus between, in panels

Pooled OLS mixes two sources of variation: within-unit, the deviations from a unit's own mean, and between-unit, the differences across unit means. Which one dominates is not chosen; it follows from the same rule. Take the within side first.

Nothing new happens here — it is the opening leverage rule with a different mean subtracted. Equation 13.12 weights every observation by its squared deviation  $\tilde{V}_i^2$  from whatever the controls predict for it; Equation 13.13 weights each effect by  $w_i = \tilde{V}_i V_i$ . Simple OLS residualises against an intercept alone, so the deviation is from the grand mean and  $w_i = (x_i - \bar{x})^2$ , the leverage of Equation 13.4. Panel fixed effects residualises against unit dummies, so the deviation is from the unit's own mean and  $w_i = \sum_t (x_{it} - \bar{x}_i)^2$ , the within-unit variance in the special-cases table. Pooled OLS drops the dummies and is back to distance from the grand mean — which is exactly why the between comparison creeps back in. So the whole section is Equation 13.4 read with the unit mean in place of the grand mean.

### 13.5.1 Within-unit weighting: an extreme case

Half the units barely move ( $x$  has sd 0.2 around their own mean) and have a true slope of zero. The other half move a lot (sd 3) and have a true slope of two. The average unit's slope is one.

```
set.seed(1)
half <- N / 2
g <- data.frame(id = rep(1:N, each = T),
               move = rep(c(rep(0.2, half), rep(3, half)), each = T),
               b = rep(c(rep(0, half), rep(2, half)), each = T))
g$x <- rnorm(N * T, 0, g$move)
g$y <- g$b * g$x + rep(rnorm(N), each = T) + rnorm(N * T)

wg <- tapply(resid(feols(x ~ 1 | id, g))^2, g$id, sum)
mtab(cbind(value = c(
  `average of the two true slopes` = mean(c(0, 2)),
  `FE estimate` = coef(feols(y ~ x | id, g))["x"],
  `weight held by the steady half (%)` = 100 * sum(wg[1:half]) / sum(wg))), 3)
```

|                                    | value |
|------------------------------------|-------|
| average of the two true slopes     | 1.000 |
| FE estimate                        | 1.984 |
| weight held by the steady half (%) | 0.473 |

The fixed-effects estimate is 1.98, because the five hundred steady units hold under half of one percent of the weight. Their six observations sit almost on top of their own unit mean — with unit effects that is the relevant deviation — and squaring puts them nowhere in the sum. Nothing has gone wrong: the zero-slope units were never really in the regression.

### 13.5.2 Splitting the two channels

That result was about the *within-unit* deviation. Plain pooled OLS uses distance from the grand mean instead, and the two coincided above only because both halves were built around zero. Separate them and the algebra says exactly what happens.

Let  $G$  index groups — any number of them, any sizes — with shares  $p_g$ , group means  $\mu_g = E[x | G = g]$  and  $\bar{y}_g = E[y | G = g]$ , within-group variances  $\sigma_g^2$ , and true slopes  $\beta_g$ . The pooled slope is  $\text{plim } \hat{\beta}_{OLS} = \text{Cov}(x, y) / \text{Var}(x)$ , and the whole derivation is one move: condition on  $G$  and split numerator and denominator each into a within piece and a between piece. The laws of total covariance and total variance do this exactly,

$$\text{Cov}(x, y) = \underbrace{E[\text{Cov}(x, y | G)]}_{\text{within}} + \underbrace{\text{Cov}(E[x | G], E[y | G])}_{\text{between}} \quad (13.25)$$

$$\text{Var}(x) = \underbrace{E[\text{Var}(x | G)]}_{\text{within}} + \underbrace{\text{Var}(E[x | G])}_{\text{between}} \quad (13.26)$$

Name each variance, and define each slope as its own covariance over its own variance:

$$V_{\text{within}} = E[\text{Var}(x | G)], \quad \beta_{\text{within}} = \frac{E[\text{Cov}(x, y | G)]}{E[\text{Var}(x | G)]} \quad (13.27)$$

$$V_{\text{between}} = \text{Var}(E[x | G]), \quad \beta_{\text{between}} = \frac{\text{Cov}(E[x | G], E[y | G])}{\text{Var}(E[x | G])} \quad (13.28)$$

With those definitions the within numerator of Equation 13.25 is  $V_{\text{within}}\beta_{\text{within}}$  and the between numerator is  $V_{\text{between}}\beta_{\text{between}}$  — no algebra, just the definitions. Dividing Equation 13.25 by Equation 13.26 therefore gives the pooled slope as a convex combination,

$$\text{plim } \hat{\beta}_{OLS} = \frac{V_{\text{within}}\beta_{\text{within}} + V_{\text{between}}\beta_{\text{between}}}{V_{\text{within}} + V_{\text{between}}} \quad (13.29)$$

Nothing here used how many groups there are or that they are balanced; Equation 13.29 holds for any grouping — including  $G =$  the panel unit, with  $N$  groups of sizes  $T_i$ . The only modelling input is  $\text{Cov}(x, y | G = g) = \beta_g \sigma_g^2$ , that within a group the error is uncorrelated with  $x$ . It makes  $\beta_{\text{within}}$  an explicit share- and variance-weighted average of the group slopes, and leaves  $\beta_{\text{between}}$  as the across-group regression of the group  $y$ -means on the group  $x$ -means:

$$\beta_{\text{within}} = \frac{\sum_g p_g \beta_g \sigma_g^2}{\sum_g p_g \sigma_g^2}, \quad \beta_{\text{between}} = \frac{\sum_g p_g (\mu_g - \bar{\mu})(\bar{y}_g - \bar{\mu}_y)}{\sum_g p_g (\mu_g - \bar{\mu})^2} \quad (13.30)$$

with grand means  $\bar{\mu} = \sum_g p_g \mu_g$  and  $\bar{\mu}_y = \sum_g p_g \bar{y}_g$ .

The two-equal-group case in the simulation below is where these collapse to closed forms. Equal shares turn the averages into  $\frac{1}{2}$ 's. The two group means then form a fair two-point variable, and

such a variable taking values  $a_1, a_2$  has variance  $\frac{1}{4}(a_1 - a_2)^2$  — that is where the  $\frac{1}{4}$  comes from — so

$$V_{within} = \frac{1}{2}(\sigma_1^2 + \sigma_2^2), \quad \beta_{within} = \frac{\beta_1\sigma_1^2 + \beta_2\sigma_2^2}{\sigma_1^2 + \sigma_2^2}, \quad V_{between} = \frac{1}{4}\Delta\mu^2, \quad \beta_{between} = \frac{\Delta\bar{y}}{\Delta\mu} \quad (13.31)$$

where  $\Delta\mu = \mu_1 - \mu_2$  and  $\Delta\bar{y} = \bar{y}_1 - \bar{y}_2$ : with only two groups  $\beta_{between}$  is a single rise over run between the centroids.

$\beta_{within}$  is a variance-weighted average of the group slopes — the same  $\sigma^2$  weighting as everywhere else in this chapter, one level up. And  $\beta_{between}$  depends only on where the group centroids sit, so it can take any value at all, regardless of the within-group slopes.

The closed forms are special to two equal groups. With two *unequal* groups every  $\frac{1}{2}$  becomes a share  $p_g$  and the  $\frac{1}{4}$  becomes  $p(1 - p)$  — Angrist’s weight again; with *more than two* groups  $\beta_{between}$  is the genuine across-group regression slope of Equation 13.30, not a two-point ratio.

Fixed effects removes  $V_{between}$  by demeaning and isolates  $\beta_{within}$ . As the groups pull apart,  $\Delta\mu \rightarrow \infty$ ,  $V_{between}$  swamps  $V_{within}$  and drags  $\hat{\beta}_{OLS}$  toward  $\beta_{between}$ .

**Simulation — pulling the groups apart.** The same two half-samples as above, now shifting the steady group’s mean  $x$  by 0, 10, and 100, so only the centroid gap  $\Delta\mu$  changes while the within slopes stay fixed.

```
b_within <- (0 * 0.2^2 + 2 * 3^2) / (0.2^2 + 3^2)
V_within <- 0.5 * (0.2^2 + 3^2)
predicted <- function(shift) V_within * b_within / (V_within + 0.25 * shift^2)

park <- function(shift) {
  set.seed(1)
  half <- N / 2
  move <- c(rep(0.2, half), rep(3, half))
  b <- c(rep(0, half), rep(2, half))
  x <- rnorm(N * T, 0, rep(move, each = T)) + rep(c(rep(shift, half), rep(0, half)), each = T)
  y <- rep(b, each = T) * x + rep(rnorm(N), each = T) + rnorm(N * T)
  d <- data.frame(id = rep(1:N, each = T), x = x, y = y)
  c(`FE estimate` = coef(feols(y ~ x | id, d))["x"],
    `OLS Predicted` = predicted(shift),
    `OLS Actual` = coef(lm(y ~ x, d))["x"])
}

mtab(t(sapply(c(`shift 0` = 0, `shift 10` = 10, `shift 100` = 100), park)), 3)
```

|           | FE estimate | OLS Predicted | OLS Actual |
|-----------|-------------|---------------|------------|
| shift 0   | 1.984       | 1.991         | 1.983      |
| shift 10  | 1.984       | 0.305         | 0.311      |
| shift 100 | 1.984       | 0.004         | 0.004      |

Fixed effects stays flat near the theoretical  $\beta_{within} = 1.991$  wherever the steady group sits. Plain OLS falls from 1.983 to 0.004 on identical true slopes: as  $\Delta\mu$  grows, the between weight  $\frac{1}{4}\Delta\mu^2$  swamps the fixed within weight 4.52, and the coefficient is dragged toward  $\beta_{between} = 0$  like  $9/(4.52 + \Delta\mu^2/4)$ , the numerator  $\frac{1}{2}(\beta_1\sigma_1^2 + \beta_2\sigma_2^2) = 9$ .

```
set.seed(5)
m <- 200
draw <- function(shift, main) {
  xw <- rnorm(m, 0, 3);      yw <- 2 * xw + rnorm(m)
  xs <- rnorm(m, shift, 0.2); ys <- rnorm(m)
  x <- c(xw, xs); y <- c(yw, ys)
  plot(x, y, pch = 16, cex = .5, col = c(rep("#9E431855", m), rep("#0E6473", m)),
       xlab = "x", ylab = "y", main = main)
  abline(v = mean(x), lty = 3, col = "grey50")
  segments(mean(xw), mean(yw), mean(xs), mean(ys), lty = 2, lwd = 2, col = "grey30")
  points(c(mean(xw), mean(xs)), c(mean(yw), mean(ys)), pch = 21, bg = "white", cex = 1.4)
  abline(lm(y ~ x), lwd = 2, col = "#0E6473")
  dmu <- mean(xs) - mean(xw)
  Vb <- 0.25 * dmu^2
  Vw <- 0.5 * (var(xs) + var(xw))
  c(`Within Slope 1`      = unname(coef(lm(yw ~ xw))[2]),
    `Within Slope 2`      = unname(coef(lm(ys ~ xs))[2]),
    `Between Slope`       = unname((mean(ys) - mean(yw)) / dmu),
    `Between Weight (%)`  = 100 * Vb / (Vb + Vw),
    `OLS Estimate`        = unname(coef(lm(y ~ x))[2]))
}
op <- par(mfrow = c(1, 2), mar = c(4, 4, 2.5, 1))
a <- draw(0, "Centroids Coincide (Delta mu = 0)")
b <- draw(10, "Separated Centroids (Delta mu = 10)")
```




Figure 13.1: Pooled OLS as a within/between blend (Equation 13.29). Points are the two groups, open circles their centroids. The dashed line is the between-slope  $\beta_{\text{between}} = \Delta\bar{y}/\Delta\mu$  through the centroids; the solid line is the pooled fit  $\hat{\beta}_{OLS}$ . The within slopes are identical across panels; widening the centroid gap  $\Delta\mu$  (left to right) inflates  $V_{\text{between}}$  and rotates  $\hat{\beta}_{OLS}$  away from  $\beta_{\text{within}}$  toward  $\beta_{\text{between}}$ .

```
par(op)
mtab(cbind(`Delta mu = 0` = a, `Delta mu = 10` = b), 3)
```

|                    | Delta mu = 0 | Delta mu = 10 |
|--------------------|--------------|---------------|
| Within Slope 1     | 2.019        | 2.009         |
| Within Slope 2     | 0.007        | -0.288        |
| Between Slope      | 0.882        | -0.007        |
| Between Weight (%) | 0.026        | 85.113        |
| OLS Estimate       | 2.010        | 0.291         |

Pooled OLS as a within/between blend (Equation 13.29). Points are the two groups, open circles their centroids. The dashed line is the between-slope  $\beta_{\text{between}} = \Delta\bar{y}/\Delta\mu$  through the centroids; the solid line is the pooled fit  $\hat{\beta}_{OLS}$ . The within slopes are identical across panels; widening the centroid gap  $\Delta\mu$  (left to right) inflates  $V_{\text{between}}$  and rotates  $\hat{\beta}_{OLS}$  away from  $\beta_{\text{within}}$  toward  $\beta_{\text{between}}$ .

Read the two columns. The slopes inside each group barely change — about 2 for the wobbly cloud, indistinguishable from zero for the steady strip (the  $-0.288$  is noisy, since a regression on an  $x$  that hardly varies has a standard error near 0.35). What changes is the weight on the between comparison: 0.03% when the centres coincide, 85% when they are ten apart. So the fit lands on 2.01 in the first panel and 0.29 in the second.

The between slope reported on the left, 0.882, is meaningless — and harmlessly so: with the centres nearly on top of each other  $\Delta\mu \approx 0$ , an unstable ratio is multiplied by a vanishing weight and contributes nothing.

The strip has influence on the right despite barely moving. Influence is distance from the mean of  $x$ , and the strip sits ten units from it; what it contributes is not its slope but its position — it anchors one end of the dashed line.

A group dummy removes that channel. Demeaning subtracts each cloud's *own* mean, sliding the shifted strip back to the centre, so the right panel collapses onto the left — centroids coinciding. Run OLS on that and the solid fit swings back up to the within slope. That is all fixed effects is, and why its slope matches the left panel's regardless of how far the strip had been pushed.

So “the weight is how much  $x$  moves” is the demeaned case only. It is exact once unit effects are removed, and false in general, because with the halves at different levels of  $x$  part of the fitted line is a comparison *between* them rather than a weighting of within-unit slopes.

### 13.5.3 When the weighting changes nothing

Weighting only distorts when the effect is related to the weight. Give every unit its own slope, let units differ in how much  $x$  moves, and compare two designs: in (A) the slopes are unrelated to that movement, in (B) the slope is larger where  $x$  moves more. Both average near one. Design (A) is the  $s_i$ -versus- $b_i$  story from before, run as a panel.

```
set.seed(7)
sdX <- runif(N, 0.2, 3)

panel <- function(b) {
  h <- data.frame(id = rep(1:N, each = T),
                  x = rnorm(N * T, 0, rep(sdX, each = T)))
  h$y <- rep(b, each = T) * h$x + rep(rnorm(N), each = T) + rnorm(N * T)
  h
}

compareAB <- function(h, b) {
  w <- tapply(resid(feols(x ~ 1 | id, h))^2, h$id, sum)
  c(`FE estimate`          = coef(feols(y ~ x | id, h))["x"],
    `true slopes, plain avg` = mean(b),
    `true slopes, weighted` = sum(w * b) / sum(w))
}

b1 <- rnorm(N, 1, 1)
b2 <- 1 + 0.8 * (sdX - 1.6)
h1 <- panel(b1); h2 <- panel(b2)

mtab(cbind(`(A) unrelated` = compareAB(h1, b1),
           `(B) aligned`   = compareAB(h2, b2)), 4)
```

|                        | (A) unrelated | (B) aligned |
|------------------------|---------------|-------------|
| FE estimate            | 0.9813        | 1.5338      |
| true slopes, plain avg | 0.9716        | 1.0140      |

|                       | (A) unrelated | (B) aligned |
|-----------------------|---------------|-------------|
| true slopes, weighted | 0.9816        | 1.5422      |

In (A) all three numbers agree to within sampling noise: the weighting is there and costs nothing. In (B) the estimate is 1.53 while the plain average of the true slopes is 1.01, and the weighted average, 1.54, is what the estimator was tracking all along. None of this is a defect — a unit whose  $x$  barely moves genuinely says little about how  $y$  responds, and leaning on the units that do move is efficient. Design (B) is the continuous version of what Słoczyński describes for a binary treatment.

## 13.6 Summary

1. **Leverage weighting.** OLS weights each observation by regressor variance  $\tilde{V}_i^2 = (x_i - \bar{x})^2$  (Equation 13.4); units that vary more dominate the slope.
2. **What is being estimated.** Under heterogeneous slopes  $b_i$ , OLS estimates the variance-weighted parameter  $\sum_i w_i b_i / \sum_i w_i$  (Equation 13.10); it equals the ATE only when the  $b_i$  are uncorrelated with the weights.
3. **Between-group bias in pooled OLS.** Without fixed effects, group-mean differences  $\Delta\mu$  enter through  $V_{between} \beta_{between}$  (Equation 13.29) and distort the within-unit parameter; demeaning removes them.
4. **When it matters.** The weighting always happens. It changes the answer only when effects are heterogeneous and the heterogeneity lines up with the weight. With homogeneous effects it does no harm; with heterogeneous effects the question is not whether OLS is weighting but whether we know what it weighted.

Everything here is the *effect* reading — weights on the unobservable  $\beta_i$ , answering what  $\hat{\beta}$  estimates. The same FWL identity also weights the *observed* outcomes, a signed and fully computable combination that answers which rows drive the number and reaches, unchanged, into estimators far beyond OLS. That is the companion chapter, [Interpreting OLS: Outcome Weights](#).

# 14 Interpreting OLS: Outcome Weights

The companion chapter [Interpreting OLS: Effect Weights](#) read the OLS coefficient as a weighting of unobservable effects — it answered *whose* effect  $\hat{\beta}$  estimates. The same Frisch–Waugh–Lovell identity reads a second way, as a weighting of the *observed* outcomes. This reading answers a different question — which rows drive  $\hat{\beta}$ , and in which direction — and, unlike the first, it is computable and generalises past OLS. This chapter follows it.

## 14.1 The other reading of FWL

Residualise the regressor of interest  $V$  on the controls  $X$ , write  $\tilde{V}_i$  for the residual, and recall the residual-on-residual form of the coefficient, Equation [13.12](#):

$$\hat{\beta} = \frac{\sum_{i=1}^n \tilde{V}_i y_i}{\sum_{i=1}^n \tilde{V}_i^2}. \quad (14.1)$$

The previous chapter grouped each  $\tilde{V}_i$  with  $V_i$  and read  $\hat{\beta}$  as a weighting  $w_i = \tilde{V}_i V_i$  of effects (Equation [13.13](#)). Group  $\tilde{V}_i$  with  $y_i$  instead — pairing each residual with its own outcome in the numerator — and the *same* expression is a linear combination of the observed outcomes:

$$\hat{\beta} = \sum_{i=1}^n \omega_i y_i, \quad \omega_i = \frac{\tilde{V}_i}{\sum_{j=1}^n \tilde{V}_j^2}. \quad (14.2)$$

Equation [14.2](#) is the same algebra rearranged, but a different object from the effect weights.

|              | Weights on outcomes $\omega_i$                  | Weights on effects $w_i$                  |
|--------------|-------------------------------------------------|-------------------------------------------|
| Weights what | the observed $y_i$                              | that observation's own effect $\beta_i$   |
| Sign         | signed, sums to zero                            | non-negative (under saturation)           |
| Computable   | yes, $y$ is observed                            | no, the $\beta_i$ are never seen          |
| Holds when   | always, it is algebra                           | in expectation; per-unit needs saturation |
| Answers      | which rows drive $\hat{\beta}$ , any wrong sign | is this the effect I meant, whose effect  |

Two constraints pin  $\omega_i$  down, both from Equation [14.2](#). An intercept among the controls makes  $\sum_i \tilde{V}_i = 0$ , so the weights **sum to zero**. And  $\sum_i \tilde{V}_i V_i = \sum_i \tilde{V}_i^2$  makes the weighted regressor **sum to one**,  $\sum_i \omega_i V_i = 1$ .

Use this reading when overlap is in doubt and a row may enter with the wrong sign — which the non-negative  $w_i$  cannot show. Keep the effect reading for questions about the estimand.

## 14.2 What the outcome weights look like

Take the simple regression of  $y$  on  $x$  with an intercept, so  $\tilde{V}_i = x_i - \bar{x} = d_i$  and

$$\omega_i = \frac{d_i}{\sum_j d_j^2}. \quad (14.3)$$

Unlike the effect weights  $w_i = d_i^2$  — non-negative, growing with the *square* of distance — this is **linear in  $d_i$ , and so signed**: a point left of  $\bar{x}$  enters  $\hat{\beta}$  with a negative coefficient on its outcome, a point right of it positive. It is not a slope and does not involve  $\bar{y}$ ; it is simply a signed weight on the observed  $y_i$ .

**Simulation — simple regression.** To see it, we simulate  $n = 100$  points — regressor  $x_i \sim N(5, 2^2)$  and outcome  $y_i = 2x_i + \varepsilon_i$  with  $\varepsilon_i \sim N(0, 3^2)$  — then size each point by  $|\omega_i|$  and colour it by sign.

```
set.seed(3)
n <- 100; x <- rnorm(n, 5, 2); y <- 2 * x + rnorm(n, 0, 3)
xb <- mean(x); d <- x - xb; om <- d / sum(d^2); col <- ifelse(om >= 0, teal, rust)
R <- d > 0; L <- d < 0
xR <- sum(abs(d[R]) * x[R]) / sum(abs(d[R])); yR <- sum(abs(d[R]) * y[R]) / sum(abs(d[R]))
xL <- sum(abs(d[L]) * x[L]) / sum(abs(d[L])); yL <- sum(abs(d[L]) * y[L]) / sum(abs(d[L]))
plot(x, y, pch = 16, cex = 0.35 + 2 * abs(om) / max(abs(om)),
     col = paste0(col, "99"), xlab = "x", ylab = "y", main = "N = 100 Points")
abline(v = xb, lty = 3, col = "grey55")
abline(lm(y ~ x), lwd = 1.3, lty = 2, col = "grey55")
segments(xL, yL, xR, yR, lwd = 2.6, col = "grey15")
points(c(xL, xR), c(yL, yR), pch = 21, bg = c(rust, teal), col = "white", cex = 2.3, lwd = 2)
```




Figure 14.1: Outcome weights in a 100-point sample. Each small point's **size scales with its weight**  $|\omega_i|$ , coloured by sign — teal ( $\omega_i > 0$ ) pushes  $\hat{\beta}$  up, rust ( $\omega_i < 0$ ) pulls it down; the dotted line is the fulcrum  $x = \bar{x}$ . The two large dots are the  $|d|$ -weighted centroids of the points left and right of  $\bar{x}$ ; the solid segment joining them has slope  $\hat{\beta}$ , and the dashed line is the OLS fit (same slope, shifted through the grand mean).

```
vtab(c(`sum(omega)` = sum(om),      `sum(omega*x)` = sum(om * x),
      `positive mass` = sum(om[om > 0]), `negative mass` = sum(om[om < 0]),
      `OLS` = coef(lm(y ~ x))[[2]], `sum(omega*y)` = sum(om * y)), 4)
```

|               | value   |
|---------------|---------|
| sum(omega)    | 0.0000  |
| sum(omega*x)  | 1.0000  |
| positive mass | 0.2486  |
| negative mass | -0.2486 |
| OLS           | 1.8648  |
| sum(omega*y)  | 1.8648  |

Outcome weights in a 100-point sample. Each small point's **size scales with its weight**  $|\omega_i|$ ,

coloured by sign — teal ( $\omega_i > 0$ ) pushes  $\hat{\beta}$  up, rust ( $\omega_i < 0$ ) pulls it down; the dotted line is the fulcrum  $x = \bar{x}$ . The two large dots are the  $|d|$ -weighted centroids of the points left and right of  $\bar{x}$ ; the solid segment joining them has slope  $\hat{\beta}$ , and the dashed line is the OLS fit (same slope, shifted through the grand mean).

The printed sums confirm Equation 14.2: the weights split at  $\bar{x}$  into a positive mass ( $\approx +0.25$ ) and an equal negative one ( $\approx -0.25$ ), and  $\sum_i \omega_i x_i = 1$ .

But the figure shows more. Split the points at  $\bar{x}$ , weight each side by  $|d_i|$ , and take its centroid — the two large dots. **The segment joining them has slope  $\hat{\beta}$ :** the OLS slope is the rise over run between the low- $x$  and high- $x$  weighted centres.

The dashed fit shares that slope, shifted through the grand mean, so the two coincide only when the sides balance. This collapse to two points falls out of the slope formula.

### 14.2.1 Why it collapses to two points

Write  $d_i = x_i - \bar{x}$ , so the slope is  $\hat{\beta} = \sum_i d_i y_i / \sum_i d_i^2$ . Two elementary facts drive everything:

- $\sum_i d_i = 0$  — deviations from the mean sum to zero;
- $\sum_i d_i^2 = \sum_i d_i x_i$ , since  $\sum_i d_i x_i = \sum_i d_i (d_i + \bar{x}) = \sum_i d_i^2 + \bar{x} \sum_i d_i$ .

The second fact writes the slope as a ratio of two inner products with the *same* vector  $d$ :

$$\hat{\beta} = \frac{d \cdot y}{d \cdot x}. \quad (14.4)$$

Split the observations by the sign of  $d_i$ : let  $R = \{i : d_i > 0\}$  lie right of the mean and  $L = \{i : d_i < 0\}$  left, with  $d_i = 0$  points contributing nothing. Let  $W = \sum_{i \in R} d_i = \sum_{i \in R} |d_i|$  be the total weight on the right. Because  $\sum_i d_i = 0$ , the left must carry  $\sum_{i \in L} d_i = -W$ , that is  $\sum_{i \in L} |d_i| = W$ : **both sides hold the identical total weight  $W$ .**

Weight each side by  $|d_i|$  and form its centroid,  $\bar{v}_R = \sum_{i \in R} |d_i| v_i / W$  and  $\bar{v}_L = \sum_{i \in L} |d_i| v_i / W$  for any variable  $v$ . Using  $|d_i| = d_i$  on  $R$  and  $|d_i| = -d_i$  on  $L$ ,

$$d \cdot v = \sum_{i \in R} d_i v_i + \sum_{i \in L} d_i v_i = W \bar{v}_R - W \bar{v}_L = W(\bar{v}_R - \bar{v}_L). \quad (14.5)$$

An inner product with a zero-sum vector is a *between-group contrast* — the common weight  $W$  times the gap between the two weighted side-means. Apply it to the numerator ( $v = y$ ) and the denominator ( $v = x$ ):

$$\hat{\beta} = \frac{d \cdot y}{d \cdot x} = \frac{W(\bar{y}_R - \bar{y}_L)}{W(\bar{x}_R - \bar{x}_L)} = \frac{\bar{y}_R - \bar{y}_L}{\bar{x}_R - \bar{x}_L}. \quad (14.6)$$

The  $W$  cancels, and what remains is the plain rise over run between the two  $|d|$ -weighted centroids — the two dots in the figure.

The construction gives exactly *two* points because  $d_i$  takes two signs: a centred regressor has a left and a right of the mean, and the sign is constant within each side, so each half collapses to a single weighted centre.

The weight must be  $|d_i|$ , not a plain average, because it already sits inside the OLS sums  $d \cdot y$  and  $d \cdot x$ : the dot is a centre of mass with mass  $|d_i|$ , and the same mass fixes both coordinates — weighting  $y$  but not  $x$  gives the wrong slope.

## 14.3 When the sign flips

In simple regression the sign of  $\omega_i$  is just the side of the mean: half the weights are negative, structurally, and a low- $x$  point *should* count against a high- $x$  one. With controls  $\tilde{V}_i$  becomes a *residual*, and a row can flip to the sign its raw position denies — a hidden extrapolation the non-negative  $w_i$  can never show.

Take a binary treatment  $D$  with covariate  $x$  and thin overlap — controls low, treated high — and fit  $y \sim D + x$ . Now  $\omega_i \propto \tilde{V}_i = D_i - \hat{D}_i$ , with  $\hat{D}$  the *linear* projection of  $D$  on  $(1, x)$ . Where overlap fails, that  $\hat{D}$  leaves  $[0, 1]$ : a treated unit with  $\hat{D}_i > 1$  gets  $\tilde{V}_i < 0$ , so its outcome is *subtracted* from  $\hat{\beta}$  — a sign  $\tilde{V}_i^2 \geq 0$  hides.

**Simulation — thin overlap.** The DGP:  $n = 80$ , with 40 controls at  $x \sim N(2, 0.7^2)$  and 40 treated at  $x \sim N(5, 0.7^2)$  — barely overlapping — and  $y_i = D_i + 0.6x_i + \varepsilon_i$ ,  $\varepsilon_i \sim N(0, 0.5^2)$ , so the true effect is 1.

```
set.seed(11)
n0 <- 40; n1 <- 40
x <- c(rnorm(n0, 2.0, 0.7), rnorm(n1, 5.0, 0.7))
d <- c(rep(0, n0), rep(1, n1))
y <- 1.0 * d + 0.6 * x + rnorm(n0 + n1, 0, 0.5)
X <- cbind(1, x); dh <- as.vector(X %*% solve(crossprod(X)) %*% t(X) %*% d)
vt <- d - dh; om <- vt / sum(vt^2)
op <- par(mfrow = c(1, 2), mar = c(4, 4, 2.5, 1))
plot(x, d, pch = ifelse(d == 1, 17, 16), col = "grey45", ylim = c(-0.35, 1.35),
     xlab = "x", ylab = "D and D-hat", main = "Linear D-hat leaves [0, 1]")
abline(h = c(0, 1), lty = 3, col = "grey70"); ox <- order(x)
lines(x[ox], dh[ox], lwd = 2, col = teal)
col <- ifelse(om >= 0, teal, rust)
plot(x, y, pch = ifelse(d == 1, 17, 16), col = paste0(col, "CC"),
     cex = 0.7 + 3 * abs(om) / max(abs(om)), xlab = "x", ylab = "y",
     main = "Outcome-weight sign")
tn <- which(d == 1 & om < 0)
points(x[tn], y[tn], pch = 2, cex = 2.3, col = rust, lwd = 2)
```




Figure 14.2: Outcome weights under thin overlap. Left: the linear projection  $\hat{D}$  of treatment on  $x$  leaves  $[0, 1]$ , so residuals change sign in the extrapolation zone. Right: outcome weights coloured by sign (teal +, rust −), triangles treated and circles control, size  $\propto |\omega_i|$ . Treated units at high  $x$  (open rust triangles) carry *negative* weights — their outcomes are subtracted from  $\hat{\beta}$  — a sign the effect weights  $\tilde{V}_i^2 \geq 0$  cannot reveal.

```
par(op)
vtab(c(`treated, omega<0`      = sum(d == 1 & om < 0),
      `treated, D-hat>1`      = sum(d == 1 & dh > 1),
      `control, omega>0`      = sum(d == 0 & om > 0),
      `control, D-hat<0`      = sum(d == 0 & dh < 0),
      `beta (lm)`              = coef(lm(y ~ d + x))["d"],
      `beta (sum omega*y)`     = sum(om * y)), 4)
```

|                    | value  |
|--------------------|--------|
| treated, omega<0   | 17.000 |
| treated, D-hat>1   | 17.000 |
| control, omega>0   | 15.000 |
| control, D-hat<0   | 15.000 |
| beta (lm)          | 1.287  |
| beta (sum omega*y) | 1.287  |

Outcome weights under thin overlap. Left: the linear projection  $\hat{D}$  of treatment on  $x$  leaves  $[0, 1]$ , so residuals change sign in the extrapolation zone. Right: outcome weights coloured by sign (teal +, rust −), triangles treated and circles control, size  $\propto |\omega_i|$ . Treated units at high  $x$  (open rust triangles) carry *negative* weights — their outcomes are subtracted from  $\hat{\beta}$  — a sign the effect weights  $\tilde{V}_i^2 \geq 0$  cannot reveal.

Seventeen of the forty treated units carry a negative weight, and  $\hat{\beta}$  still equals  $\sum_i \omega_i y_i$  exactly. Nothing is wrong with the algebra: asked to compare treated and control where they barely coexist, the linear model extrapolates, and the extrapolation surfaces as flipped signs.

Which negatives matter? Two kinds.

*Structural*: regressed on the treatment alone, every treated unit is positive and every control negative, and  $\hat{\beta} = \bar{y}_1 - \bar{y}_0$ .

*Wrong-sign*: the opposite pairing — a treated unit with  $\omega_i < 0$ , or a control with  $\omega_i > 0$ . This cannot happen without covariates. It appears only when the residual  $\tilde{V}_i = D_i - \hat{D}_i$  carries a fitted  $\hat{D}_i$  outside  $[0, 1]$ , and the counts above confirm the treated units with  $\omega_i < 0$  are exactly those with  $\hat{D}_i > 1$ .

So the test is one line: **fit the treatment on the controls; any  $\hat{D}_i \notin [0, 1]$  is an extrapolated, wrong-sign unit** — a point where treated and control covariates fail to overlap. (With a continuous regressor there is no bounded range to leave: a negative  $\omega_i$  just means “below the conditional mean”, and overlap is read from leverage, not sign.)

A wrong-sign weight is extrapolation: the model reaches outside common support, and  $\hat{\beta}$  can leave the convex hull of the observed outcomes — outside the range of any real comparison.

This is no failure of OLS; it locates the reliance on functional form. The responses are the usual ones: trim to common support (the negatives then vanish), interpolate instead (matching, or overlap weights that are non-negative by design), or report how far  $\hat{\beta}$  moves when the extrapolated units are dropped.

This is the implied-weights diagnostic of Chattopadhyay and Zubizarreta (2023); their `lmw` package computes the weights and the balance, extrapolation, and influence diagnostics built on them, and `OutcomeWeights` (Knaus 2024) carries it to machine-learning estimators. On the data above, `lmw`’s negative weights are exactly the units our  $\hat{D} \notin [0, 1]$  rule flags:

```
library(lmw)
df <- data.frame(y, d, x)
# lmw's implied weights for the same regression (y ~ d + x, ATE)
L <- lmw(~ d + x, data = df, treat = "d", estimand = "ATE", method = "URI")
manual <- (d == 1 & dh > 1) | (d == 0 & dh < 0) # our flag: linear D-hat outside [0, 1]
knitr::kable(
  data.frame(
    quantity = c("lmw negative weights", "our D-hat outside [0, 1]", "same units flagged"),
    value     = c(as.character(sum(L$weights < 0)),
                  as.character(sum(manual)),
                  as.character(all((L$weights < 0) == manual))),
    col.names = c("", "value"), align = c("l", "r"))
```

|                          | value |
|--------------------------|-------|
| lmw negative weights     | 32    |
| our D-hat outside [0, 1] | 32    |
| same units flagged       | TRUE  |

The mainstream overlap check takes the other route, through the propensity score. `WeightIt` estimates it and reports the effective sample size and weight range, `cobalt::bal.plot` draws the two groups' score densities, and the fixes are trimming (Crump et al. 2009) or overlap weights, the "ATO" estimand (Li, Morgan, and Zaslavsky 2018).

On this data both routes agree — the logistic score *perfectly separates*:

```
w <- WeightIt::weightit(d ~ x, data = df, method = "glm", estimand = "ATE")
vtab(c(`max control PS` = max(w$ps[d == 0]),
      `min treated PS` = min(w$ps[d == 1])), 3)
```

|                | value |
|----------------|-------|
| max control PS | 0     |
| min treated PS | 1     |

Both routes point to the same problem — the groups barely overlap — using different tools: the regression's own implied weights (`lmw`), or a separate propensity-score model (`WeightIt`, `cobalt`). Here overlap fails completely, so neither can help; there is no shared region to weight toward.

When overlap is partial, the diagnostic becomes a sensitivity check: trim the flagged units, refit, and see whether the estimate moves. Here it moves by a lot. The outcome is convex in  $x$ , controls sit low on the curve and treated sit high, and the straight  $y \sim D + x$  reports  $\hat{\beta} = 0.69$  — a third short of the true effect of 1, an artifact of one line forced through a curve the groups do not share.

**Simulation — fragile estimate under misspecification.** The DGP: 150 controls near  $x = 2.5$  and 150 treated near  $x = 4.5$  (both sd 1.1), with  $y_i = D_i + 0.8(x_i - 3.5)^2 + 0.3x_i + \varepsilon_i$ ,  $\varepsilon_i \sim N(0, 0.5^2)$  — a constant true effect of 1 sitting on a convex baseline.

```
set.seed(7)
n <- 150
x <- c(rnorm(n, 2.5, 1.1), rnorm(n, 4.5, 1.1))
d <- c(rep(0, n), rep(1, n))
y <- 1.0 * d + 0.8 * (x - 3.5)^2 + 0.3 * x + rnorm(2 * n, 0, 0.5)
df <- data.frame(y, d, x)
X <- cbind(1, x); dh <- as.vector(X %*% solve(crossprod(X)) %*% t(X) %*% d)
flag <- (d == 1 & dh > 1) | (d == 0 & dh < 0)
pch <- ifelse(d == 1, 17, 16); col <- ifelse(d == 1, teal, rust)
plot(x, y, type = "n", xlab = "x", ylab = "y")
xs <- seq(min(x), max(x), length = 200)
mc <- function(z) 0.8 * (z - 3.5)^2 + 0.3 * z # true control mean E[y | x, D = 0]
lines(xs, mc(xs), col = rust, lwd = 2.4) # control mean
lines(xs, mc(xs) + 1, col = teal, lwd = 2.4, lty = 3) # treated mean = control + true eff
points(x, y, pch = pch, col = paste0(col, "88"), cex = 0.85) # all points, by gr
points(x[flag], y[flag], pch = 1, col = "grey15", cex = 1.7, lwd = 1.4) # ring the flagged u
arrows(3.5, mc(3.5), 3.5, mc(3.5) + 1, code = 3, length = 0.05, angle = 90, lwd = 2.4, col = "g
segments(3.5, mc(3.5) + 0.5, 2.05, 13.3, col = "grey55", lwd = 1)
text(2.0, 14, "gap = true effect = 1", pos = 4, cex = 0.82, col = "grey10")
```

```

legend("bottomright", bty = "n", cex = 0.82,
      legend = c("control", "treated", "flagged (trimmed)", "control mean", "treated mean"),
      pch = c(16, 17, 1, NA, NA), lty = c(NA, NA, NA, 1, 3),
      col = c(rust, teal, "grey15", rust, teal))

```




Figure 14.3: A fragile OLS gap. The two curves are the true means — control (rust, solid) and treated (teal, dotted), a constant 1 apart: that gap is the true effect (arrow). Controls sit low on the curve and treated high, so the groups share only the middle. Ringed points are the extrapolating units the one-line rule flags, where the groups do not overlap. Forcing a straight  $y \sim D + x$  across the whole range reports  $\hat{\beta} = 0.69$ ; dropping the flagged units, or reweighting to the overlap, restores  $\approx 1$ .

```

w_ato <- WeightIt::weightit(d ~ x, data = df, method = "glm", estimand = "ATO")
vtab(c(`flagged units`      = sum(flag),
      `beta (full)`        = coef(lm(y ~ d + x))["d"],
      `beta (trimmed)`     = coef(lm(y ~ d + x, data = df[!flag, ]))["d"],
      `beta (overlap-wtd)` = coef(lm(y ~ d + x, weights = w_ato$weights))["d"]), 3)

```

|                    | value  |
|--------------------|--------|
| flagged units      | 37.000 |
| beta (full)        | 0.688  |
| beta (trimmed)     | 0.991  |
| beta (overlap-wtd) | 0.979  |

A fragile OLS gap. The two curves are the true means — control (rust, solid) and treated (teal, dotted), a constant 1 apart: that gap is the true effect (arrow). Controls sit low on the curve and treated high, so the groups share only the middle. Ringed points are the extrapolating units the one-line rule flags, where the groups do not overlap. Forcing a straight  $y \sim D + x$  across the whole range reports  $\hat{\beta} = 0.69$ ; dropping the flagged units, or reweighting to the overlap, restores  $\approx 1$ .

The outcome weights find the artifact. Thirty-seven units carry the wrong sign, all in the non-overlapping tails. Trim them and refit and  $\hat{\beta}$  returns to 0.99; reweight to the overlap instead — the "ATO" weights — and it returns 0.98. Either way the coefficient had leaned on the extrapolation, and the diagnostic showed it, read off the regression itself with no separate model.

Whether it matters depends on the data. Here trimming moved  $\hat{\beta}$  a lot, so the estimate leaned on the extrapolation. Had the outcome been linear — the two groups sharing one straight line — the same trim would barely move  $\hat{\beta}$ : the negative weights are still there, but dropping those units changes almost nothing. The check tells you which case you are in.

And it survives where the propensity score cannot. On the perfectly-separated data above the logistic model fails to converge and inverse-probability weighting collapses, yet the implied weights still flag every extrapolating unit.

## 14.4 Out of OLS: the same weights beyond the linear model

The outcome reading has one advantage the effect reading lacks: it survives leaving OLS. Any estimator that can be written  $\hat{\tau} = \sum_i \omega_i Y_i = \omega'Y$  — a linear combination of the observed outcomes — has outcome weights, whatever machinery produced them.

Knaus (2024) makes this the organising idea, and his `OutcomeWeights` package computes  $\omega$  for AIPW, double machine learning, and the causal forest.

### 14.4.1 The weights, derived

The OLS outcome weights (Equation 14.2) followed from one fact: in the FWL numerator  $\sum_i \tilde{V}_i y_i$ , the outcome  $y_i$  appears once, linearly, so  $\hat{\beta} = \sum_i \omega_i y_i$  with  $\omega_i = \tilde{V}_i / \sum_j \tilde{V}_j^2$ . The same fact holds whenever the outcome model is a *smoother* — a map  $\hat{m} = SY$  that is linear in the training outcomes. Substituting  $SY$  for  $\hat{m}$  keeps  $Y$  linear in the numerator, and reading off its coefficient gives  $\omega$ . We carry this out for PLR and AIPW.

*PLR.* Take the partially-linear estimator (the R-learner behind the causal forest). It is FWL with the projections done by a flexible learner: residualise treatment and outcome on the covariates, then regress one residual on the other,

$$\hat{\tau} = \frac{\sum_i (D_i - \hat{e}_i)(Y_i - \hat{m}_i)}{\sum_i (D_i - \hat{e}_i)^2}, \quad (14.7)$$

with  $\hat{e} = \hat{E}[D \mid X]$  and  $\hat{m} = \hat{E}[Y \mid X]$  the fitted nuisances.

The outcome nuisance is a *smoother*: its fitted values are a linear map of the training outcomes,  $\hat{m} = SY$  — true for OLS, ridge, kernels, and honest forests alike. So the outcome residual is  $Y - \hat{m} = (\mathbf{I} - S)Y$ . Writing  $\tilde{D} = D - \hat{e}$  for the treatment residual and substituting,

$$\hat{\tau} = \frac{\tilde{D}'(\mathbf{I} - S)Y}{\tilde{D}'\tilde{D}}. \quad (14.8)$$

Now  $Y$  appears once, linearly, in the numerator. Read off its coefficient and the weights are done:

$$\hat{\tau} = \omega'Y, \quad \omega' = \frac{\tilde{D}'(\mathbf{I} - S)}{\tilde{D}'\tilde{D}}. \quad (14.9)$$

That is Equation 14.2 — the OLS weights — with the projection annihilator  $M_X$  replaced by the smoother annihilator  $\mathbf{I} - S$ . OLS is the case of a linear learner,  $S = P_X$ : there  $S$  is a projection,  $\mathbf{I} - S = M_X$ , and Equation 14.9 is Equation 14.2 exactly. Component by component, the weight on observation  $i$  is

$$\omega_i = \frac{\tilde{D}_i - \sum_j S_{ji} \tilde{D}_j}{\sum_k \tilde{D}_k^2} : \quad (14.10)$$

the treatment residual of unit  $i$ , minus how much other units' treatment residuals leak to  $i$  through the outcome smoother, normalised by the total treatment-residual variance.

The outcome weights tell which rows drive the estimate; the *effect* weights tell whose effect it estimates. Substitute  $Y_i = \tau_i D_i + m(X_i) + \varepsilon_i$  into the PLR numerator. With consistent nuisances  $\hat{m} \approx m$ , the noise vanishes in expectation and the numerator reduces to  $\sum_i \tilde{D}_i D_i \tau_i$ , so the effect weight on  $\tau_i$  is  $\tilde{D}_i D_i = (D_i - e_i)D_i$ . Since  $D$  is binary ( $D_i^2 = D_i$ ),

$$E[\tilde{D}_i D_i \mid X_i] = E[D_i \mid X_i] - e_i E[D_i \mid X_i] = e_i - e_i^2 = e_i(1 - e_i). \quad (14.11)$$

The denominator  $\tilde{D}'\tilde{D} = \sum_i \tilde{D}_i^2$  has the same expectation,  $\sum_i e_i(1 - e_i)$ , since  $\text{Var}(D_i \mid X_i) = e_i(1 - e_i)$ . Both numerator and denominator weight by propensity variance, and the ratio is the variance-weighted average treatment effect,

$$\tau_{\text{VWATE}} = \frac{\sum_i e(X_i)(1 - e(X_i)) \tau(X_i)}{\sum_i e(X_i)(1 - e(X_i))}, \quad (14.12)$$

not the plain ATE  $\bar{\tau} = n^{-1} \sum_i \tau_i$ . The two coincide only when effects are constant or treatment is randomly assigned. The weighting is the same  $e(1 - e)$  that appeared in the previous chapter's Equation 13.17, where Angrist (1998) showed OLS with saturated strata weights each stratum by

$n_x \hat{D}(x)(1 - \hat{D}(x))$ . That result required discrete strata and saturated controls; Equation 14.12 is the same propensity-variance weighting with continuous covariates and a flexible learner. The individual  $\tau(X_i)$  are never observed or estimated — they are the structural unit-level effects in  $Y_i = \tau(X_i)D_i + m(X_i) + \varepsilon_i$ . PLR returns one number,  $\hat{\tau}$ , that converges to their  $e(1 - e)$ -weighted average; to estimate  $\tau(X_i)$  itself one needs a CATE estimator (a causal forest, say).

Three moves did it: write the estimator, substitute the smoother  $\hat{m} = SY$ , read off the coefficient on  $Y$ . They work for any estimator built from smoothers; only the operator on  $Y$  changes. AIPW is worth carrying through, because its operator is more than  $\mathbf{I} - S$ .

*AIPW*. With outcome smoothers  $\hat{\mu}_1 = S_1 Y$ ,  $\hat{\mu}_0 = S_0 Y$  and propensity  $\hat{e}$ , the score of unit  $i$  is

$$\psi_i = \hat{\mu}_{1i} - \hat{\mu}_{0i} + \frac{D_i}{\hat{e}_i} (Y_i - \hat{\mu}_{1i}) - \frac{1 - D_i}{1 - \hat{e}_i} (Y_i - \hat{\mu}_{0i}), \quad \hat{\tau} = \frac{1}{n} \sum_i \psi_i. \quad (14.13)$$

Each  $\hat{\mu}_{ji}$  is  $(S_j Y)_i$ , so  $Y$  again enters only linearly; collecting its coefficient gives  $\hat{\tau} = \omega' Y$  with

$$T = S_1 - S_0 + \text{diag}\left(\frac{D}{\hat{e}}\right)(\mathbf{I} - S_1) - \text{diag}\left(\frac{1-D}{1-\hat{e}}\right)(\mathbf{I} - S_0), \quad \omega' = \frac{\mathbf{1}' T}{n}. \quad (14.14)$$

That is,  $\omega_j$  is the  $j$ -th column-sum of  $T$  divided by  $n$ : observation  $j$ 's outcome enters the ATE through its direct IPW term and through its indirect contribution via the two outcome smoothers  $S_1$ ,  $S_0$  to every other unit's augmentation correction.

Comparing the two:

|               | PLR (Equation 14.9)                                                           | AIPW                                   |
|---------------|-------------------------------------------------------------------------------|----------------------------------------|
| Matrix form   | $\omega' = \tilde{D}'(\mathbf{I} - S) / \tilde{D}' \tilde{D}$                 | $\omega' = \mathbf{1}' T / n$          |
| Per unit      | $\omega_i = (\tilde{D}_i - \sum_j S_{ji} \tilde{D}_j) / \sum_k \tilde{D}_k^2$ | $\omega_j = \frac{1}{n} \sum_i T_{ij}$ |
| Smoother      | one $S$ (outcome on $X$ )                                                     | two: $S_1$ (treated), $S_0$ (control)  |
| Normalisation | treatment-residual variance                                                   | sample size $n$                        |

Both follow the same recipe — substitute the smoother  $\hat{m} = SY$ , read off  $Y$ 's coefficient — and differ only in the operator acting on  $Y$ .

**Simulation — heterogeneous effects (DML).** Under a heterogeneous effect the two estimators diverge, and the weights show why. We simulate  $n = 2000$  units: five standard-normal covariates  $X$ , selection  $W_i \sim \text{Bernoulli}(\Lambda(-1 + 1.2 X_{i1}))$  with  $\Lambda$  the logistic function (treatment is scarce and tied to  $X_1$ ), and outcome  $Y_i = \tau_i W_i + X_{i1} + 0.5 X_{i2} + \varepsilon_i$  with an effect  $\tau_i = 1 + X_{i1}$  that varies across units — a  $W \times X_1$  interaction. The average effect is  $E[\tau] = 1$ , but units with larger  $X_1$  have larger effects. `dml_with_smoother` returns a partially-linear (PLR) and an AIPW estimator; `get_outcome_weights` extracts the weight matrix  $\omega$ , one row per estimator.

```
library(OutcomeWeights)
set.seed(1)
n <- 2000; p <- 5
X <- matrix(rnorm(n * p), n, p)
etrue <- plogis(-1 + 1.2 * X[, 1]) # asymmetric selection on X1
```

```

W <- rbinom(n, 1, etrue)
tau <- 1 + X[, 1] # heterogeneous effect: W x X1
Y <- tau * W + X[, 1] + 0.5 * X[, 2] + rnorm(n)
dml <- dml_with_smoother(Y, W, X, estimators = c("PLR", "AIPW_ATE"))
om <- get_outcome_weights(dml) # om$omega: 2 x n, om$treat = W
invisible(capture.output(s <- summary(dml)))
W1 <- om$treat == 1
tab <- t(sapply(rownames(om$omega), function(nm) {
  w <- om$omega[nm, ]
  c(estimate = s[nm, 1], `omega %*% Y` = sum(w * Y),
    treated = sum(w[W1]), control = sum(w[!W1]))
}))
ov <- etrue * (1 - etrue) # e(1-e): PLR's per-unit weight
# ATE = plain mean of tau; overlap-ATE = e(1-e)-weighted mean of tau (PLR's target)
cat(sprintf("true ATE = %.2f    true overlap-ATE = %.2f\n", mean(tau), sum(ov * tau) / sum(ov)))

```

```
true ATE = 0.99    true overlap-ATE = 1.31
```

```
mtab(tab, 3)
```

|          | estimate | omega %*% Y | treated | control |
|----------|----------|-------------|---------|---------|
| PLR      | 1.287    | 1.287       | 0.997   | -0.997  |
| AIPW-ATE | 1.074    | 1.074       | 1.000   | -1.000  |

Two things stand out. First, the estimators now *disagree* — PLR reports 1.29, AIPW 1.07 — and neither is wrong; they target different populations. PLR targets the VWATE (Equation 14.12), not the ATE. Here treatment selection loads on  $X_1$  through  $e = \Lambda(-1 + 1.2 X_1)$ , and  $\tau = 1 + X_1$  also rises in  $X_1$ , so the units PLR overweights — those with  $e$  near  $\frac{1}{2}$ , around  $X_1 \approx 0.8$  — carry above-average effects, pushing PLR's estimate above the true ATE. AIPW instead reweights by  $1/e$  and  $1/(1 - e)$  to undo selection, returning the plain average, the ATE ( $\approx 1$ ). Same outcomes  $Y$ , weighted two ways, two population-averages of  $\tau$  — and the weights are what make the difference visible.

Second, the weights still reproduce each estimate exactly and sum to zero overall. They part on the other constraint from §1,  $\sum_i \omega_i V_i = 1$ : for a binary treatment that constraint is the treated weights summing to +1 (and, with the zero total, controls to -1). AIPW keeps it exactly; the partially-linear estimator sums to  $\pm 0.99$ , quietly breaking the within-group budget that OLS and AIPW hold.

That divergence shows what the representation is for. In Knaus's (2024) framing  $\omega$  does three things: it puts OLS, IPW, AIPW, DML, and the causal forest on **one footing** ( $\omega'Y$ ); it **characterises the target population** each reports — the PLR-versus-AIPW gap above; and it **checks covariate balance** via `cobalt`, flagging when weights break a property usually assumed — the sum-to- $\pm 1$  within treated and control groups that PLR violates here.

The package does exactly this — reads  $(S, S_0, S_1, \hat{e})$  from the fitted object, forms these operators, and reads off  $\omega$ . The formulas above take a few lines to compute by hand; here we do it for both PLR and AIPW and check they match the package's weights:

```
n <- length(Y)
S <- dml$NuPa.hat$smoothers$S[1, , ]      # outcome smoother E[Y|X]
S1 <- dml$NuPa.hat$smoothers$S.d1[1, , ]   # E[Y|X, D=1] smoother
S0 <- dml$NuPa.hat$smoothers$S.d0[1, , ]   # E[Y|X, D=0] smoother
e <- dml$NuPa.hat$predictions$D.hat[, 1]   # fitted propensity

# PLR: omega = (I - S)'(W - e) / ||W - e||^2
Wtil <- W - e
om_plr <- as.numeric((Wtil - crossprod(S, Wtil)) / sum(Wtil^2))

# AIPW: omega = 1'T / n, T = S1 - S0 + diag(W/e)(I - S1) - diag((1-W)/(1-e))(I - S0)
Imat <- diag(n)
Tm <- S1 - S0 + (W / e) * (Imat - S1) - ((1 - W) / (1 - e)) * (Imat - S0)
om_aipw <- colSums(Tm) / n

vtab(c(`PLR: max |by hand - package|` = max(abs(om_plr - om$omega["PLR", ])),
      `PLR: omega' Y` = sum(om_plr * Y),
      `AIPW: max |by hand - package|` = max(abs(om_aipw - om$omega["AIPW-ATE", ])),
      `AIPW: omega' Y` = sum(om_aipw * Y)), 4)
```

|                              | value  |
|------------------------------|--------|
| PLR: max  by hand - package  | 0.0000 |
| PLR: omega' Y                | 1.2873 |
| AIPW: max  by hand - package | 0.0000 |
| AIPW: omega' Y               | 1.0738 |

Agreement to machine precision: both by-hand  $\omega$  vectors match the package, and each  $\omega'Y$  returns its point estimate — the derivation above, in code.

### 14.4.2 Two uses: balance and extrapolation

Having  $\omega$ , two checks run directly on it. Does the estimator's implied weighting balance the covariates? Hand  $\omega$  to `cobalt` and read the largest standardised mean difference. And is it extrapolating? Count the wrong-sign weights — a control with  $\omega_i > 0$  or a treated unit with  $\omega_i < 0$ , the §3 signature of a comparison reaching off common support.

```
library(cobalt)
uses <- t(sapply(rownames(om$omega), function(nm) {
  w <- om$omega[nm, ]
  c(`max |SMD|` = max(abs(bal.tab(as.data.frame(X), treat = W, weights = w)$Balance$Dif:
    `wrong-sign units` = sum((om$treat == 1 & w < 0) | (om$treat == 0 & w > 0)))
```

```
}))
mtab(uses, 3)
```

|          | max  SMD | wrong-sign units |
|----------|----------|------------------|
| PLR      | 0.025    | 21               |
| AIPW-ATE | 0.091    | 0                |

Raw, the groups are badly imbalanced — the largest covariate SMD is 1.11, since treatment was tied to  $X_1$ . Both weightings pull it down to near zero (0.03 for PLR, 0.09 for AIPW): each balances the covariates, which for a causal forest you could not otherwise check.

They differ on extrapolation. PLR carries 21 wrong-sign controls — it leans slightly off common support — while AIPW carries none. Neither is disqualifying here; the point is that balance and extrapolation are now *checkable*, on estimators that never exposed a weighting.

## 14.5 The DiD setting

When the regressor is a treatment indicator, every unit in a group gets the same weight, so  $\hat{\beta}$  reduces to a contrast of group means. Regressed on the treatment alone it is the difference in means between treated and control. Add a time period and its interaction with treatment and it becomes a difference-in-differences across four group-by-period cells. We take these in turn, then bring in covariates.

**Simulation — treatment alone.** Regress  $y$  on the treatment  $D$  alone — here 30 treated drawn  $y \sim N(5, 1)$  and 50 controls  $y \sim N(2, 1)$ . The residual  $\tilde{V}_i = D_i - \bar{D}$  is constant within each group, so the weight is uniform within a group and  $\hat{\beta}$  reduces to the two group means:  $\hat{\beta} = \bar{y}_1 - \bar{y}_0$ , the difference in means. Every treated unit then carries positive weight, every control negative — the structural pattern of the previous section, with the wrong-sign off-diagonal empty until covariates enter.

```
set.seed(7)
n1 <- 30; n0 <- 50
d <- c(rep(1, n1), rep(0, n0)); y <- c(rnorm(n1, 5), rnorm(n0, 2))
dev <- d - mean(d); om <- dev / sum(dev^2); col <- ifelse(om >= 0, teal, rust)
y1 <- mean(y[d == 1]); y0 <- mean(y[d == 0])
plot(jitter(d, 0.4), y, pch = 16, col = paste0(col, "55"), cex = 0.9,
     xlim = c(-0.3, 1.3), xaxt = "n", xlab = "", ylab = "y", main = "Regression on treatment")
axis(1, at = c(0, 1), labels = c("control", "treated"))
segments(0, y0, 1, y1, lwd = 2.6, col = "grey15")
points(c(0, 1), c(y0, y1), pch = 21, bg = c(rust, teal), col = "white", cex = 2.5, lwd = 2)
```




Figure 14.4: Regression on the treatment alone. Points are the two groups (rust control, teal treated); the large dots are the group means, and the segment joining them has slope  $\hat{\beta} = \bar{y}_1 - \bar{y}_0$ , the difference in means.

```
vtab(c(`beta (lm)` = coef(lm(y ~ d))["d"], `y1 - y0` = y1 - y0,
      `off-diagonal weights` = sum((d == 1 & om < 0) | (d == 0 & om > 0))), 3)
```

|                      | value |
|----------------------|-------|
| beta (lm)            | 3.318 |
| y1 - y0              | 3.318 |
| off-diagonal weights | 0.000 |

Regression on the treatment alone. Points are the two groups (rust control, teal treated); the large dots are the group means, and the segment joining them has slope  $\hat{\beta} = \bar{y}_1 - \bar{y}_0$ , the difference in means.

The estimate equals  $\bar{y}_1 - \bar{y}_0$ , and no weight carries the wrong sign — the off-diagonal count is zero.

**Simulation —  $2 \times 2$  difference-in-differences.** Add a period and interact it with the group, and the design has four cells. Each collapses to its mean, and the interaction coefficient is their double contrast,

$$\widehat{\text{DiD}} = (\bar{y}_{T,\text{post}} - \bar{y}_{T,\text{pre}}) - (\bar{y}_{C,\text{post}} - \bar{y}_{C,\text{pre}}). \quad (14.15)$$

Now there are four cell means, and with them four slopes — the two group trends over time and the two cross-sectional gaps. The DiD is the difference between either parallel pair, trend minus trend or gap minus gap. We simulate 200 units in each group×period cell from  $y = 3 + 2g + p + 2(g \cdot p) + \varepsilon$ , so the true DiD (the interaction) is 2:

```
set.seed(3)
g <- rep(c(1, 1, 0, 0), each = 200); p <- rep(c(0, 1, 0, 1), each = 200)
y <- 3 + 2 * g + 1 * p + 2 * (g * p) + rnorm(800)
m <- tapply(y, list(g, p), mean)
Tpre <- m["1", "0"]; Tpost <- m["1", "1"]; Cpre <- m["0", "0"]; Cpost <- m["0", "1"]
cf <- Tpre + (Cpost - Cpre)
plot(NA, xlim = c(-0.15, 1.25), ylim = range(y), xaxt = "n", xlab = "", ylab = "y",
     main = "Difference-in-differences")
axis(1, at = c(0, 1), labels = c("pre", "post"))
segments(0, Tpre, 1, Tpost, lwd = 2.4, col = teal)
segments(0, Cpre, 1, Cpost, lwd = 2.4, col = rust)
segments(0, Tpre, 1, cf, lty = 2, lwd = 1.6, col = "grey55")
segments(1, cf, 1, Tpost, lwd = 4, col = "grey15")
points(c(0, 1, 0, 1), c(Tpre, Tpost, Cpre, Cpost),
       pch = 21, bg = c(teal, teal, rust, rust), col = "white", cex = 2.3, lwd = 2)
text(1.02, (cf + Tpost) / 2, "DiD", pos = 4, cex = 0.9, col = "grey15")
```




Figure 14.5: The 2×2 difference-in-differences: four cell means (dots), the treated (teal) and control (rust) trends, the parallel-trends counterfactual (dashed), and the DiD (black bar) as the gap between the treated post-mean and that counterfactual.

```
vtab(c(`treated trend` = Tpost - Tpre, `control trend` = Cpost - Cpre,
      `pre gap` = Tpre - Cpre, `post gap` = Tpost - Cpost,
      `DiD (interaction)` = coef(lm(y ~ g * p))["g:p"]), 3)
```

|                   | value |
|-------------------|-------|
| treated trend     | 3.029 |
| control trend     | 0.897 |
| pre gap           | 2.019 |
| post gap          | 4.151 |
| DiD (interaction) | 2.132 |

The  $2 \times 2$  difference-in-differences: four cell means (dots), the treated (teal) and control (rust) trends, the parallel-trends counterfactual (dashed), and the DiD (black bar) as the gap between the treated post-mean and that counterfactual.

The estimated interaction recovers the true DiD of 2, and the four cell-mean weights carry the structural signs (+, −, −, +) as long as every cell is populated.

The clean  $2 \times 2$  has trivial weights — four cell means, each  $\pm 1$ . Real difference-in-differences carries covariates: parallel trends holds only *conditional* on  $X$ , and the comparison has to be modelled. That model is a weighting of outcomes too.

The doubly-robust DiD of Sant’Anna and Zhao (2020) is AIPW read on the outcome *change*  $\Delta Y = Y_1 - Y_0$ : a propensity for treatment, a control-group regression of  $\Delta Y$  on  $X$ , combined. By the recipe of the last section its estimate is  $\widehat{ATT} = \omega' \Delta Y$ , with  $\omega' = a'(\mathbf{I} - S)$  — the control-trend smoother  $S$  and the group weights  $a = w_{\text{treat}} - w_{\text{cont}}$ .

**Simulation — doubly-robust DiD.** Two periods and  $n = 2000$ : treatment is selected on  $x_1$  with thin overlap,  $\Pr(D = 1) = \Lambda(-0.5 + 2x_1)$ ; the control trend depends on  $x_1$ , so trends are parallel only given  $X$ ; the true ATT is 1. We fit DR-DiD with DRDID, rebuild its weights by hand, and check.

```
library(DRDID)
set.seed(2)
n <- 2000
x1 <- rnorm(n); x2 <- rnorm(n); Xcov <- cbind(1, x1, x2)
D <- rbinom(n, 1, plogis(-0.5 + 2 * x1)) # selection into treatment on x1 (thin overlap)
y0 <- 1 + x1 + 0.5 * x2 + rnorm(n) # pre-period outcome
y1 <- y0 + (0.5 + 0.4 * x1) + 1 * D + rnorm(n) # control trend depends on x1; true ATT = 1
dY <- y1 - y0
att <- drdid_panel(y1 = y1, y0 = y0, D = D, covariates = Xcov, boot = FALSE)$ATT

# rebuild the DR-DiD weights: omega = (I - S)' a, acting on the outcome change dY
e <- pmin(as.vector(glm(D ~ x1 + x2, family = binomial)$fitted), 1 - 1e-6)
keep <- rep(TRUE, n); keep[D == 0] <- e[D == 0] < 0.995 # DRDID's control trimming
Xc <- Xcov[D == 0, ]
S <- matrix(0, n, n); S[, D == 0] <- Xcov %*% solve(crossprod(Xc)) %*% t(Xc)
```

```

w.t <- keep * D
w.c <- keep * e * (1 - D) / (1 - e)
a <- w.t / sum(w.t) - w.c / sum(w.c) # group weights, treated minus reweighted controls
omega <- as.vector(a - crossprod(S, a))

Xk <- cbind(x1, x2); sds <- apply(Xk, 2, sd)
vtab(c(`DRDID ATT` = att,
      `omega' dY` = sum(omega * dY),
      `wrong-sign weights` = sum((D == 1 & omega < 0) | (D == 0 & omega > 0)),
      `covariate imbalance, raw` = max(abs((colMeans(Xk[D == 1, ]) - colMeans(Xk[D == 0, ]))),
      `covariate imbalance, wtd` = max(abs(as.vector(a %*% Xk) / sds))), 3)

```

|                          | value   |
|--------------------------|---------|
| DRDID ATT                | 0.980   |
| omega' dY                | 0.980   |
| wrong-sign weights       | 151.000 |
| covariate imbalance, raw | 1.198   |
| covariate imbalance, wtd | 0.086   |

$\omega' \Delta Y$  reproduces the DRDID estimate exactly — the difference-in-differences is a signed weighting of observed outcome *changes*, and the diagnostics of this chapter apply to it. The weighting balances the covariate that drove selection (its standardised imbalance falls from 1.20 to 0.09). And 151 units carry a wrong-sign weight — a treated unit subtracted, or a control added — the §3 signature of a comparison reaching past common support.

### 14.5.1 Staggered designs: TWFE at the observation level

Negative weights in staggered designs are usually described at the level of two-by-two comparisons: TWFE averages many such comparisons, and some enter with a negative weight — the forbidden comparisons that use already-treated units as controls (de Chaisemartin and D'Haultfœuille 2020; Goodman-Bacon 2021). That is the effect-weight reading. The outcome-weight reading goes one level down, to the individual  $y_{it}$ .

TWFE is still FWL: residualise  $D_{it}$  on unit and time fixed effects to get  $\tilde{D}_{it}$ , and  $\hat{\beta} = \sum_{it} \omega_{it} y_{it}$  with  $\omega_{it} = \tilde{D}_{it} / \sum_{jt} \tilde{D}_{jt}^2$ . A treated observation with  $\omega_{it} < 0$  is a cell whose outcome is subtracted from the estimated effect — used as a control even though it is treated. An untreated observation with  $\omega_{it} > 0$  enters as though it were treated. Both are the §3 sign-flip, caused by the unit-and-time residualisation pushing  $\tilde{D}_{it}$  past zero: treated cells in late periods of early cohorts, where the fixed-effect prediction of  $D_{it}$  exceeds 1, and untreated cells in early periods of never-treated units, where it falls below 0.

**Simulation — staggered panel.** Three cohorts: never-treated (20%), early-treated at period 3 with effect  $\tau = 3$  (40%), and late-treated at period 6 with effect  $\tau = 1$  (40%), observed across  $T = 8$  periods;  $y_{it} = \alpha_i + \gamma_t + \tau_g D_{it} + \varepsilon_{it}$ . The high treated share makes the forbidden comparison visible.

```

library(fixest)
set.seed(5)
N <- 300; TT <- 8
g <- sample(c(0, 3, 6), N, TRUE, prob = c(.2, .4, .4))
panel <- data.frame(id = rep(1:N, each = TT), t = rep(1:TT, N))
panel$g <- g[panel$id]
panel$D <- as.integer(panel$g > 0 & panel$t >= panel$g)
tau <- ifelse(panel$g == 3, 3, ifelse(panel$g == 6, 1, 0))
alpha <- rnorm(N); gamma <- seq(0, 1.4, length.out = TT)
panel$y <- alpha[panel$id] + gamma[panel$t] + tau * panel$D + rnorm(nrow(panel), 0, 0.5)

Dtil <- resid(feols(D ~ 1 | id + t, data = panel))
omega_tw <- Dtil / sum(Dtil^2)
n1 <- sum(panel$D)
vtab(c(`true ATT`           = mean(tau[panel$D == 1]),
      `TWFE`               = coef(feols(y ~ D | id + t, panel))["D"],
      `omega' y`           = sum(omega_tw * panel$y),
      `treated, omega < 0` = sum(panel$D == 1 & omega_tw < 0),
      `untreated, omega > 0` = sum(panel$D == 0 & omega_tw > 0)), 3)

```

|                      | value   |
|----------------------|---------|
| true ATT             | 2.426   |
| TWFE                 | 1.915   |
| omega' y             | 1.915   |
| treated, omega < 0   | 384.000 |
| untreated, omega > 0 | 551.000 |

The TWFE coefficient sits well below the true ATT, and  $\sum \omega_{it} y_{it}$  reproduces it exactly. Hundreds of treated cells carry a negative weight: these are the early cohort's cells in periods 6–8, the ones where the unit-and-time prediction of  $D$  exceeds 1 and the residual  $\tilde{D}_{it}$  flips sign. Their outcomes are subtracted from  $\hat{\beta}$  — the observation-level signature of the Bacon forbidden comparison. (Untreated cells with  $\omega > 0$  also appear, a consequence of the same overshooting in the opposite direction; these are less interpretable as a diagnostic, since they arise from the fixed-effect structure on imbalanced data rather than from any comparison being used backwards.)

### 14.5.2 Imputation as a smoother

The imputation estimator of Borusyak, Jaravel, and Spiess (2024) fits the outcome model on untreated observations alone and imputes each treated cell's counterfactual. With unit and time fixed effects the imputed value is  $\hat{Y}_{it}(0) = \hat{\alpha}_i + \hat{\gamma}_t$ , the coefficients from  $Y_{jt} = \alpha_j + \gamma_t + \varepsilon_{jt}$  run on  $\{(j, t) : D_{jt} = 0\}$ .

That regression is a smoother:  $\hat{Y}(0) = SY$  with  $S = X(X'_0 X_0)^{-1} X'_0 P_0$ , where  $X_0$  and  $P_0$  restrict to untreated rows. The ATT averages the treated residuals,

$$\widehat{\text{ATT}} = \frac{1}{n_1} \sum_{D_{it}=1} (Y_{it} - \hat{Y}_{it}(0)) = a'(\mathbf{I} - S)Y = \omega'Y \quad (14.16)$$

with  $a_{it} = 1/n_1$  for treated cells and zero otherwise — the recipe of §4. Because the smoother draws only on untreated outcomes — column  $j$  of  $S$  is zero when  $D_j = 1$  — the weight on every treated cell reduces to  $\omega_{it} = 1/n_1$ : **each treated observation carries the same positive weight**. Untreated cells carry varying weights that sum to  $-1$ ; individual cells can be positive or negative, since the OLS fit on unbalanced untreated data couples the fixed effects.

```
unt <- panel$D == 0
X <- model.matrix(~ factor(id) + factor(t), data = panel)
X0 <- X[unt, ]
coefs <- solve(crossprod(X0), crossprod(X0, panel$y[unt]))
panel$y0hat <- as.vector(X %*% coefs)
att_imp <- mean(panel$y[panel$D == 1] - panel$y0hat[panel$D == 1])

a <- ifelse(panel$D == 1, 1 / n1, 0)
v <- solve(crossprod(X0), crossprod(X, a))
Sta <- numeric(nrow(panel)); Sta[unt] <- as.vector(X0 %*% v)
omega_imp <- a - Sta

vtab(c(`imputation ATT`           = att_imp,
      `omega' y`                 = sum(omega_imp * panel$y),
      `treated, omega < 0`       = sum(panel$D == 1 & omega_imp < 0),
      `untreated, omega > 0`     = sum(!panel$D & omega_imp > 0),
      `max |treated wt - 1/n1|` = max(abs(omega_imp[panel$D == 1] - 1/n1))), 3)
```

|                        | value   |
|------------------------|---------|
| imputation ATT         | 2.408   |
| omega' y               | 2.408   |
| treated, omega < 0     | 0.000   |
| untreated, omega > 0   | 551.000 |
| max  treated wt - 1/n1 | 0.000   |

The imputation ATT recovers the true value,  $\omega'Y$  reproduces it, and no treated cell carries a negative weight — the key diagnostic. The maximum deviation of a treated weight from  $1/n_1$  is at machine precision, confirming the structural guarantee. Untreated cells with  $\omega > 0$  still appear (as in TWFE), but these reflect the coupling in the fixed-effect estimation, not a comparison running backwards.

```
cols <- c("grey50", teal, rust)
pchs <- c(16, 17, 15)
cell_tw <- tapply(omega_tw, list(panel$g, panel$t), sum)
cell_imp <- tapply(omega_imp, list(panel$g, panel$t), sum)
ylim <- range(cell_tw, cell_imp)
```

```

plot_cells <- function(cells, main) {
  plot(NA, xlim = c(0.5, 8.5), ylim = ylim, xlab = "period",
       ylab = "total cell weight", main = main, xaxt = "n")
  axis(1, at = 1:8)
  abline(h = 0, lty = 3, col = "grey55")
  for (k in 1:3) {
    lines(1:8, cells[k, ], col = cols[k], lwd = 2)
    points(1:8, cells[k, ], pch = pchs[k], col = cols[k], cex = 1.4)
  }
}

op <- par(mfrow = c(1, 2), mar = c(4, 4.5, 2.5, 1))
plot_cells(cell_tw, "TWFE")
legend("bottomleft", legend = c("never (g=0)", "early (g=3)", "late (g=6)"),
      pch = pchs, col = cols, lty = 1, lwd = 2, bty = "n", cex = 0.85)
plot_cells(cell_imp, "Imputation")

```




Figure 14.6: Cell-level outcome weights in the staggered panel, summed within each cohort-period cell. TWFE (left): the early cohort's treated cells (teal triangles) turn negative in periods 6–8 — those outcomes are subtracted from  $\hat{\beta}$ , the observation-level signature of the forbidden comparison. Imputation (right): every treated cell stays positive; no treated observation can count as a control.

```
par(op)
```

The two panels locate the structural difference. TWFE residualises the treatment indicator over the full panel, so when the early cohort's treated cells sit in a period where many units are treated, their residual turns negative and their outcomes count as controls. The imputation estimator's smoother, fitted on untreated cells alone, cannot reach treated outcomes, so no treated cell can count as a

control — the forbidden-comparison problem that defines staggered TWFE bias is structurally absent.

The wrong-sign diagnostic in a panel is by observation  $(i, t)$ , not by unit. A single unit can have  $\omega_{it} > 0$  in some periods and  $\omega_{it} < 0$  in others — in the simulation above, every early-cohort unit is positive in periods 3–5 and negative in periods 6–8. This changes the response. In the cross-sectional case of §3, the fix was to trim wrong-sign observations and refit. In a panel, trimming means dropping specific  $(i, t)$  cells, which removes observations that anchor that unit’s fixed effect and distorts the estimates for its remaining cells. The response in the staggered setting is not to trim but to switch estimators — imputation, Callaway and Sant’Anna, or Sun and Abraham — that avoid the forbidden comparison by construction.

### 14.5.3 Implied versus chosen

One distinction cuts across everything above. OLS, AIPW, DML, TWFE, and the imputation estimator all *imply* weights — they fall out of the estimation and are diagnosed afterward. Synthetic controls and synthetic difference-in-differences *choose* weights as part of the design, selecting  $\omega$  to balance pre-treatment outcomes or covariates. For those estimators  $\hat{\tau} = \omega'Y$  is the definition, not a decomposition; the balance and extrapolation checks of this chapter still apply, but what they diagnose is the designer’s choice rather than a side effect of the estimator.

A closing caution. These weights diagnose *estimation* — which comparisons a difference-in-differences rests on, and whether it leans on non-overlapping or forbidden ones — not *identification*. Parallel trends is an assumption about the treated group’s unobserved untreated path, and no weighting of observed outcomes can test it: a negative weight and a failure of parallel trends are independent, either able to occur without the other. Pre-trends are probed separately, by the event-study leads.

## 14.6 Summary

1. **A signed weighting of the outcomes.** FWL writes  $\hat{\beta} = \sum_i \omega_i y_i$  with  $\omega_i = \tilde{V}_i / \sum_j \tilde{V}_j^2$  (Equation 14.2): a signed, computable weighting of the observed outcomes, different from the effect weights  $w_i = \tilde{V}_i V_i$ .
2. **Two constraints.** The weights sum to zero and the weighted regressor sums to one:  $\sum_i \omega_i = 0$  and  $\sum_i \omega_i V_i = 1$ .
3. **The sign shows direction.** Because  $\omega_i$  is proportional to  $\tilde{V}_i$ , not to  $\tilde{V}_i^2$ , it carries a sign: it shows which rows push  $\hat{\beta}$  up and which pull it down, and, under controls with thin overlap, which enter with the sign their raw position would not predict.
4. **The same weights beyond OLS.** Any estimator that can be written  $\hat{\tau} = \omega'Y$  has outcome weights, so the same object and its checks reach AIPW, double machine learning, and causal forests (Knaus 2024). PLR’s denominator  $\tilde{D}'\tilde{D}$  reveals its estimand as the VWATE (Equation 14.12), not the ATE; and the sum-to-zero of OLS need not become a sum to  $\pm 1$  within the treated and control groups.
5. **From difference in means to difference-in-differences.** Regressed on the treatment alone,  $\hat{\beta}$  is the difference in group means; a  $2 \times 2$  difference-in-differences is a signed contrast of four cell means. With covariates, a doubly-robust DiD weights the outcome changes,  $\widehat{ATT} = \omega'\Delta Y$  (Sant’Anna and Zhao 2020), open to the same checks.

6. **Staggered designs and imputation.** In staggered TWFE the outcome weights flag both treated cells used as controls and untreated cells used as treated — the observation-level version of the Bacon (2021) forbidden comparisons. The imputation estimator of Borusyak, Jaravel, and Spiess (2024) is a smoother on untreated outcomes (Equation 14.16) and cannot assign negative weight to any treated cell: every treated cell carries the same weight  $1/n_1$ , so the forbidden-comparison diagnostic is clean by construction.

# References

This guide draws on a handful of textbooks, the original papers behind the estimators discussed, and package documentation. The list below collects the works cited across the chapters.

## Textbooks and lecture notes

- Baltagi, B. H. *Econometric Analysis of Panel Data*. Wiley.
- Baum, C. F. *An Introduction to Modern Econometrics Using Stata*. Stata Press.
- Davidson, R., and J. G. MacKinnon. *Econometric Theory and Methods*. Oxford University Press.
- Greene, W. H. *Econometric Analysis*. Prentice Hall.
- Hamilton, J. D. (1994). *Time Series Analysis*. Princeton University Press.
- Wooldridge, J. M. *Econometric Analysis of Cross Section and Panel Data*. MIT Press.

## Core papers

- Anderson, T. W., and C. Hsiao (1981). “Estimation of Dynamic Models with Error Components.” *Journal of the American Statistical Association* 76: 598–606.
- Andrews, D. W. K., M. J. Moreira, and J. H. Stock (2006). “Optimal Two-Sided Invariant Tests for Instrumental Variables Regression.” *Econometrica* 74(3): 715–752.
- Angrist, J. D. (1998). “Estimating the Labor Market Impact of Voluntary Military Service Using Social Security Data on Military Applicants.” *Econometrica* 66(2): 249–288.
- Arellano, M., and S. Bond (1991). “Some Tests of Specification for Panel Data.” *Review of Economic Studies* 58: 277–297.
- Baltagi, B. H., and D. Levin (1992). “Cigarette Taxation: Raising Revenues and Reducing Consumption.” *Structural Change and Economic Dynamics* 3(2): 321–335.
- Blundell, R., and S. Bond (1998). “Initial Conditions and Moment Restrictions in Dynamic Panel Data Models.” *Journal of Econometrics* 87: 115–143.
- Cameron, A. C., and P. K. Trivedi (2005). *Microeconometrics: Methods and Applications*. Cambridge University Press.
- Chattopadhyay, A., and J. R. Zubizarreta (2023). “On the implied weights of linear regression for causal inference.” *Biometrika* 110(3): 615–629. R package `lmw`.
- Cox, D. R. (1972). “Regression Models and Life-Tables.” *Journal of the Royal Statistical Society B* 34: 187–220.
- Crump, R. K., V. J. Hotz, G. W. Imbens, and O. A. Mitnik (2009). “Dealing with Limited Overlap in Estimation of Average Treatment Effects.” *Biometrika* 96(1): 187–199.
- de Chaisemartin, C., and X. D’Haultfoeuille (2020). “Two-Way Fixed Effects Estimators with Heterogeneous Treatment Effects.” *American Economic Review* 110(9): 2964–2996.

- Donald, S. G., and K. Lang (2007). “Inference with Difference-in-Differences and Other Panel Data.” *Review of Economics and Statistics* 89(2): 221–233.
- Frisch, R., and F. V. Waugh (1933). “Partial Time Regressions as Compared with Individual Trends.” *Econometrica* 1(4): 387–401.
- Goodman-Bacon, A. (2021). “Difference-in-Differences with Variation in Treatment Timing.” *Journal of Econometrics* 225(2): 254–277.
- Grogger, J. T., and R. T. Carson (1991). “Models for Truncated Counts.” *Journal of Applied Econometrics* 6(3): 225–238.
- Hazlett, C., and T. Shinkre (2024). “Demystifying and Avoiding the OLS ‘Weighting Problem’: Unmodeled Heterogeneity and Straightforward Solutions.” arXiv:2403.03299.
- Heckman, J. J. (1979). “Sample Selection Bias as a Specification Error.” *Econometrica* 47: 153–161.
- Judson, R. A., and A. L. Owen (1999). “Estimating Dynamic Panel Data Models: A Guide for Macroeconomists.” *Economics Letters* 65: 9–15.
- Knaus, M. C. (2024). “Treatment Effect Estimators as Weighted Outcomes.” arXiv:2411.11559. R package `OutcomeWeights`.
- Li, F., K. L. Morgan, and A. M. Zaslavsky (2018). “Balancing Covariates via Propensity Score Weighting.” *Journal of the American Statistical Association* 113(521): 390–400.
- Lovell, M. C. (1963). “Seasonal Adjustment of Economic Time Series and Multiple Regression Analysis.” *Journal of the American Statistical Association* 58: 993–1010.
- Moreira, M. J. (2003). “A Conditional Likelihood Ratio Test for Structural Models.” *Econometrica* 71: 1027–1048.
- Moulton, B. R. (1990). “An Illustration of a Pitfall in Estimating the Effects of Aggregate Variables on Micro Units.” *Review of Economics and Statistics* 72(2): 334–338.
- Mundlak, Y. (1978). “On the Pooling of Time Series and Cross Section Data.” *Econometrica* 46(1): 69–85.
- Newey, W. K., and K. D. West (1987). “A Simple, Positive Semi-Definite, Heteroskedasticity and Autocorrelation Consistent Covariance Matrix.” *Econometrica* 55: 703–708.
- Newey, W. K., and K. D. West (1994). “Automatic Lag Selection in Covariance Matrix Estimation.” *Review of Economic Studies* 61: 631–653.
- Nickell, S. (1981). “Biases in Dynamic Models with Fixed Effects.” *Econometrica* 49: 1417–1426.
- Sargan, J. D. (1958). “The Estimation of Economic Relationships Using Instrumental Variables.” *Econometrica* 26: 393–415.
- Singer, J. D., and J. B. Willett (1993). “It’s About Time: Using Discrete-Time Survival Analysis to Study Duration and the Timing of Events.” *Journal of Educational Statistics* 18: 155–195.
- Słoczyński, T. (2022). “Interpreting OLS Estimands When Treatment Effects Are Heterogeneous: Smaller Groups Get Larger Weights.” *Review of Economics and Statistics* 104(3): 501–509.
- Stock, J. H., and M. Yogo (2005). “Testing for Weak Instruments in Linear IV Regression.” In *Identification and Inference for Econometric Models*, ed. Andrews and Stock. Cambridge University Press.
- Wooldridge, J. M. (1999). “Distribution-Free Estimation of Some Nonlinear Panel Data Models.” *Journal of Econometrics* 90: 77–97.

## Missing data

- Rubin, D. B. (1987). *Multiple Imputation for Nonresponse in Surveys*. Wiley.
- van Buuren, S. (2018). *Flexible Imputation of Missing Data*. 2nd ed. CRC Press. (MICE.)

## Software

- R Core Team. *R: A Language and Environment for Statistical Computing*. R packages used in the code examples include `AER`, `ivreg`, `gmm`, `survival`, `MASS`, `mvtnorm`, `strucchange`, and `car`.
- StataCorp. *Stata Statistical Software*.

